Working in the (Azure Data) Factory the right way

Riccardo Perico

rperico@lucient.com

Riccardo Perico

BI & Power BI Engineer @ Lucient

PBIUG Italy Co-leader

Microsoft Data Platform MVP

Linkedin: https://linkedin.com/in/riccardo-perico

Email: rperico@lucient.com

Slideshare: https://www.slideshare.net/RiccardoPerico

GitHub: https://github.com/R1k91

Partners

UNIVERSITÀ POLITECNICA DELLE MARCHE

Where are we? We're in the factory

*"**Big data** requires a service that can **orchestrate** and **operationalize** processes to refine these **enormous stores of raw data** into actionable business insights. Azure Data Factory is a **managed cloud service** that's built for these complex hybrid extract-transform-load (**ETL**), extract-load-transform (**ELT**), and **data integration** projects."*

What about this session?

Starting to work in ADF it's easy but…

You'll learn by doing especially best practices and patterns

Let me share some lessons learnt from by background

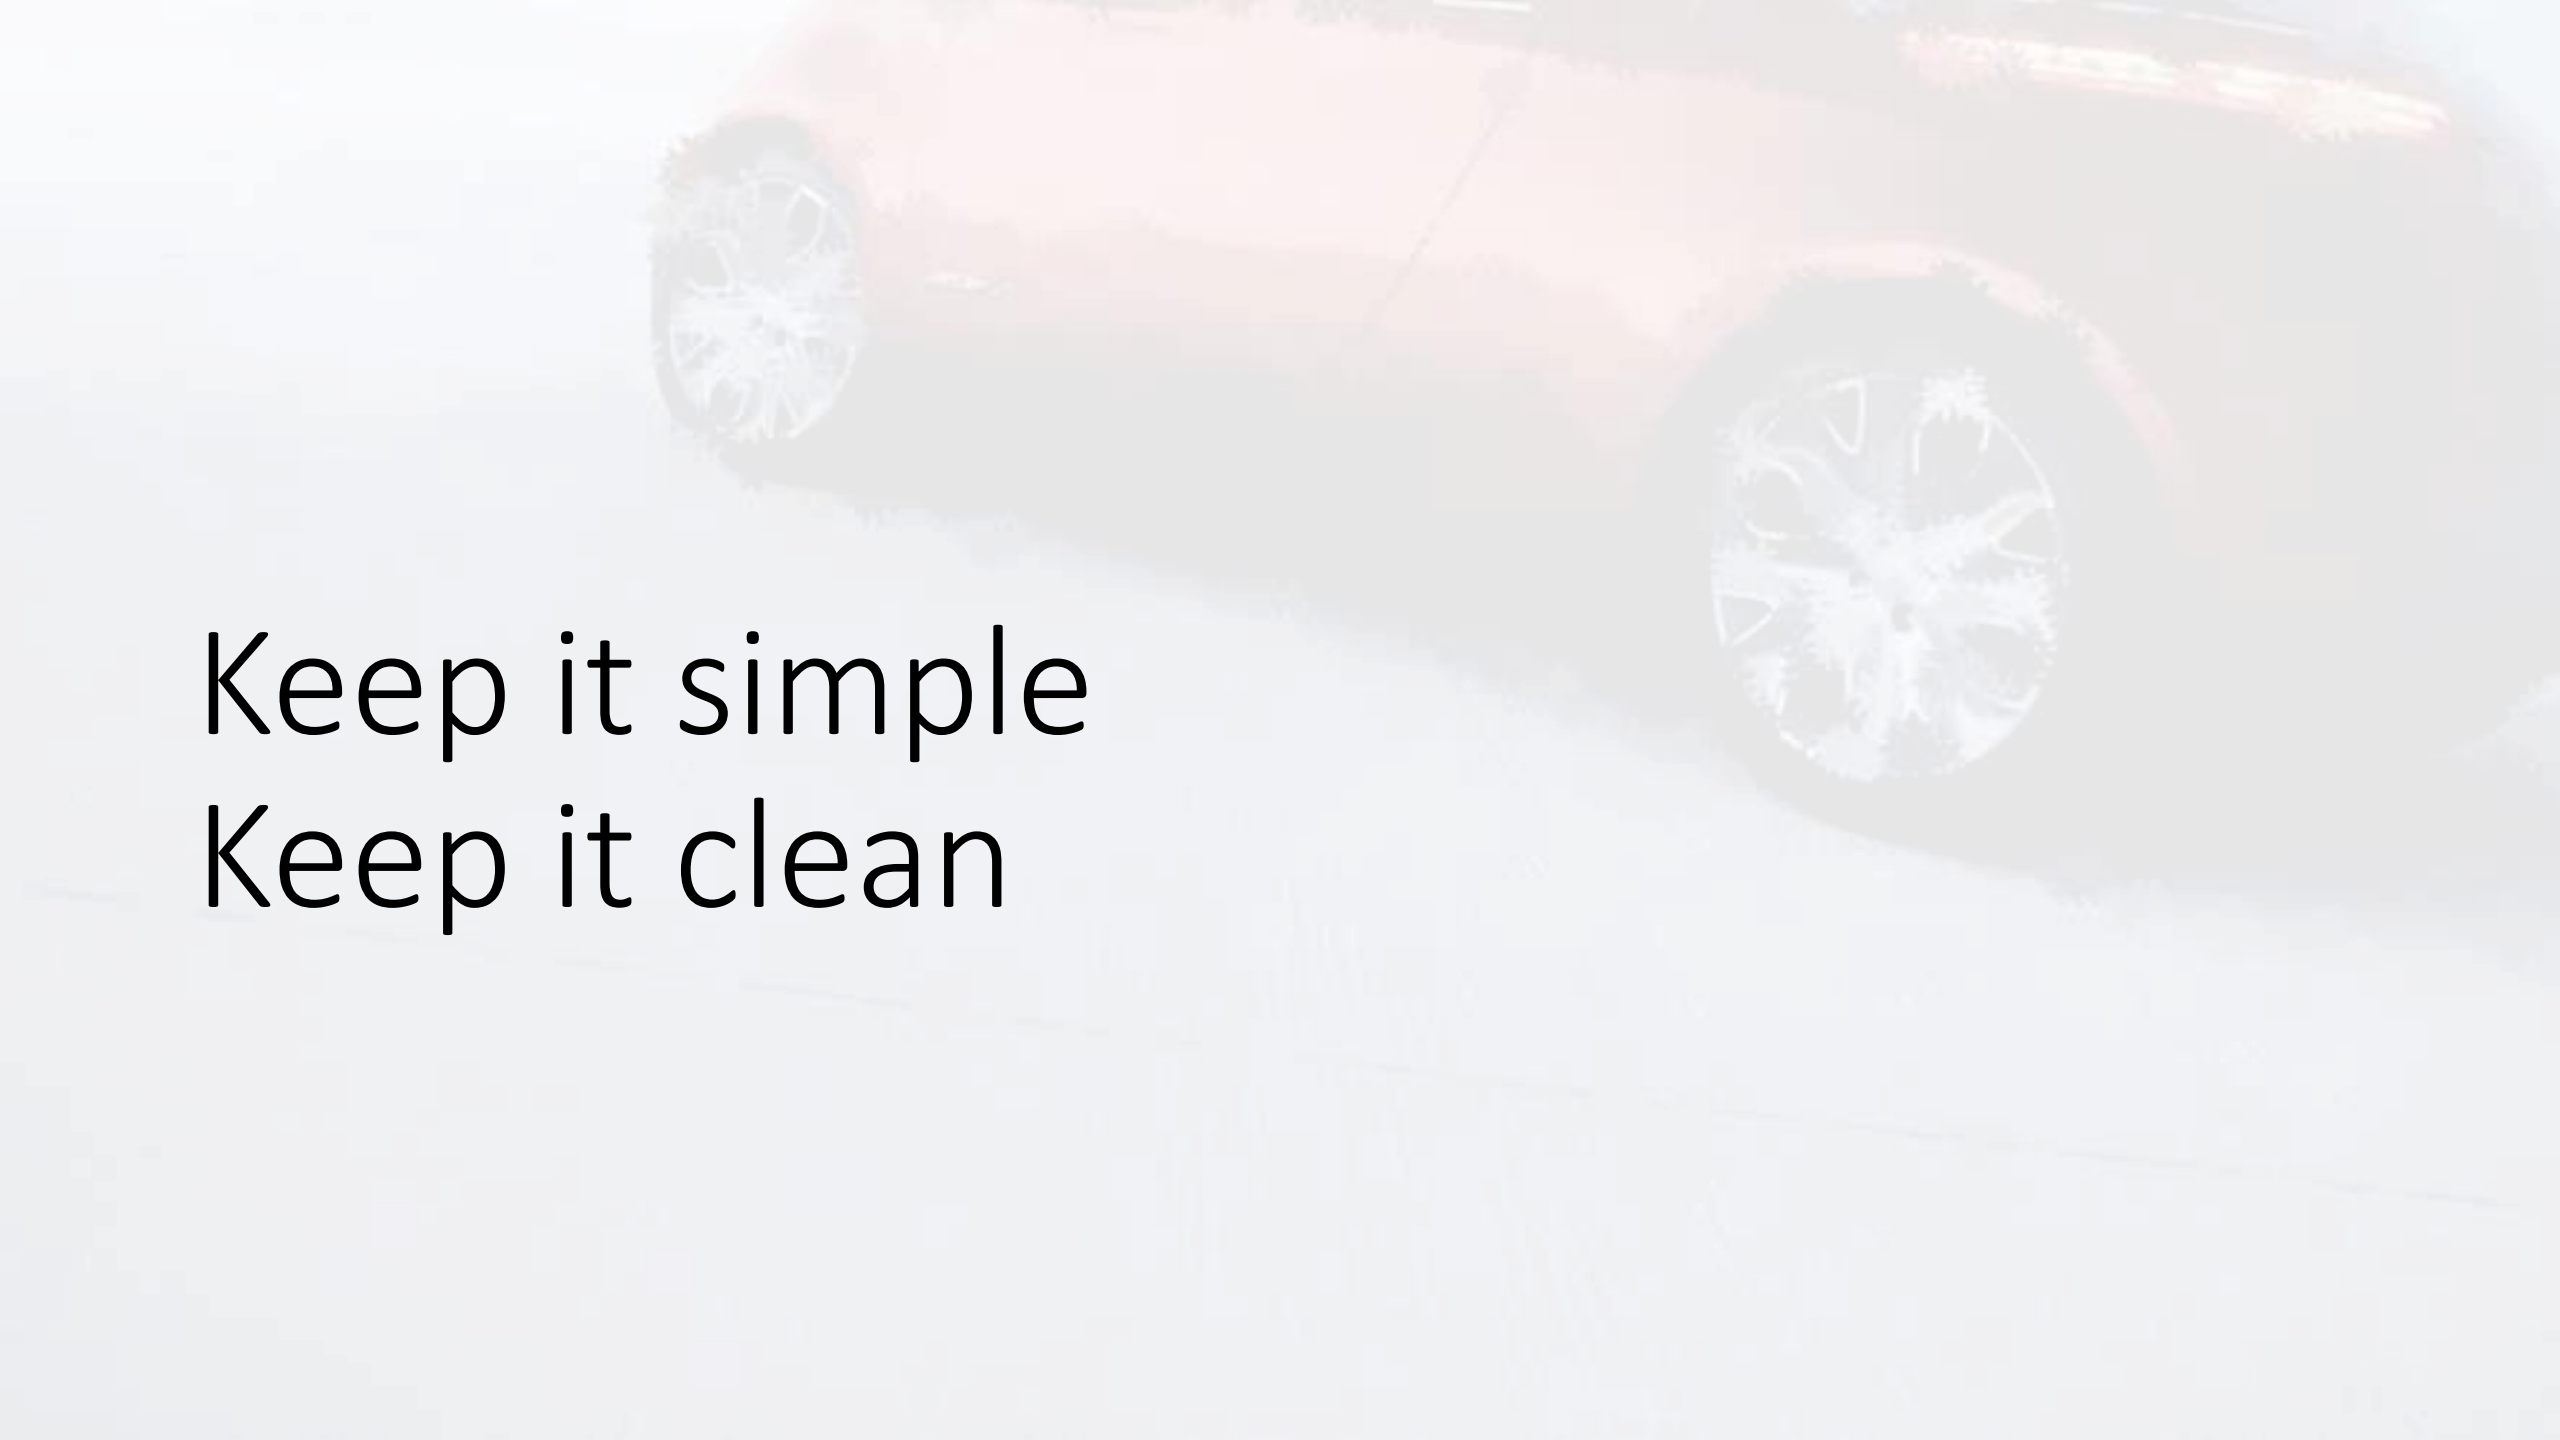

Keep it simple
Keep it clean

General rules

- Give objects (all the objects!!!) a meaningful name
- Use a “pipeline approach”
- Organize objects accordingly
- …and many others

Amarcord

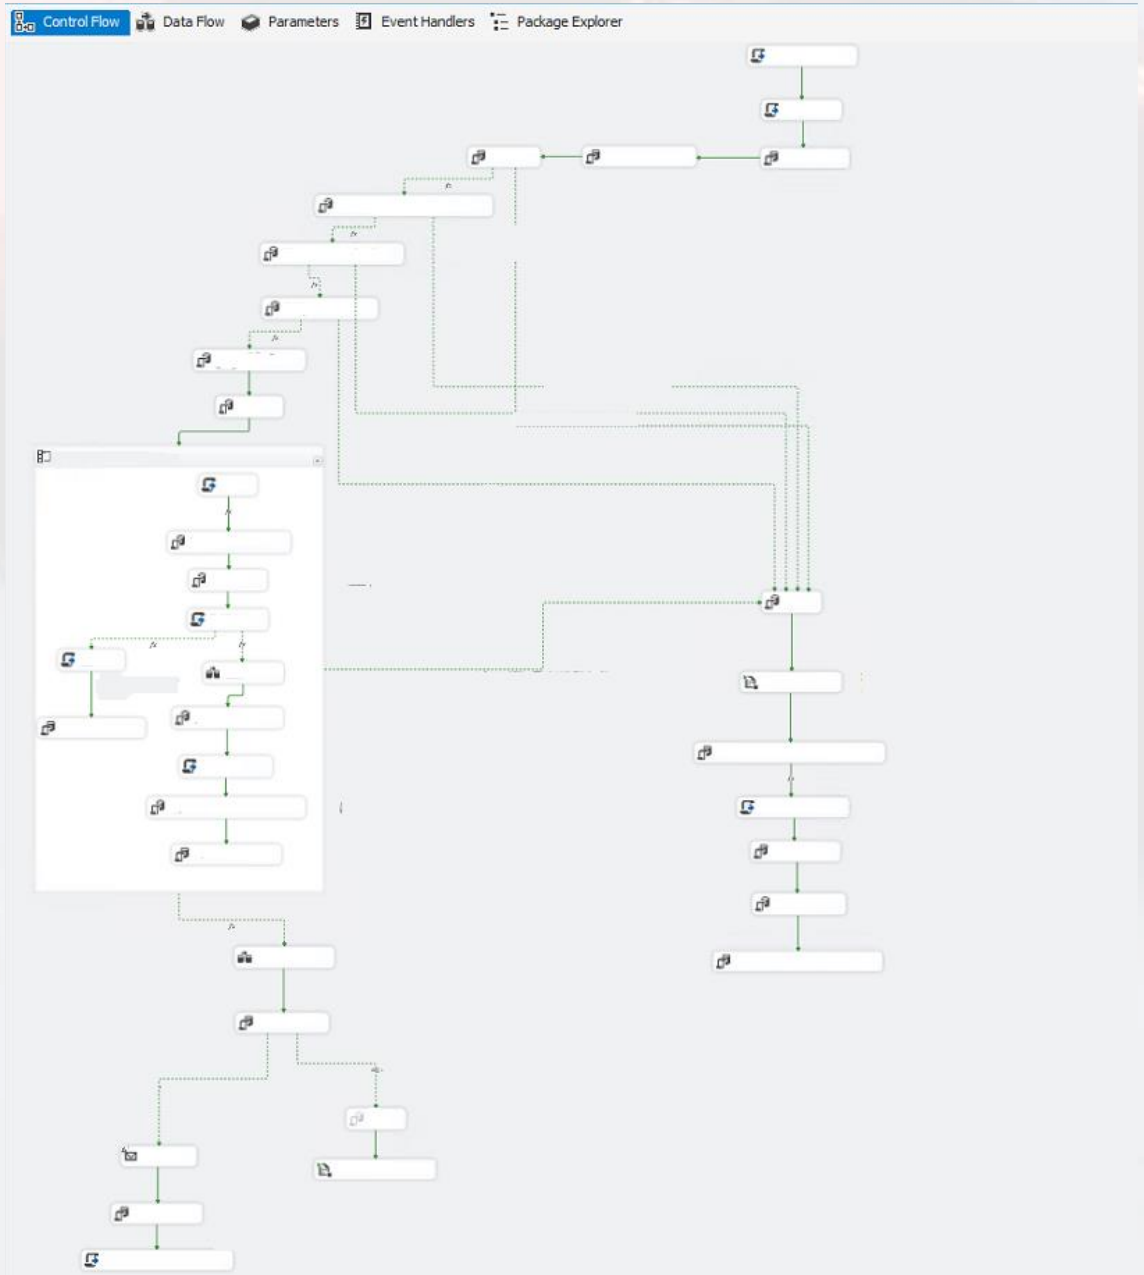
Control Flow
Data Flow
Parameters
Event Handlers
Package Explorer

New tools same pitfalls

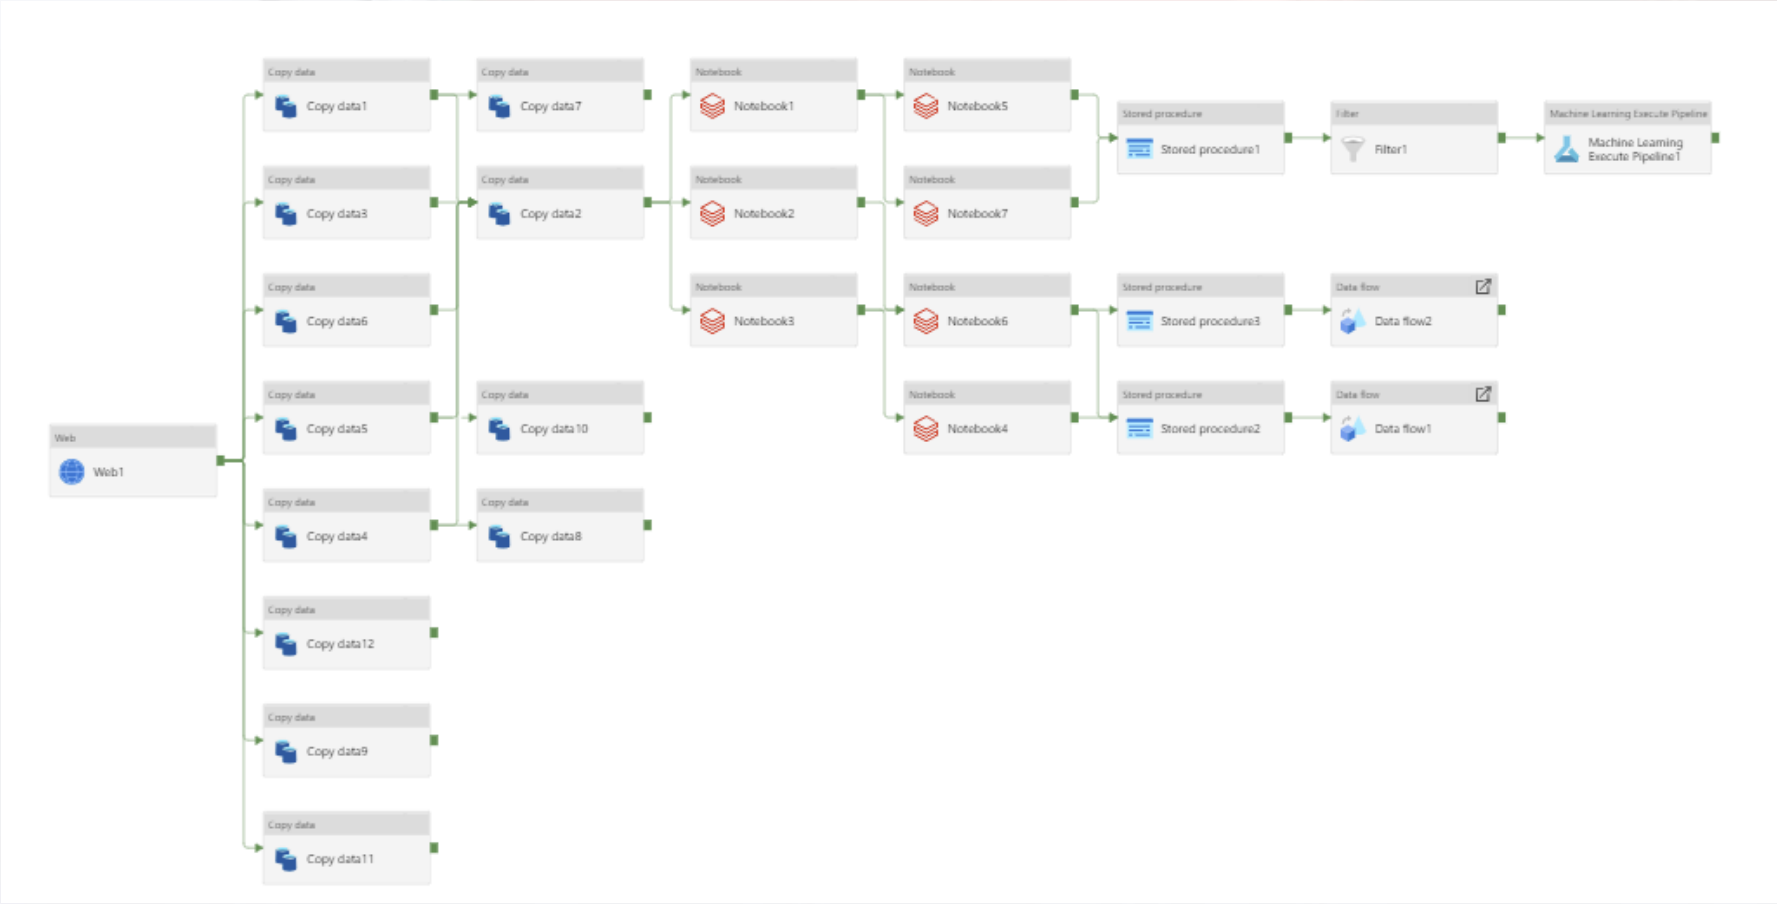

“Pipeline approach”: it flows in one way

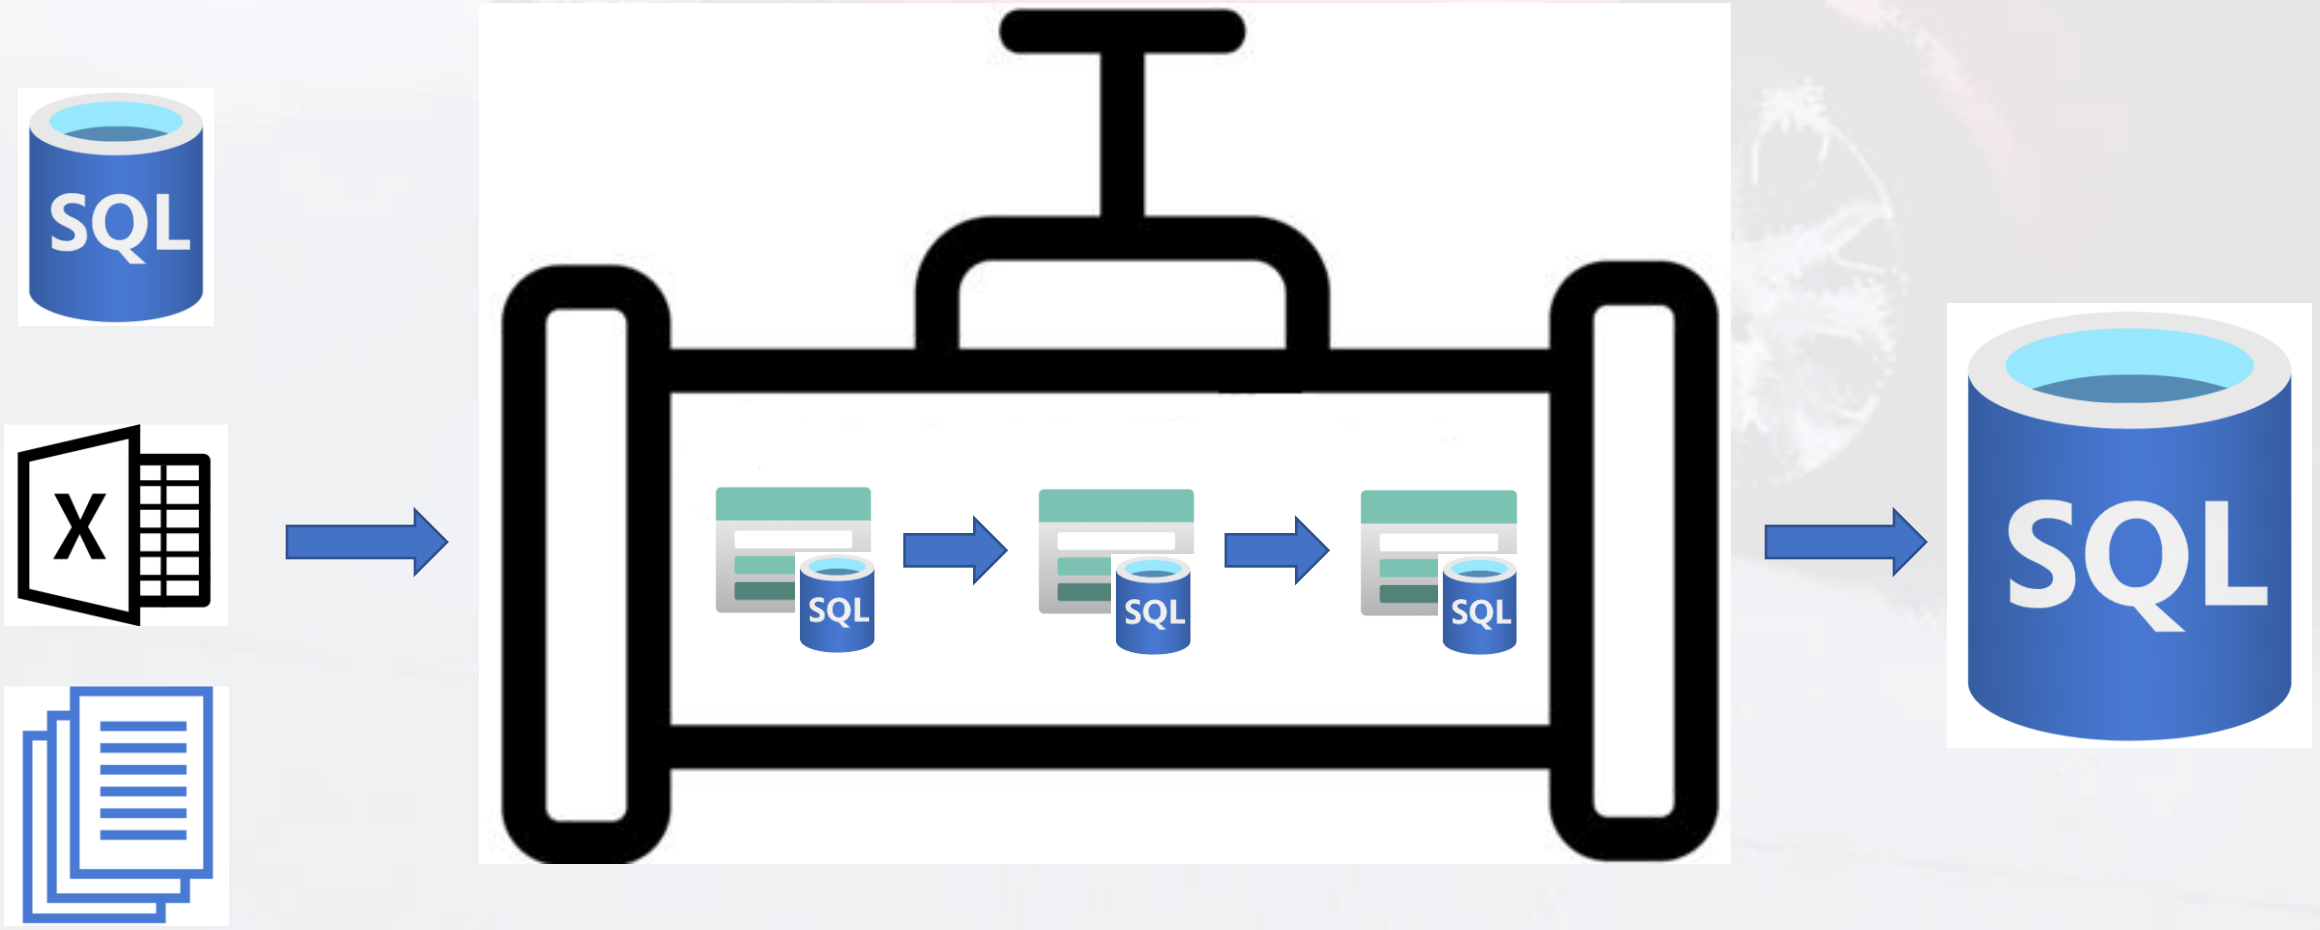

Pipeline approach: not that chaos… please

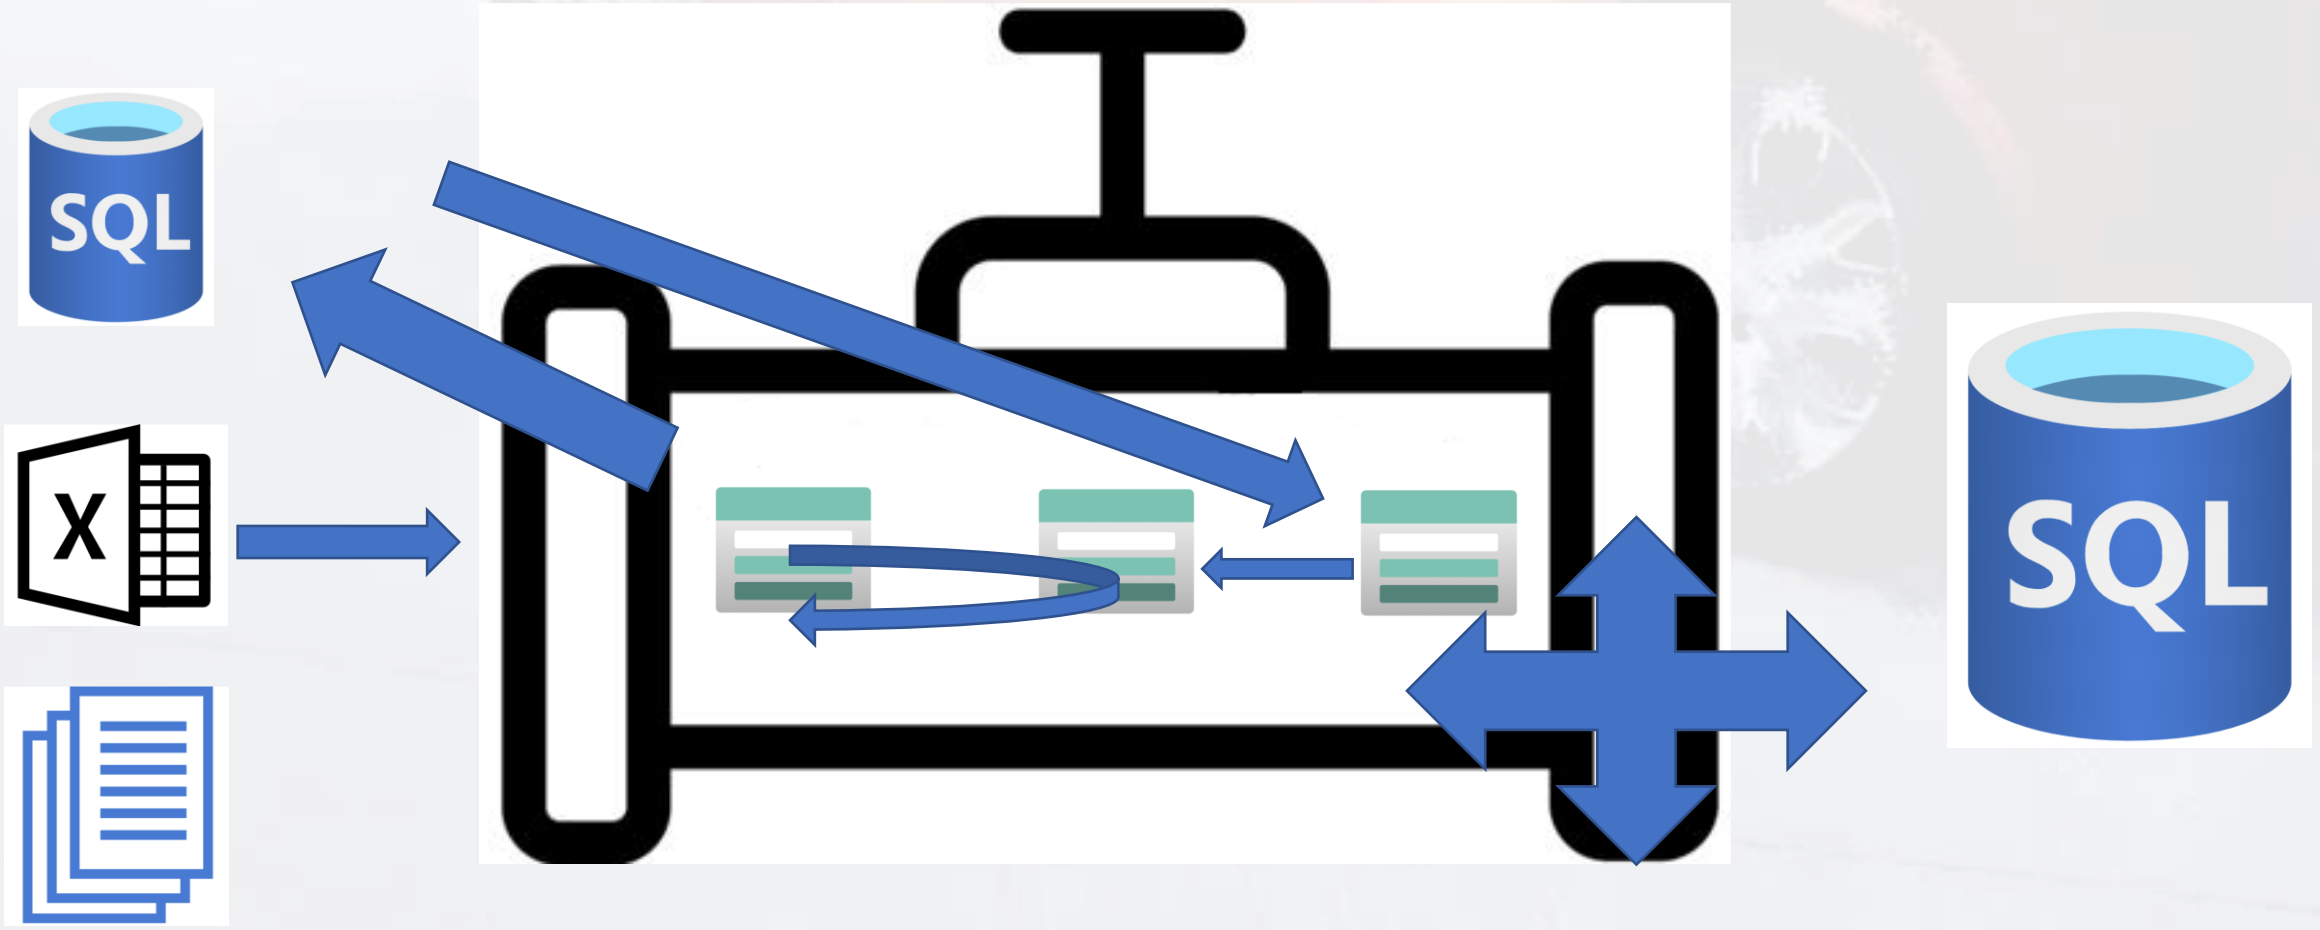

Keep it clean

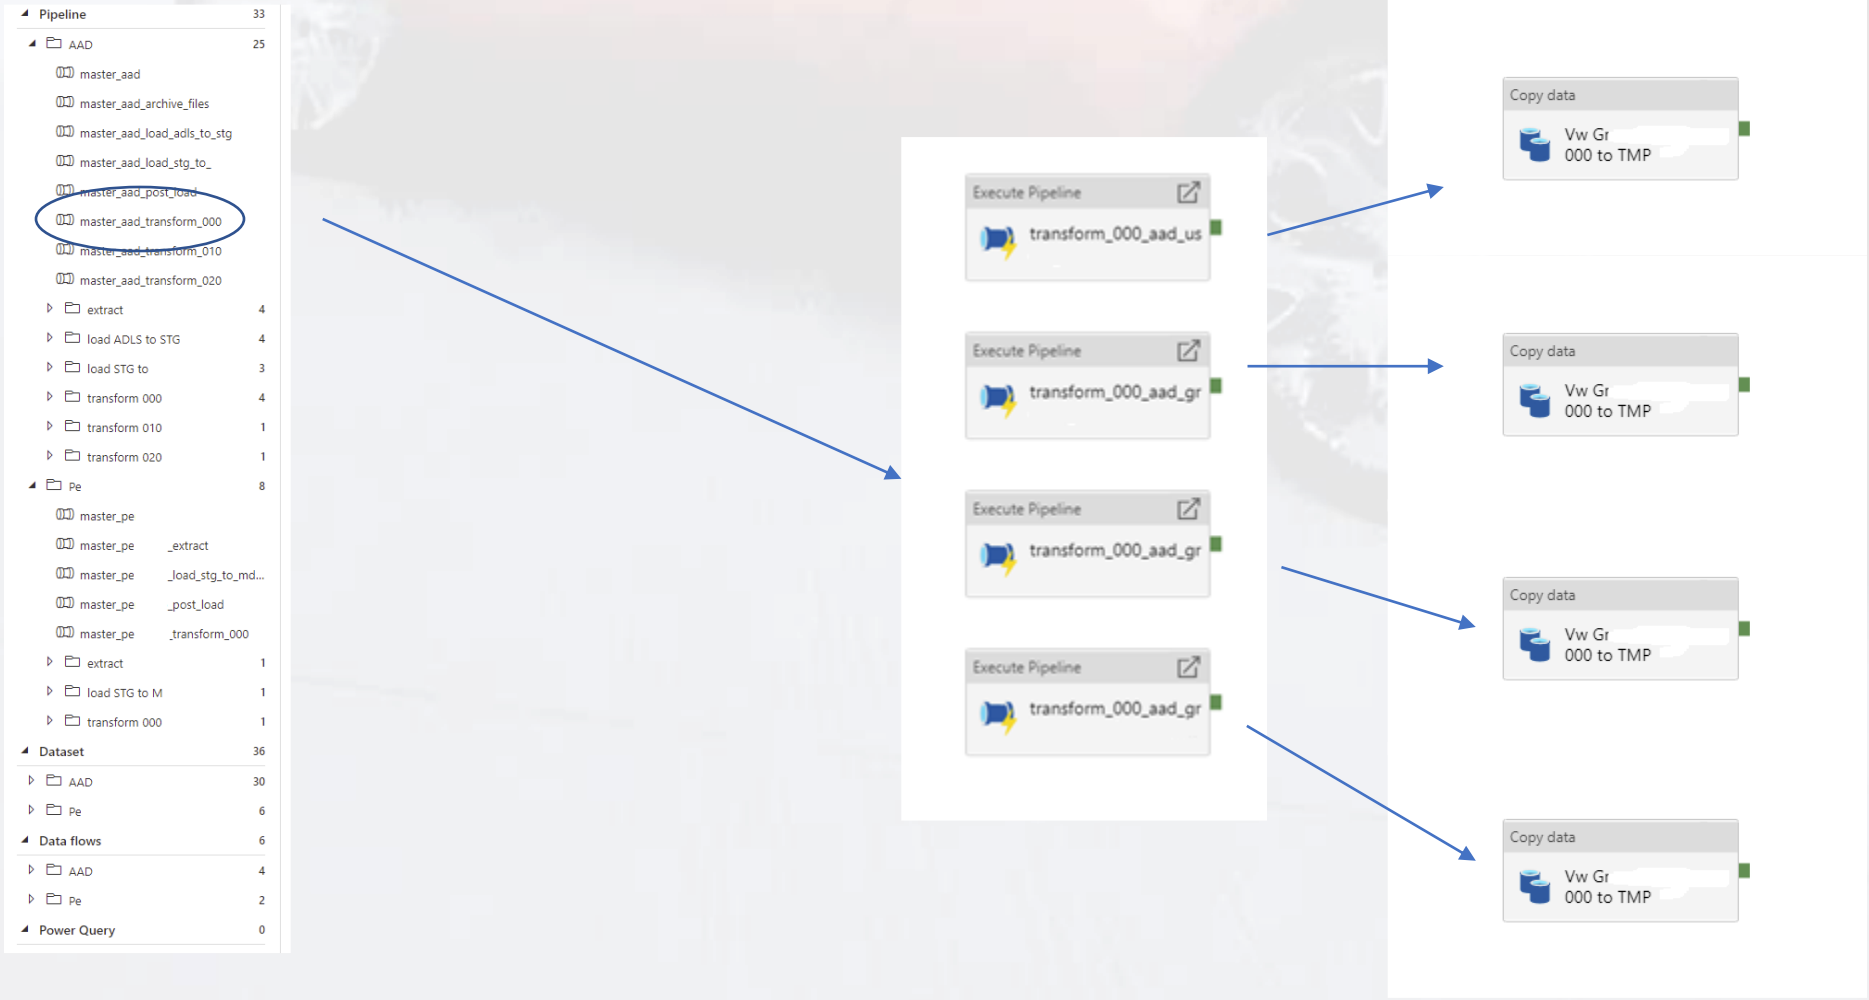

Global Parameters

Before and After

Parameter Name	Parameter Value
pContainerName	myContainer
pArchiveFiles	True

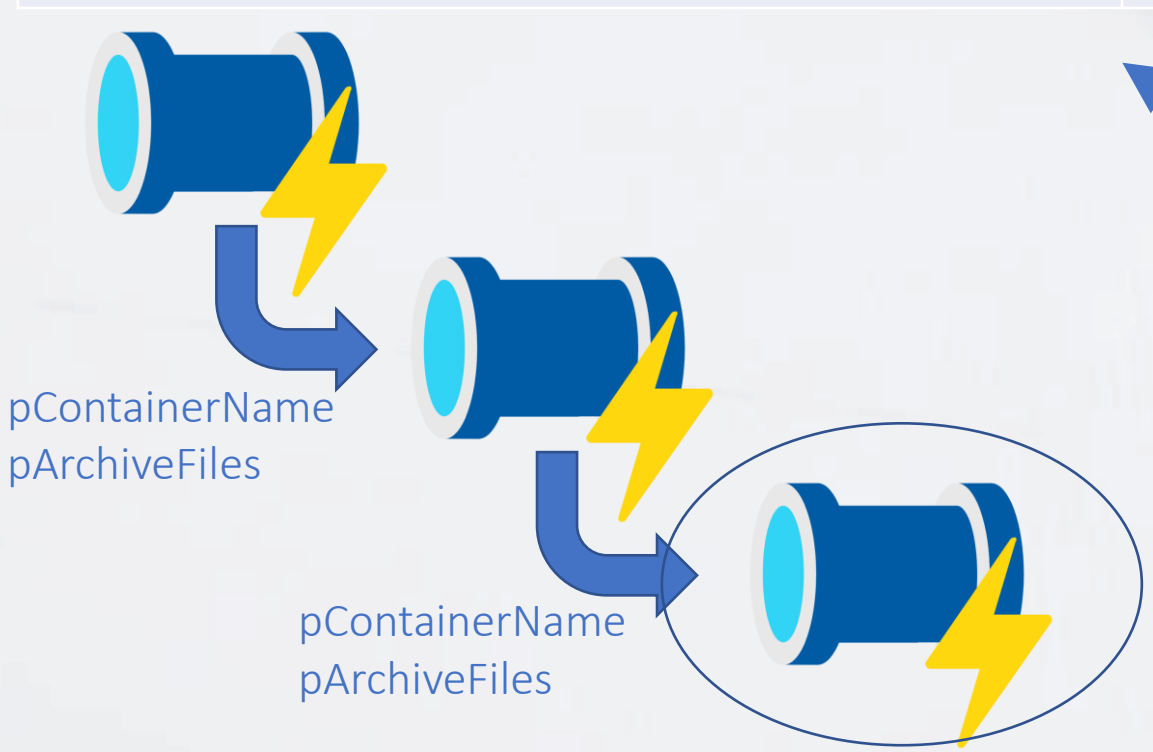

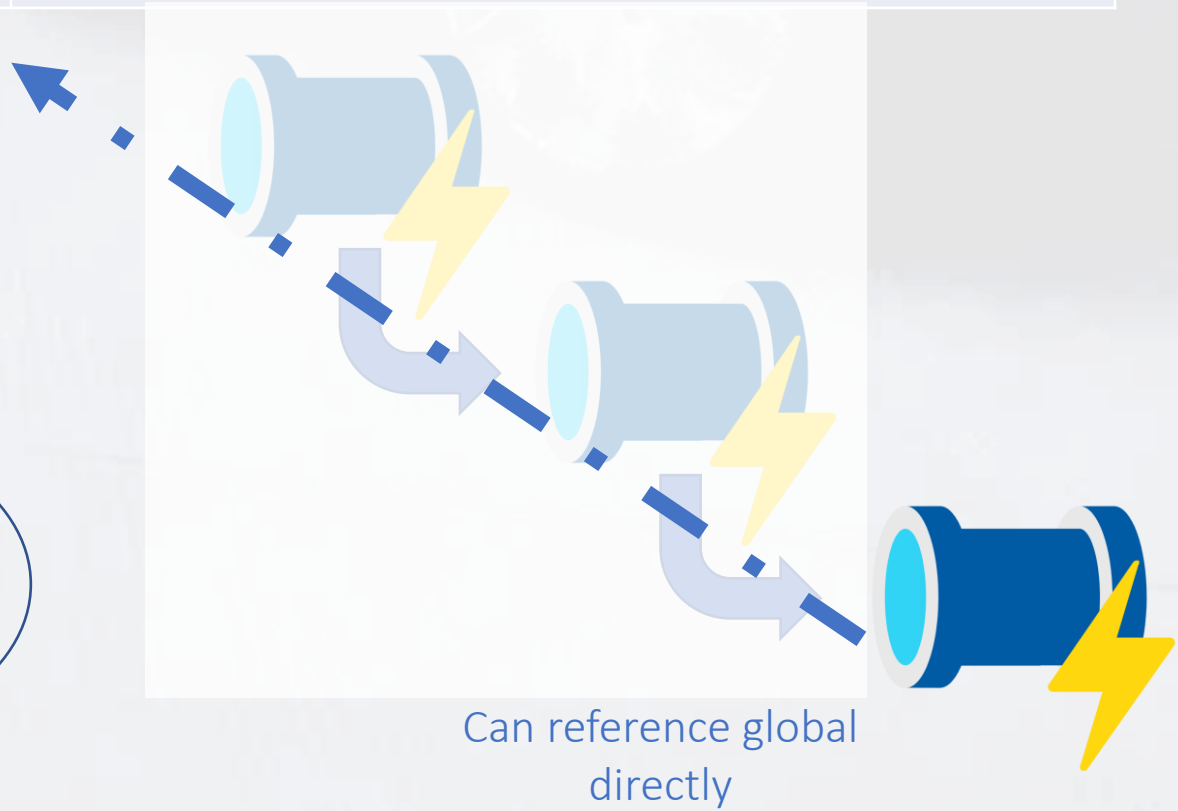

One drawback

- Only pipelines can use them

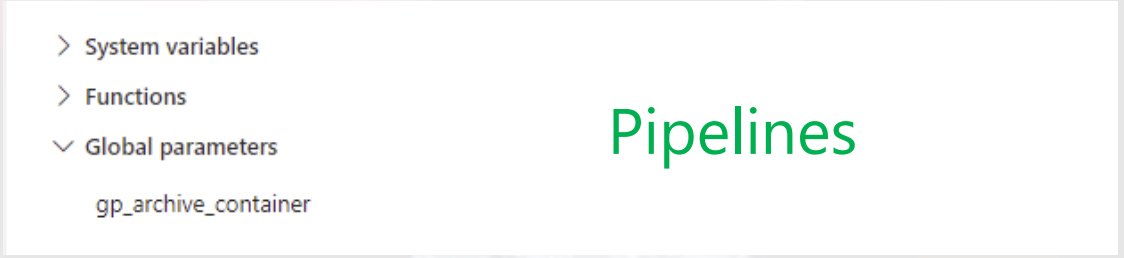

Global Parameters

Demo

Metadata Driven Approach

What is a Metadata-driven approach?

- Don’t do it manually
- Make it flexible and dynamic
- Make it configurable

Is there anything out-of-the-box?

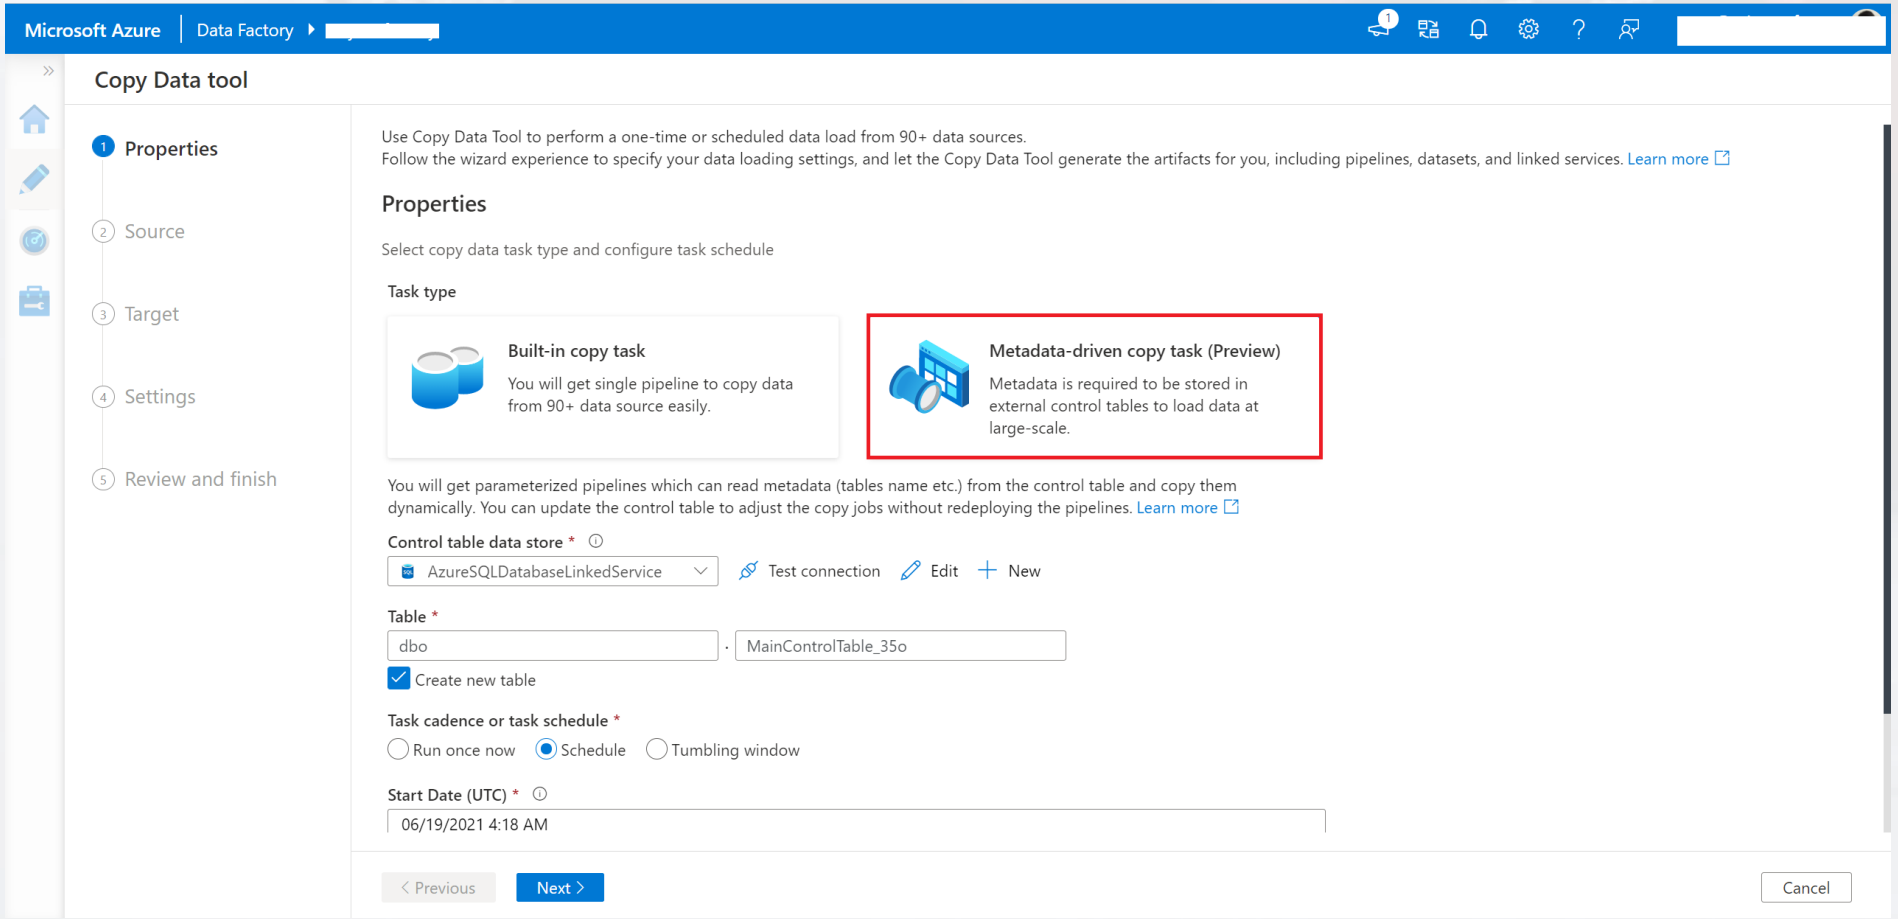

Metadata-driven Approach

Demo

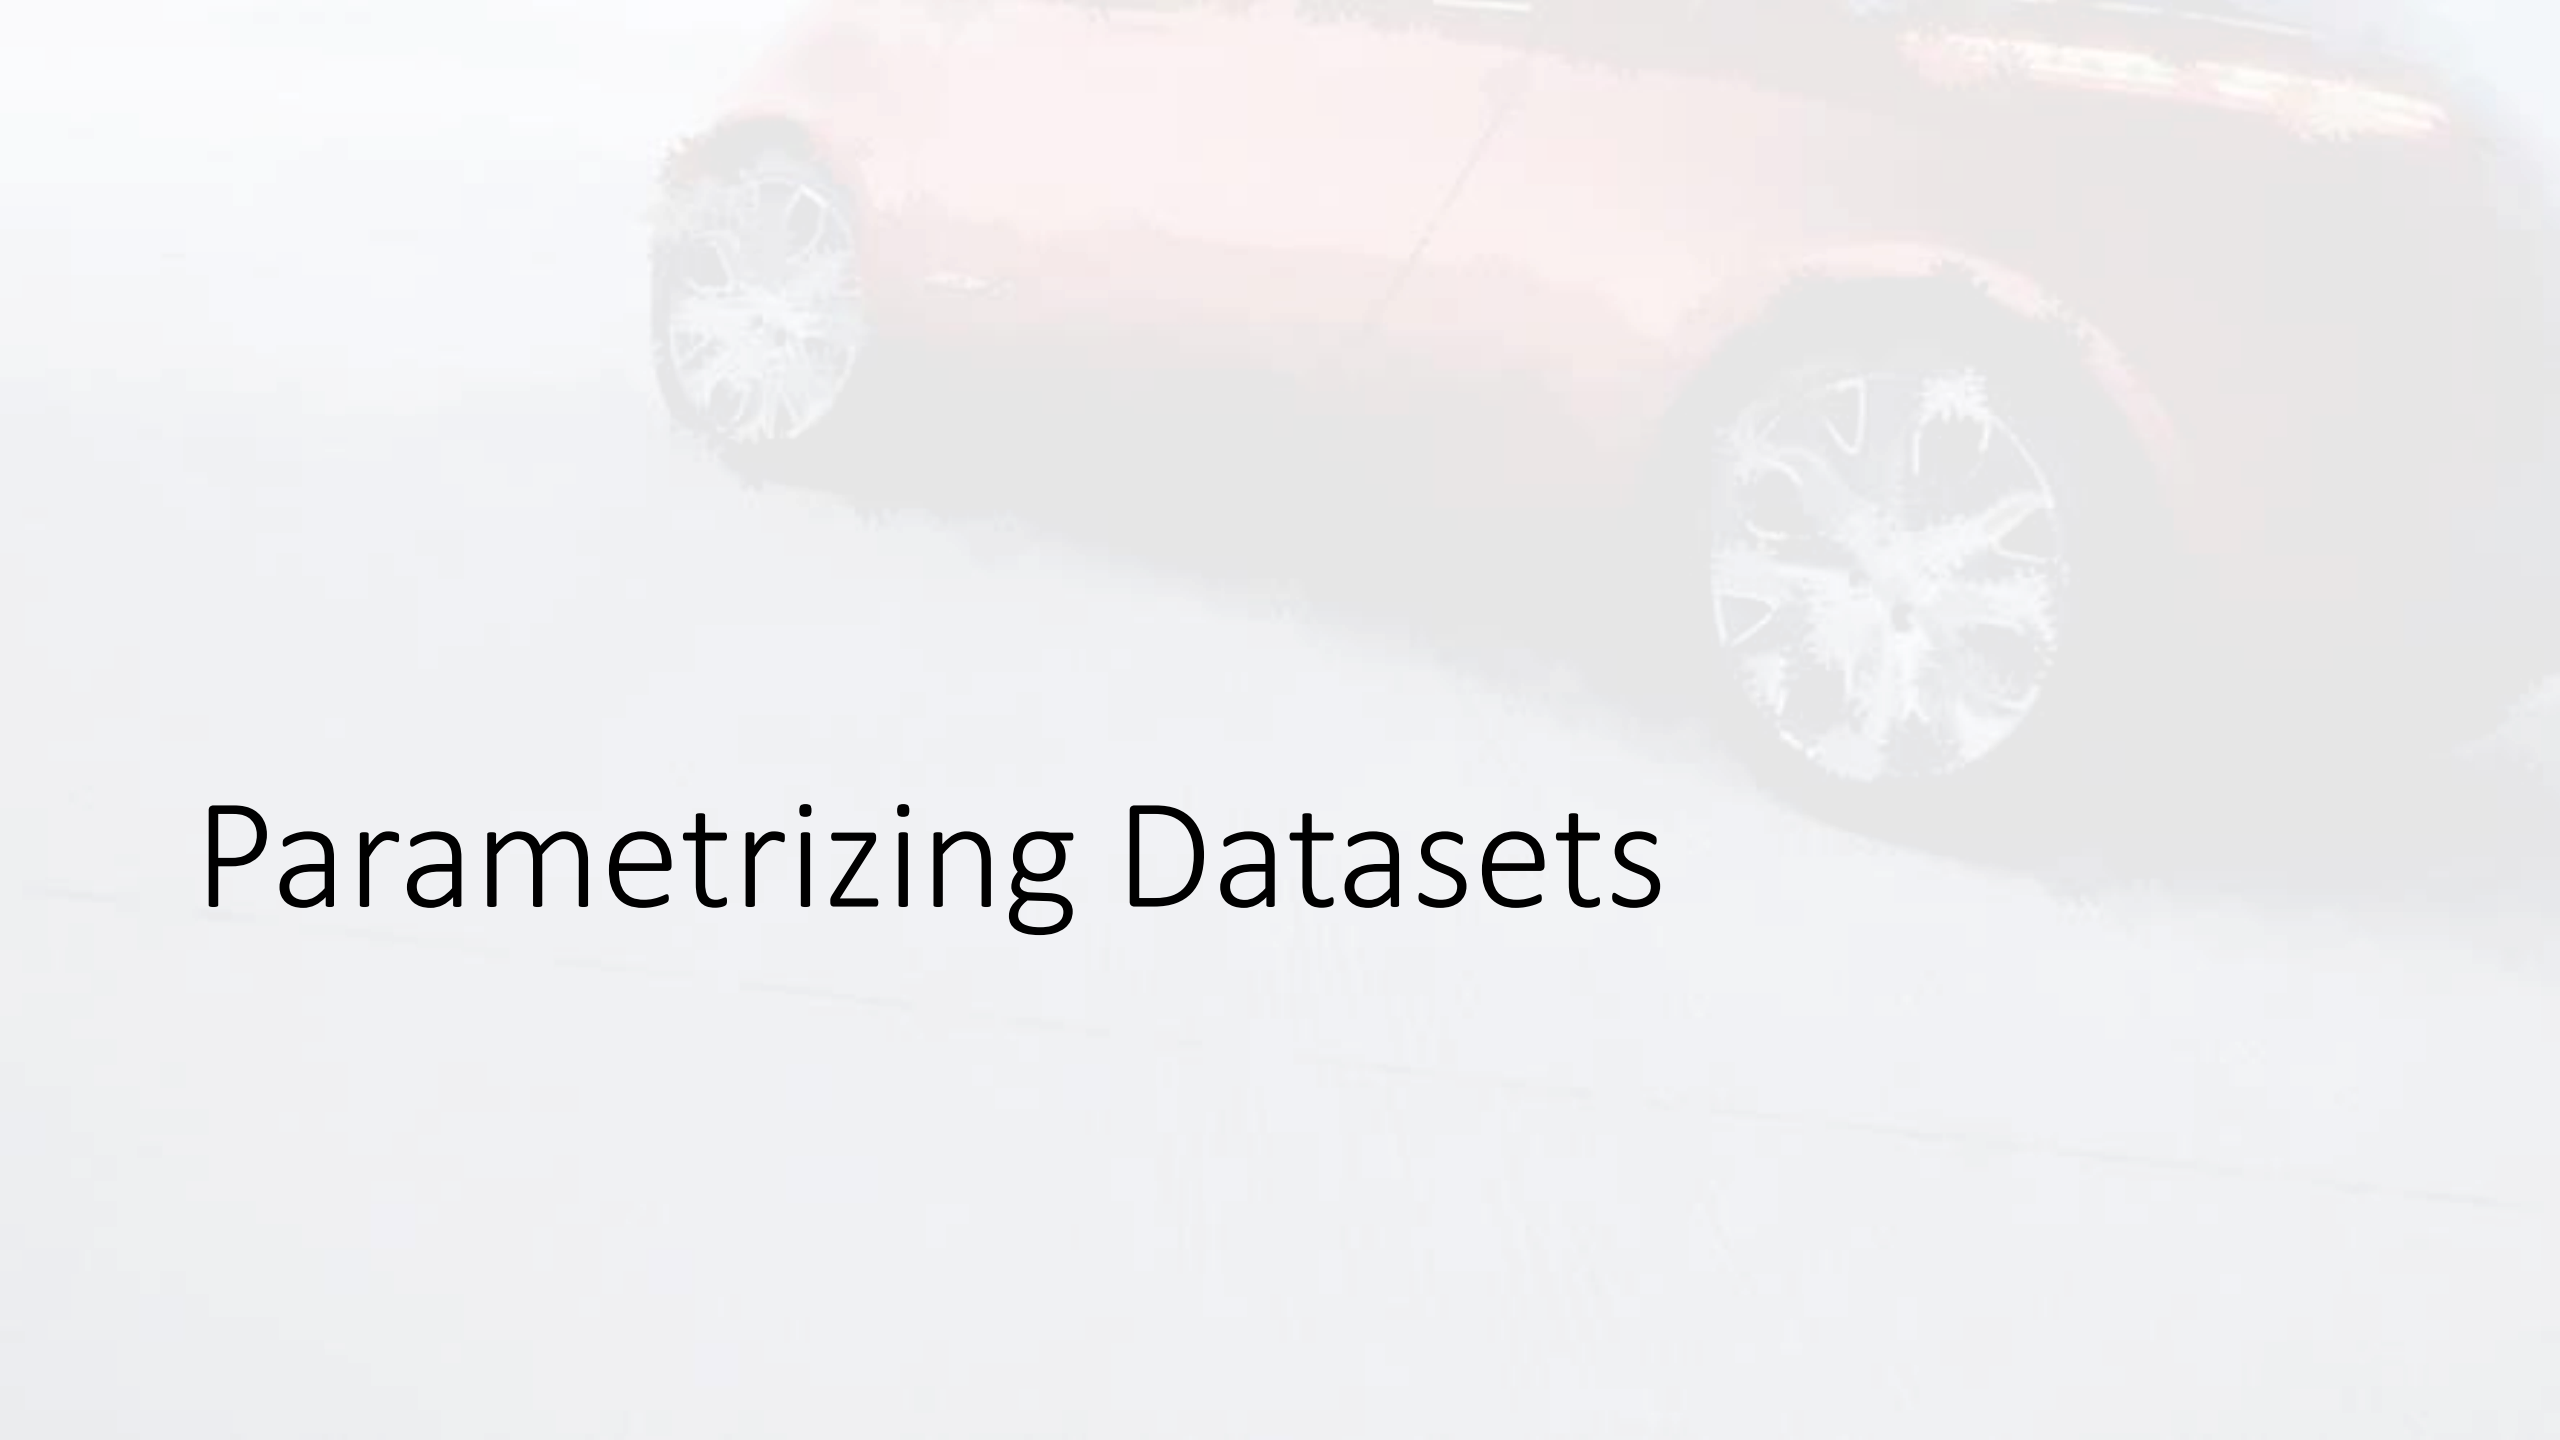

Parametrizing Datasets

Parametrizing Datasets

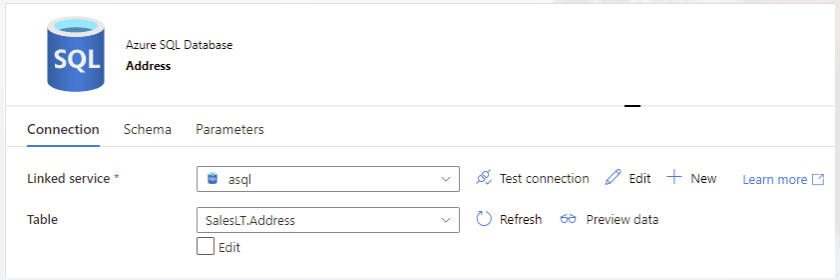

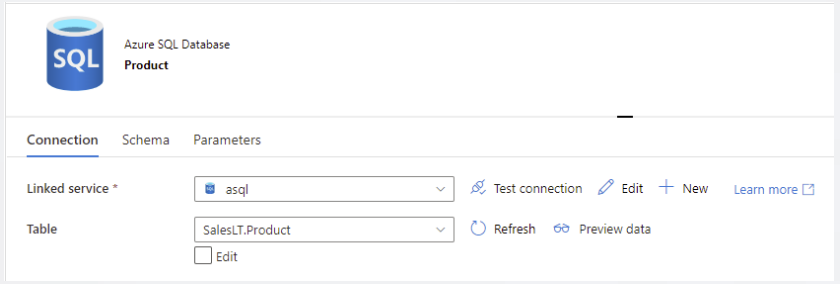

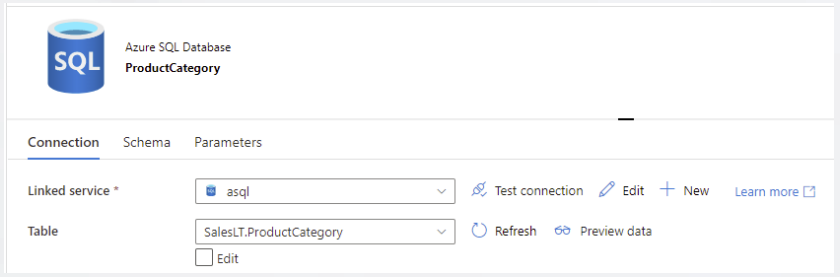

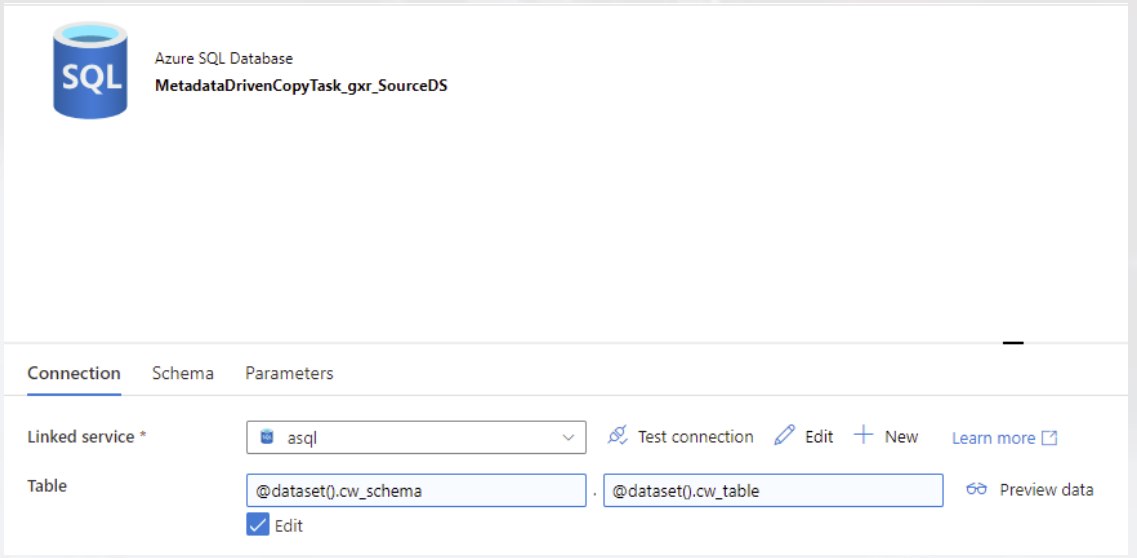

Moreover…


Trust your cloud.
Datasoluzioni
Mattia Cavallotti · 1°
Azure Solutions Architect | MCSE | Microsoft Partner
Datasoluzioni

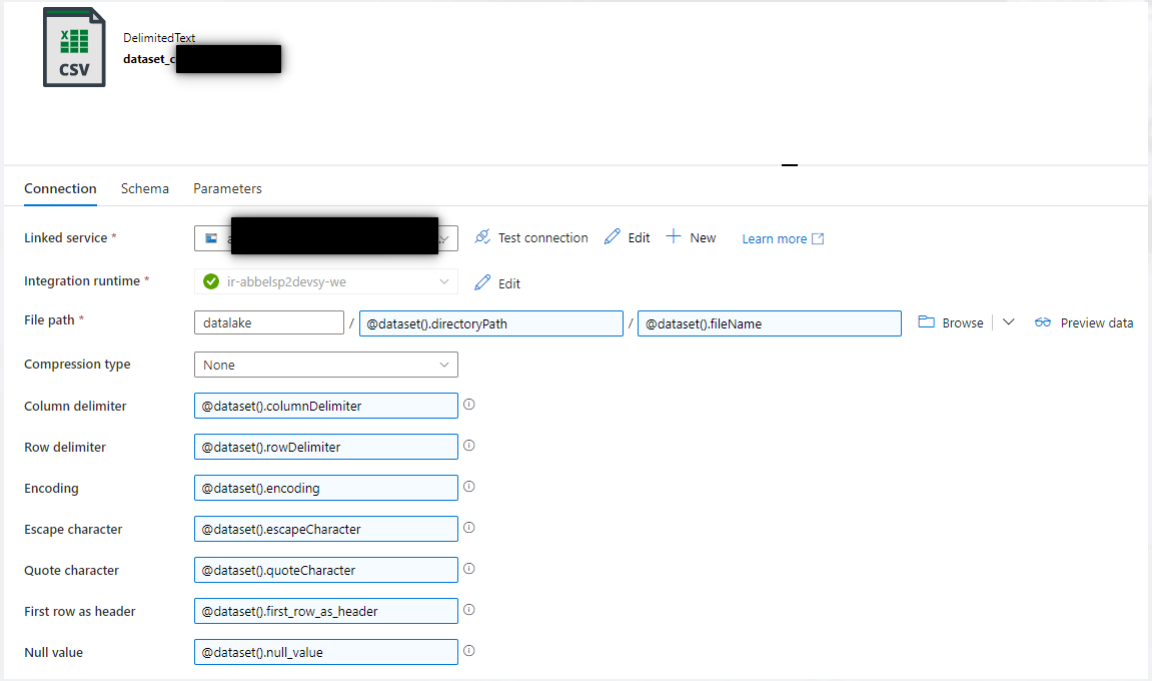
DelimitedText
Connection
Schema
Parameters
Linked service *
Test connection
Edit
New
Learn more
Integration runtime *
ir-abbelsp2devsy-we
Edit
File path *
datalake
@dataset().directoryPath
@dataset().fileName
Browse
Preview data
Compression type
None
Column delimiter
@dataset().columnDelimiter
Row delimiter
@dataset().rowDelimiter
Encoding
@dataset().encoding
Escape character
@dataset().escapeCharacter
Quote character
@dataset().quoteCharacter
First row as header
@dataset().first_row_as_header
Null value
@dataset().null_value

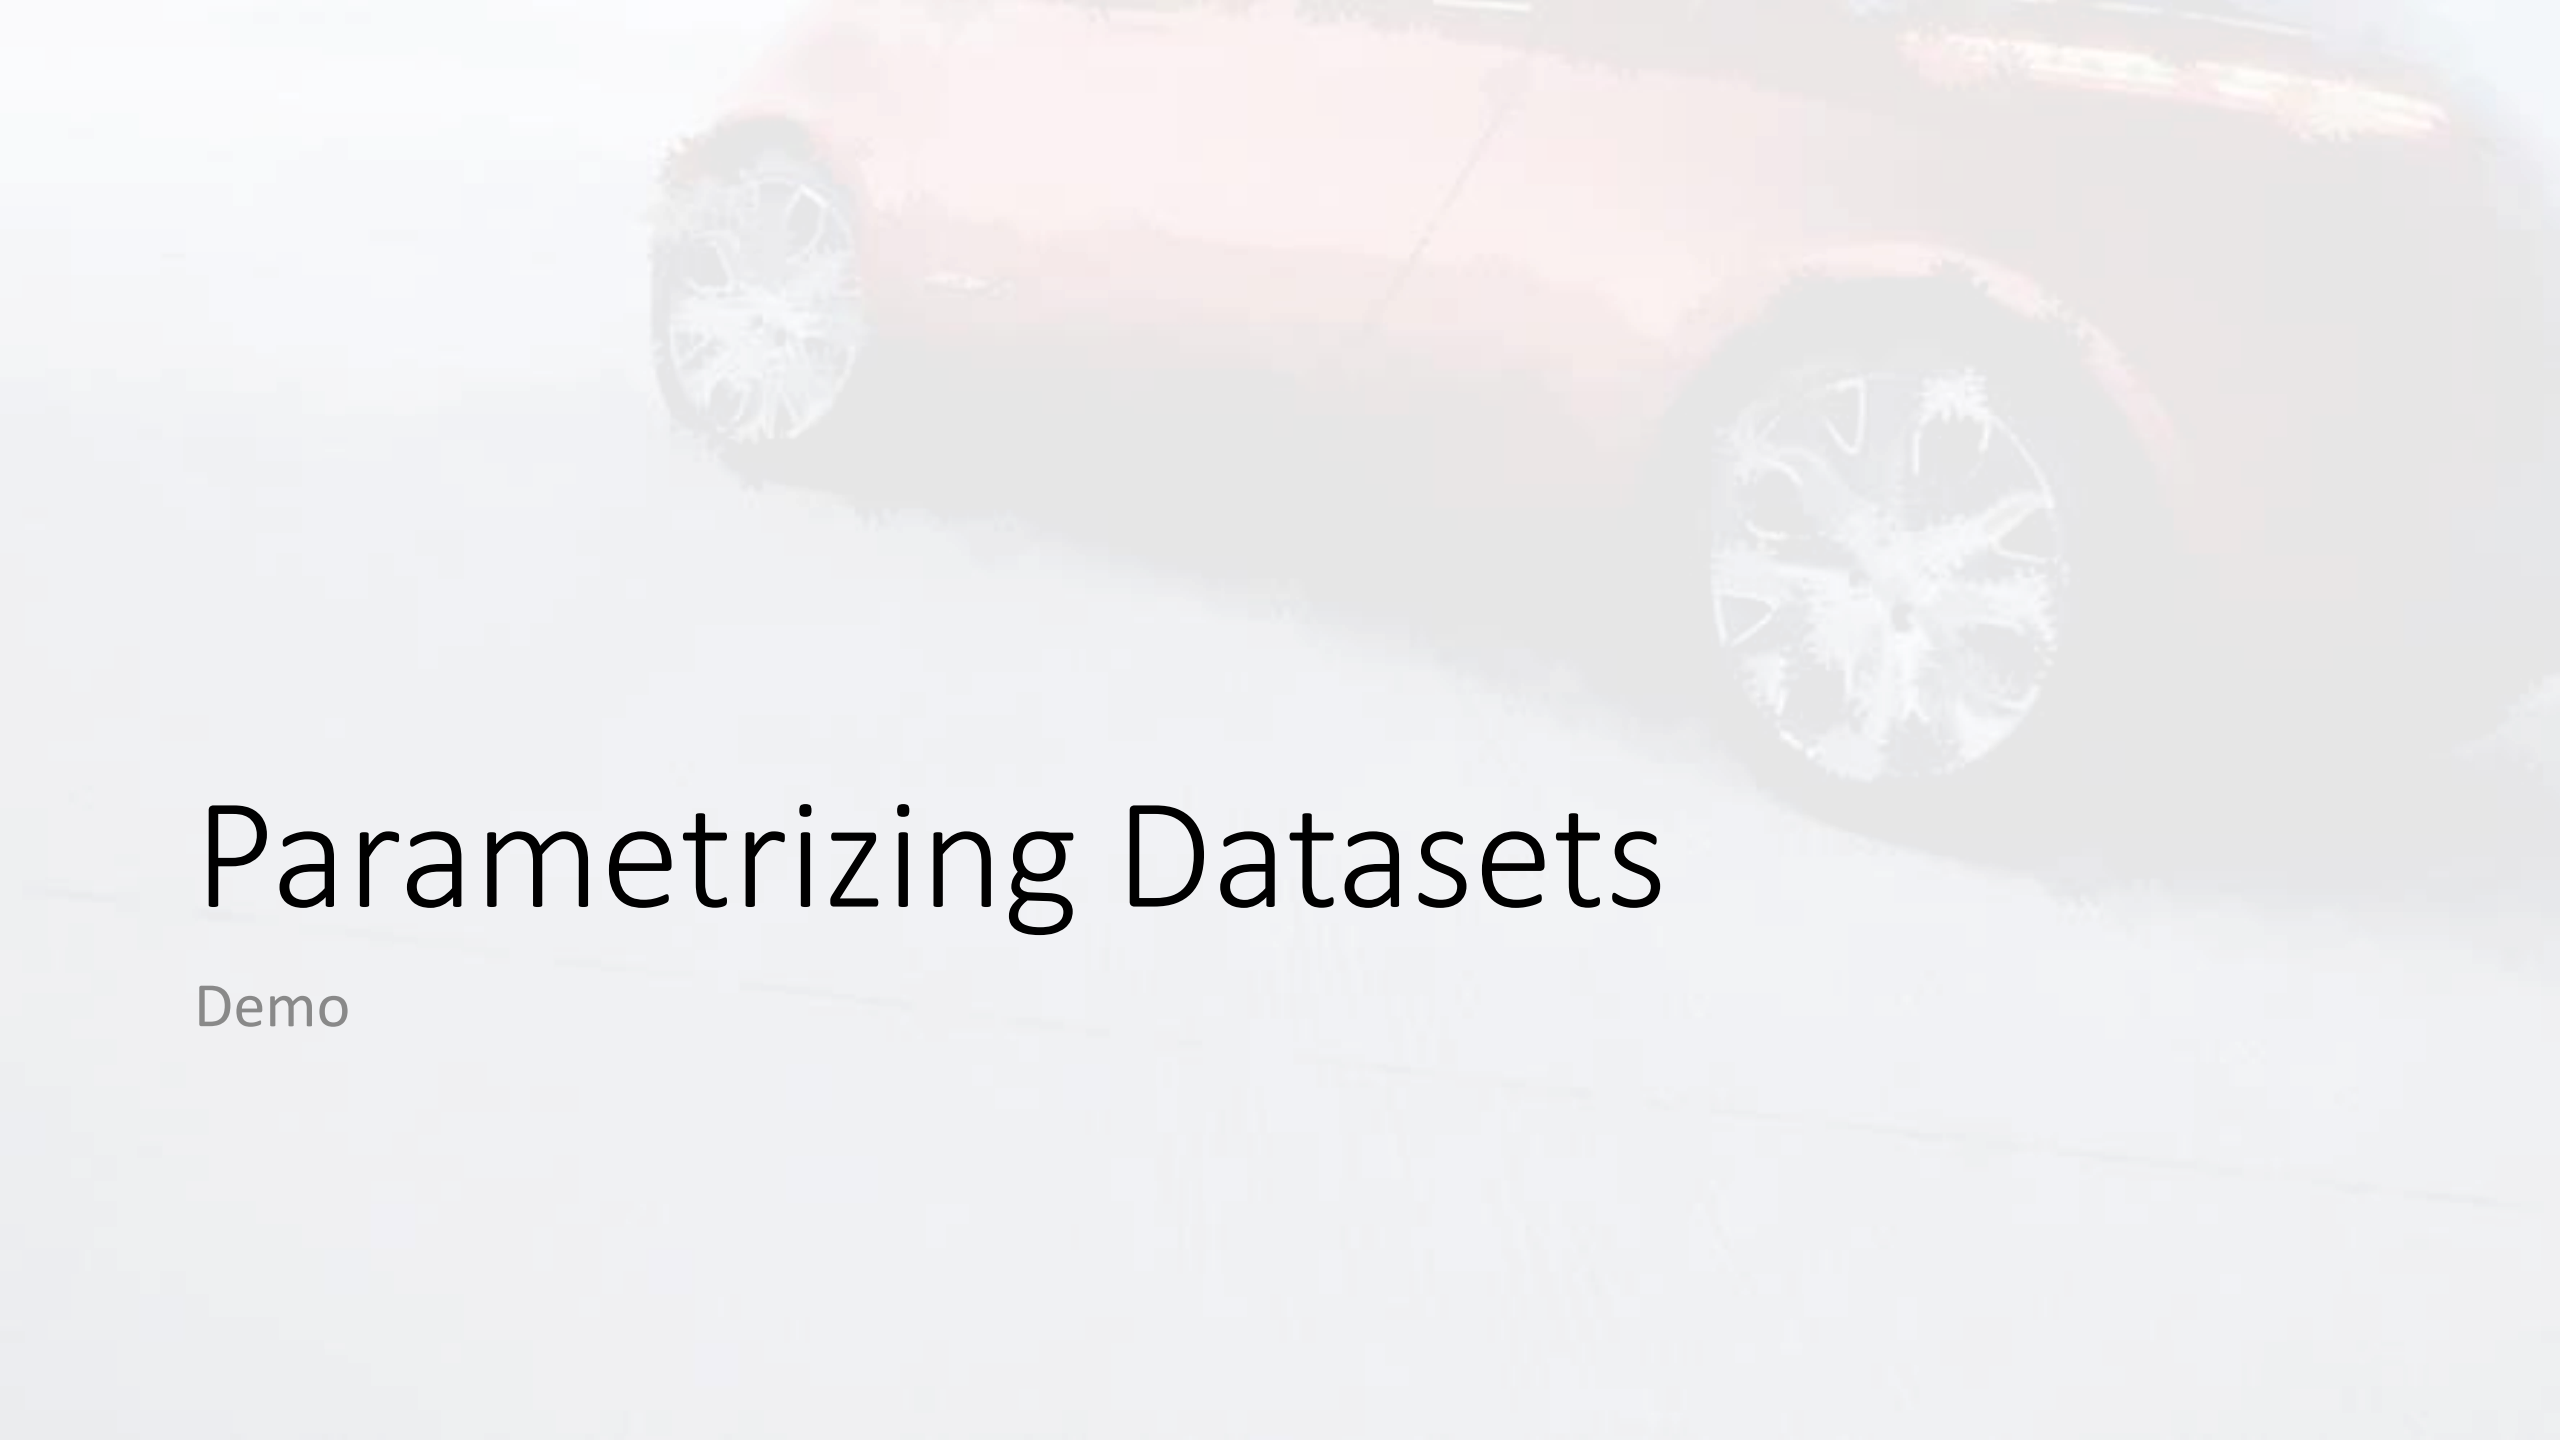

Parametrizing Datasets

Demo

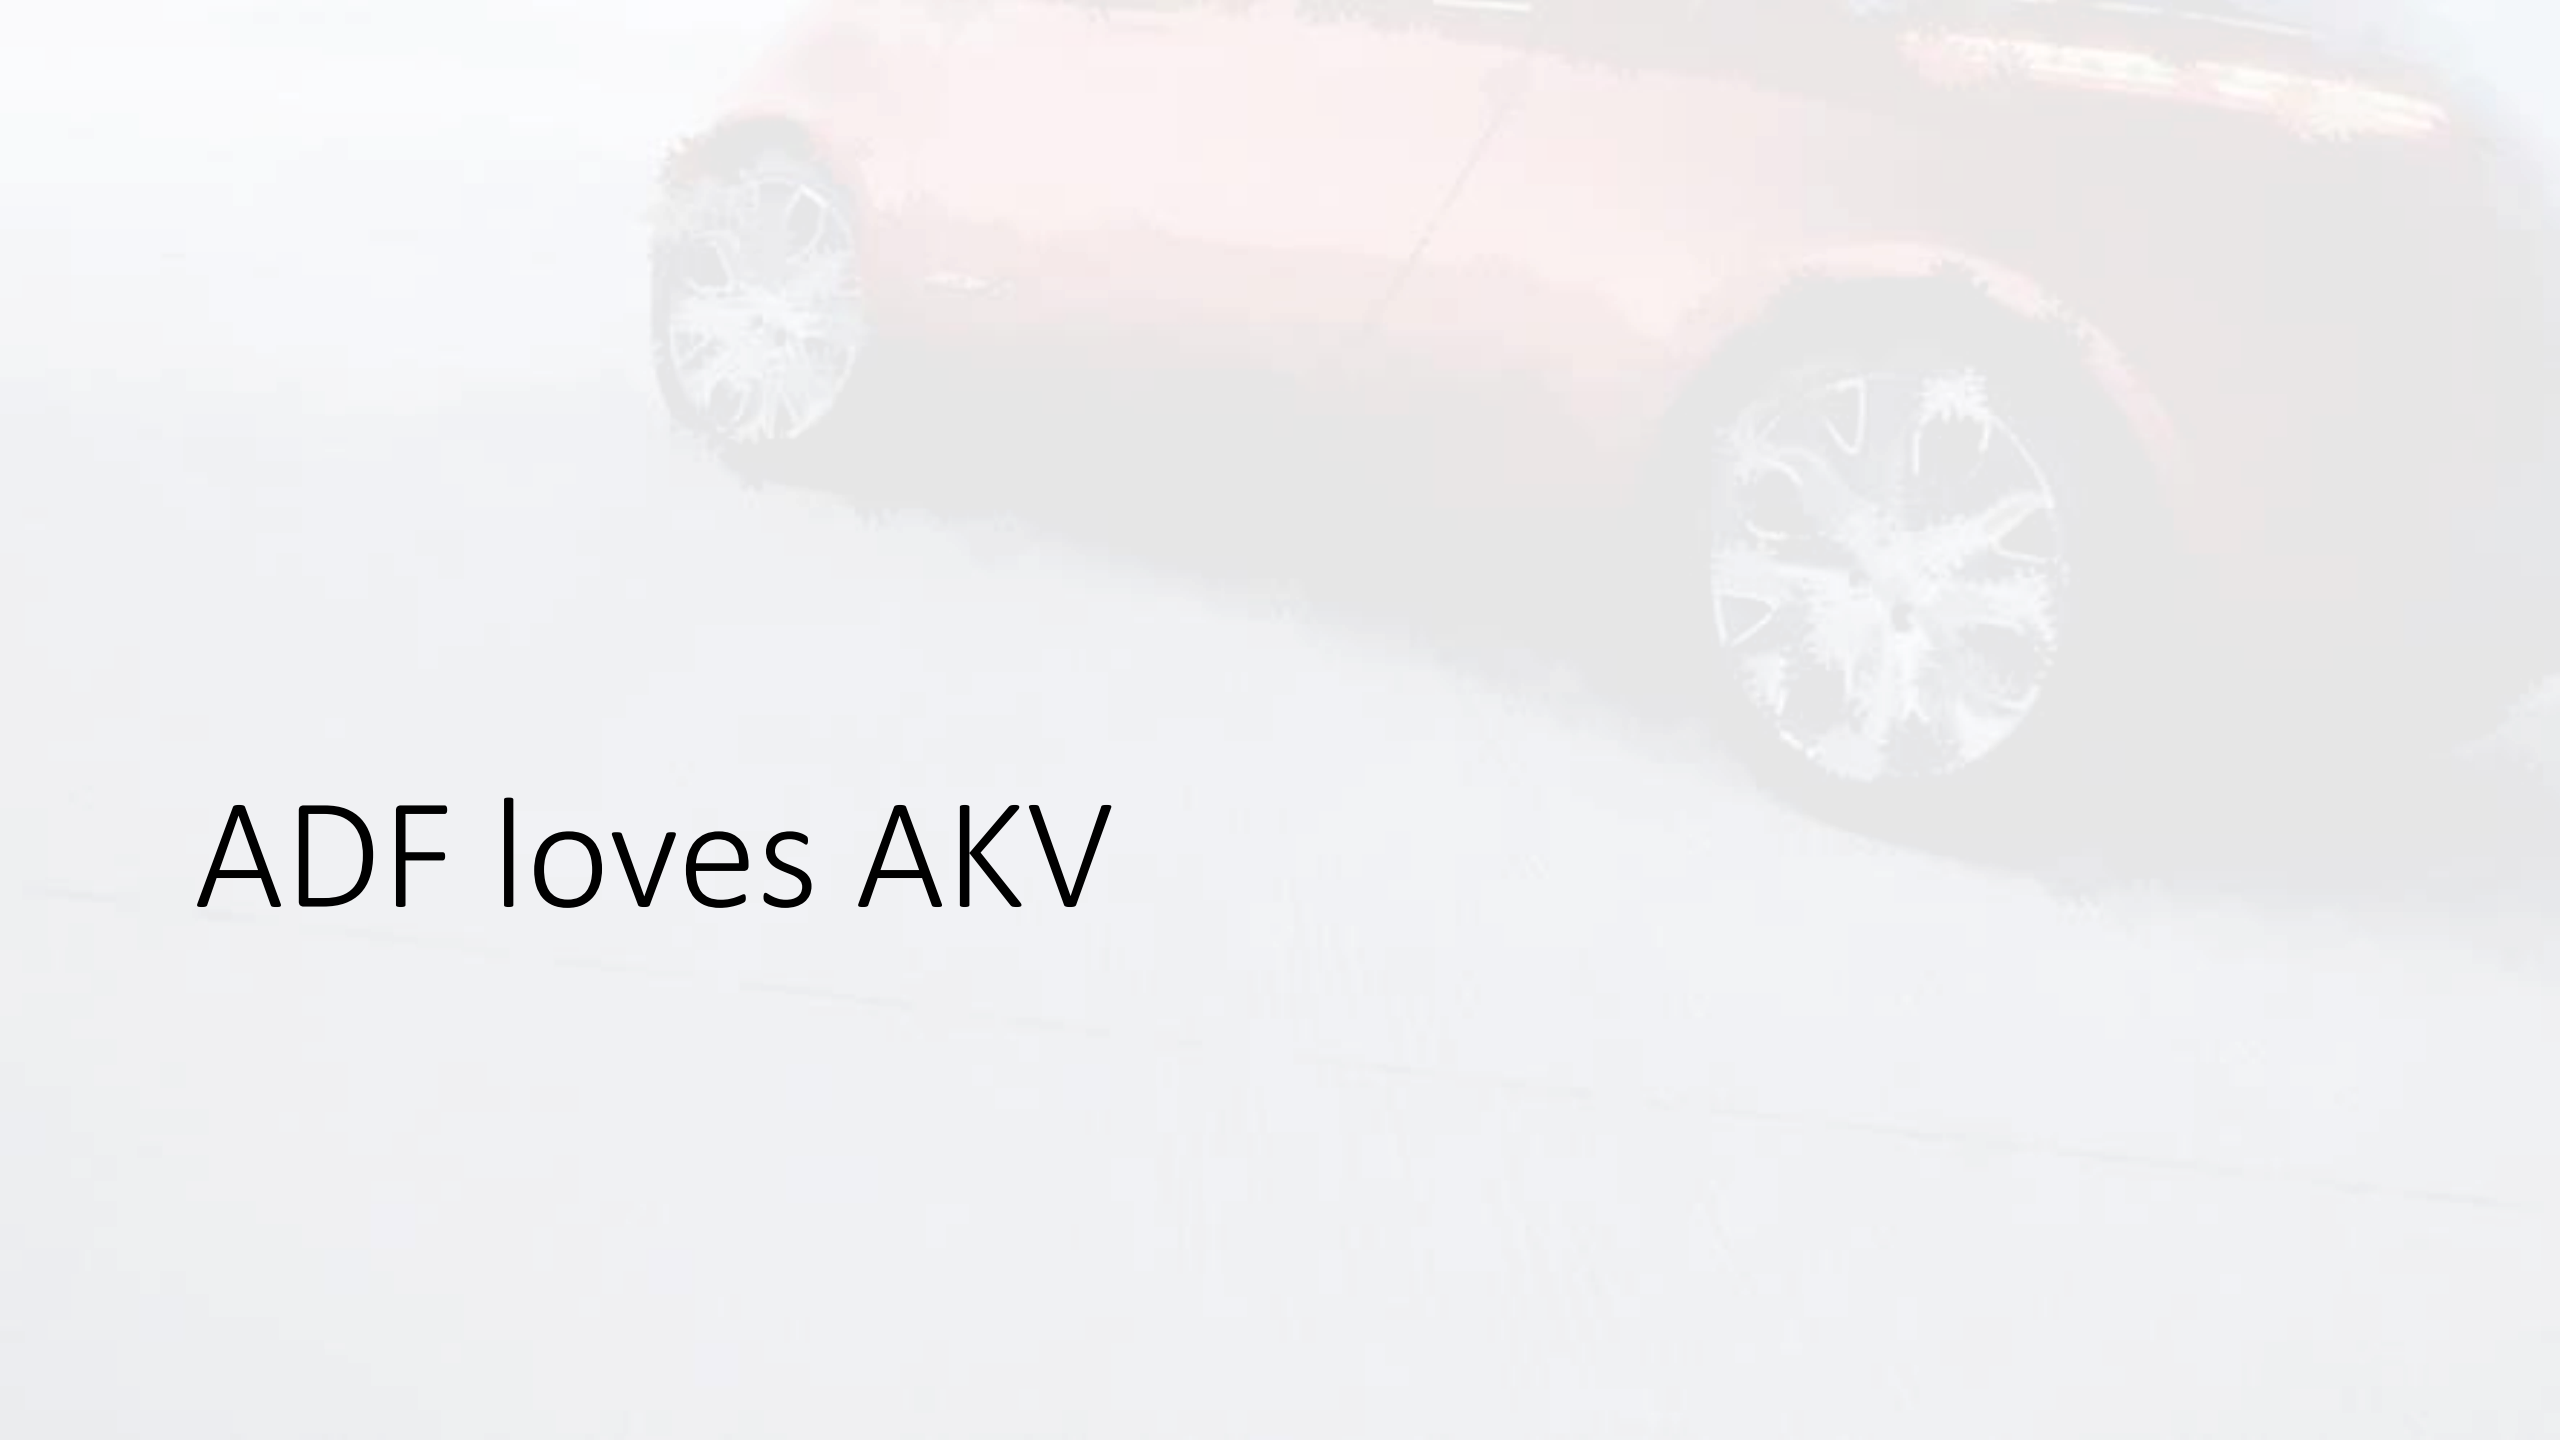

ADF loves AKV

ADF loves AKV

- Best practice: always store **secrets outside ADF**
- Azure Key Vault securely stores and gives access to secrets
- PROS:
  - Developers can work **without knowing secrets**
  - Administrators can **setup and rotate secrets** without accessing ADF
- What should we store?
  - Passwords
  - Users Id / Identities
  - Servers' names / Services' names
  - …



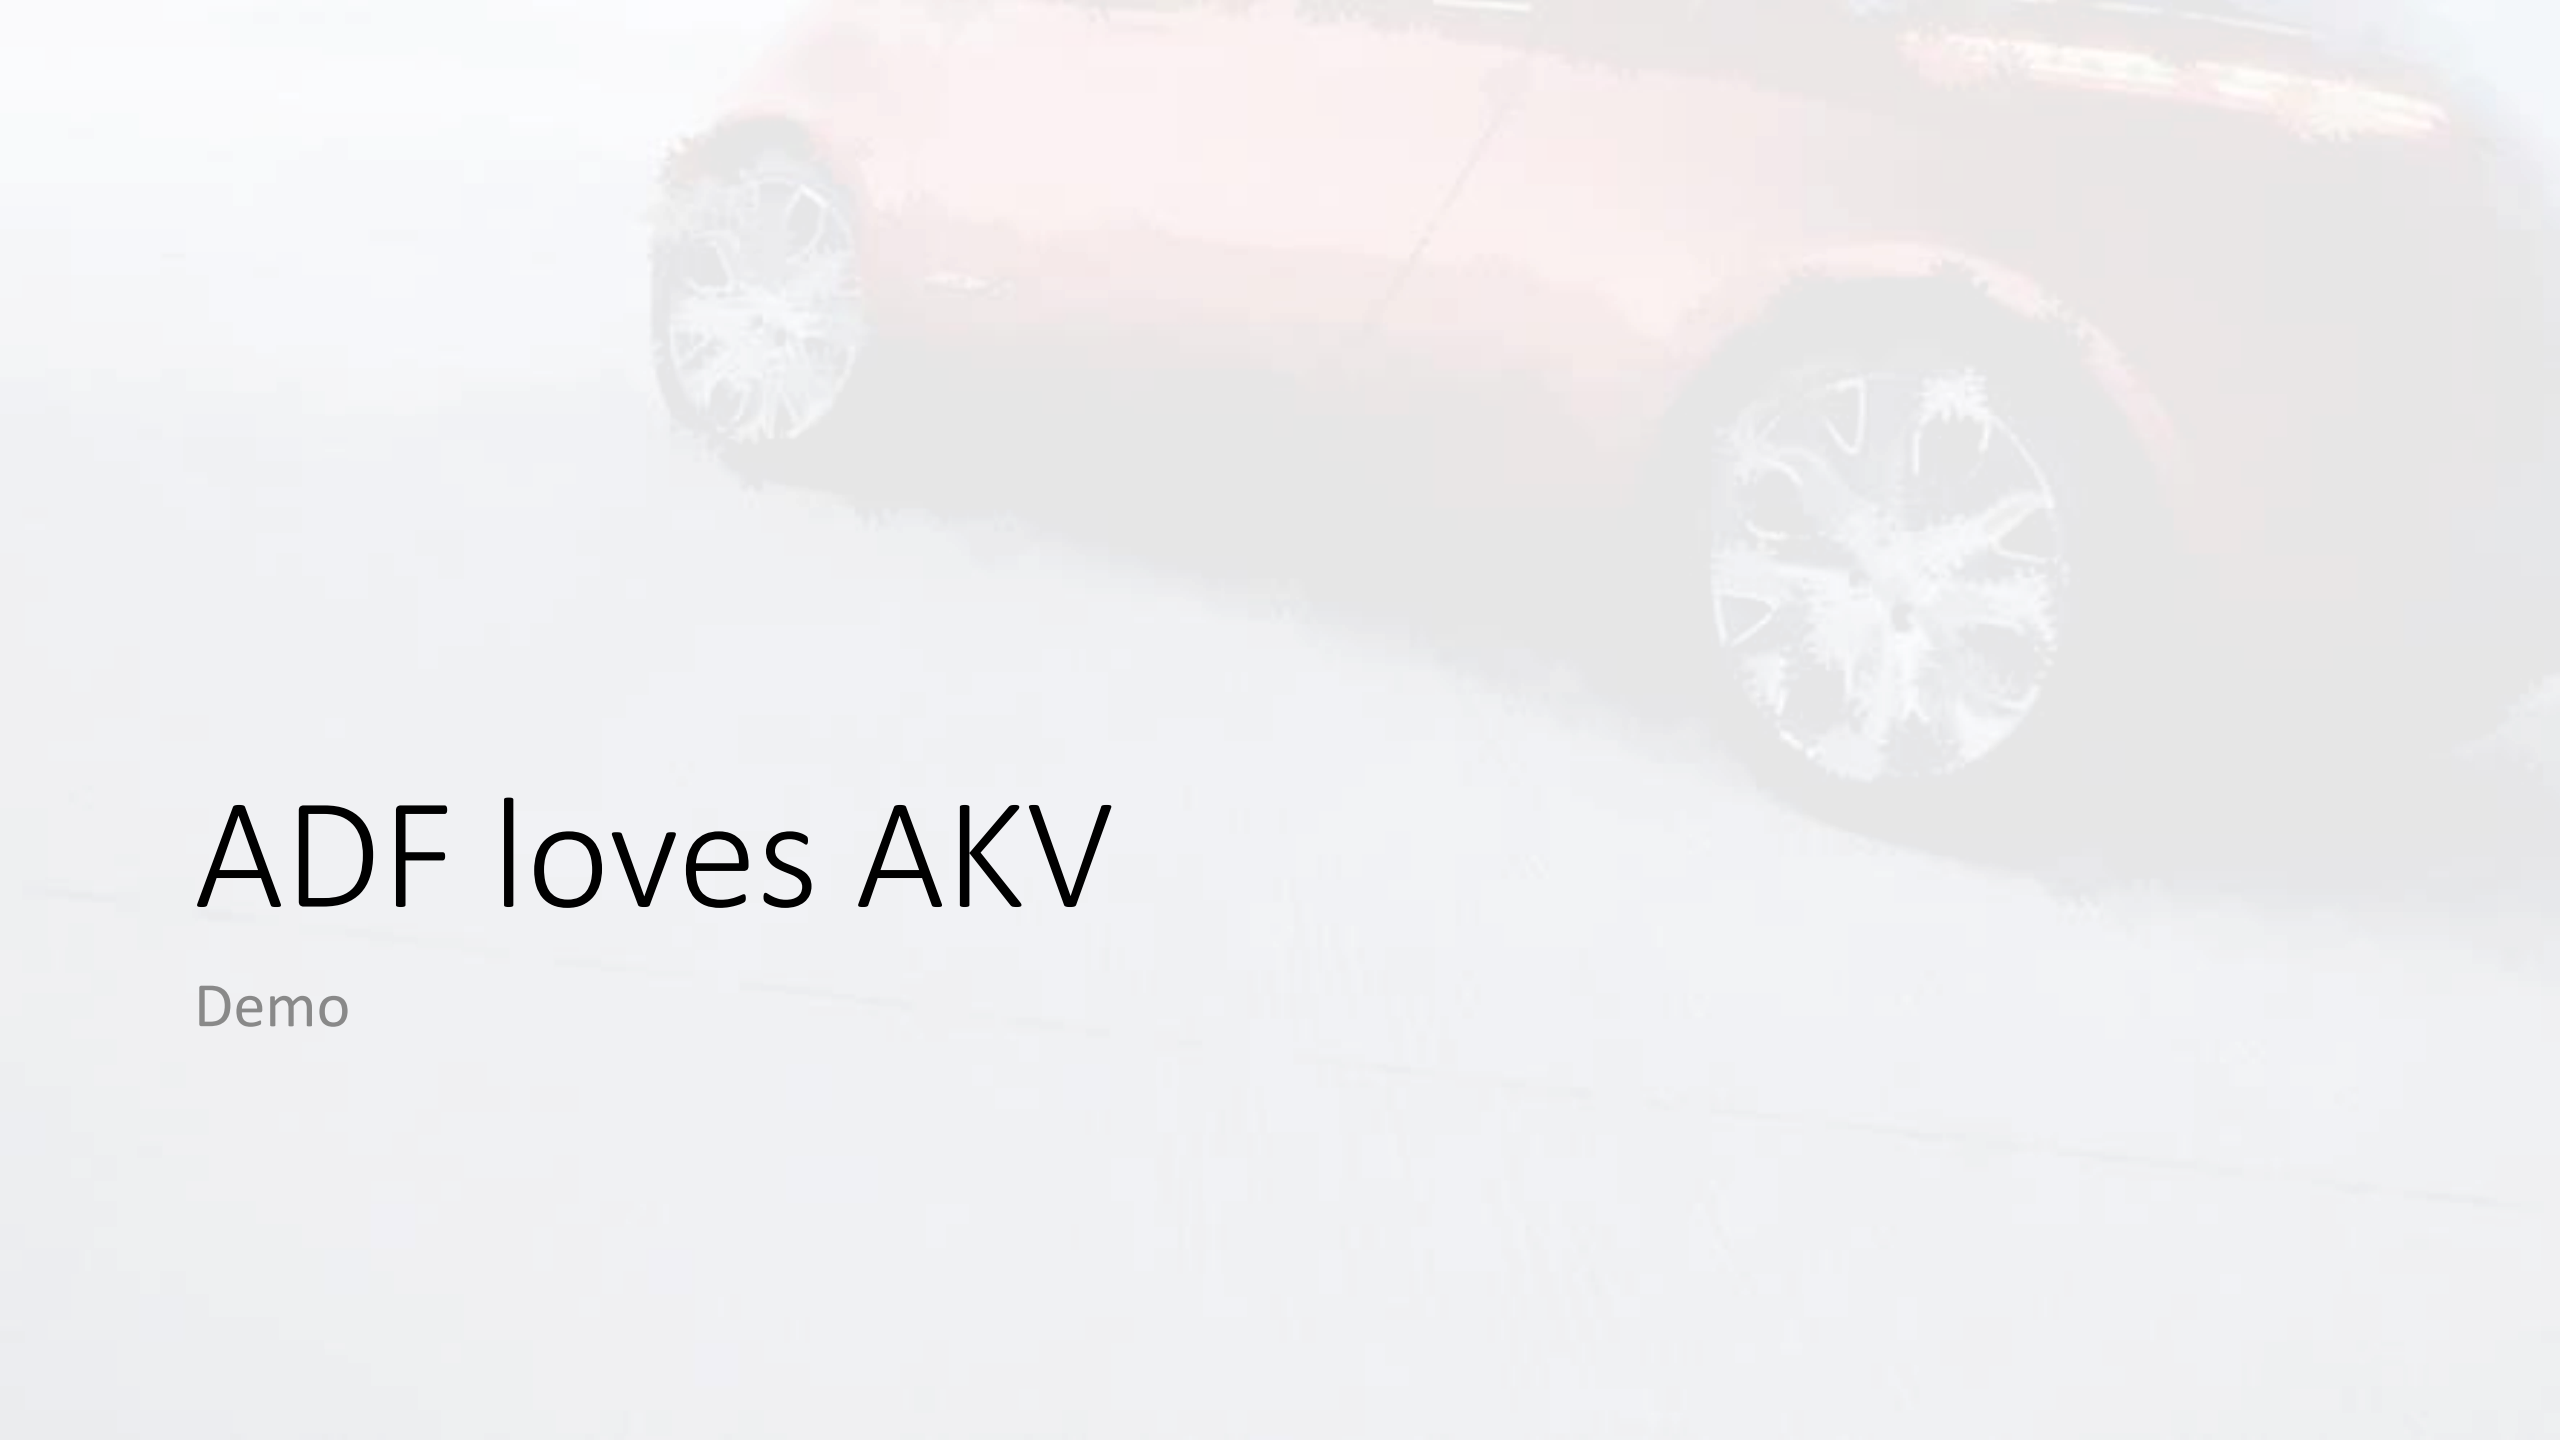

ADF loves AKV

Demo

Managed Identities to get access

What is this?

*"**Managed identities** provide an **identity for applications** to use when **connecting to resources** that support Azure Active Directory (Azure AD) authentication"*

2 types and both supported by ADF

Property	System-assigned managed identity	User-assigned managed identity
Creation	Created as part of an Azure resource (for example, Azure Virtual Machines or Azure App Service).	Created as a stand-alone Azure resource.
Life cycle	Shared life cycle with the Azure resource that the managed identity is created with. When the parent resource is deleted, the managed identity is deleted as well.	Independent life cycle. Must be explicitly deleted.
Sharing across Azure resources	Can't be shared. It can only be associated with a single Azure resource.	Can be shared. The same user-assigned managed identity can be associated with more than one Azure resource.
Common use cases	Workloads that are contained within a single Azure resource. Workloads for which you need independent identities. For example, an application that runs on a single virtual machine.	Workloads that run on multiple resources and can share a single identity. Workloads that need pre-authorization to a secure resource, as part of a provisioning flow. Workloads where resources are recycled frequently, but permissions should stay consistent. For example, a workload where multiple virtual machines need to access the same resource.

Connect to ASQL via Managed Identity

```
CREATE USER [my-adf-name] FOR EXTERNAL PROVIDER;
GO

GRANT CONNECT TO [my-adf-name];
GO

ALTER ROLE [db_owner] ADD MEMBER [my-adf-name];
GO
```

TTL integration runtime

You definitely should use it

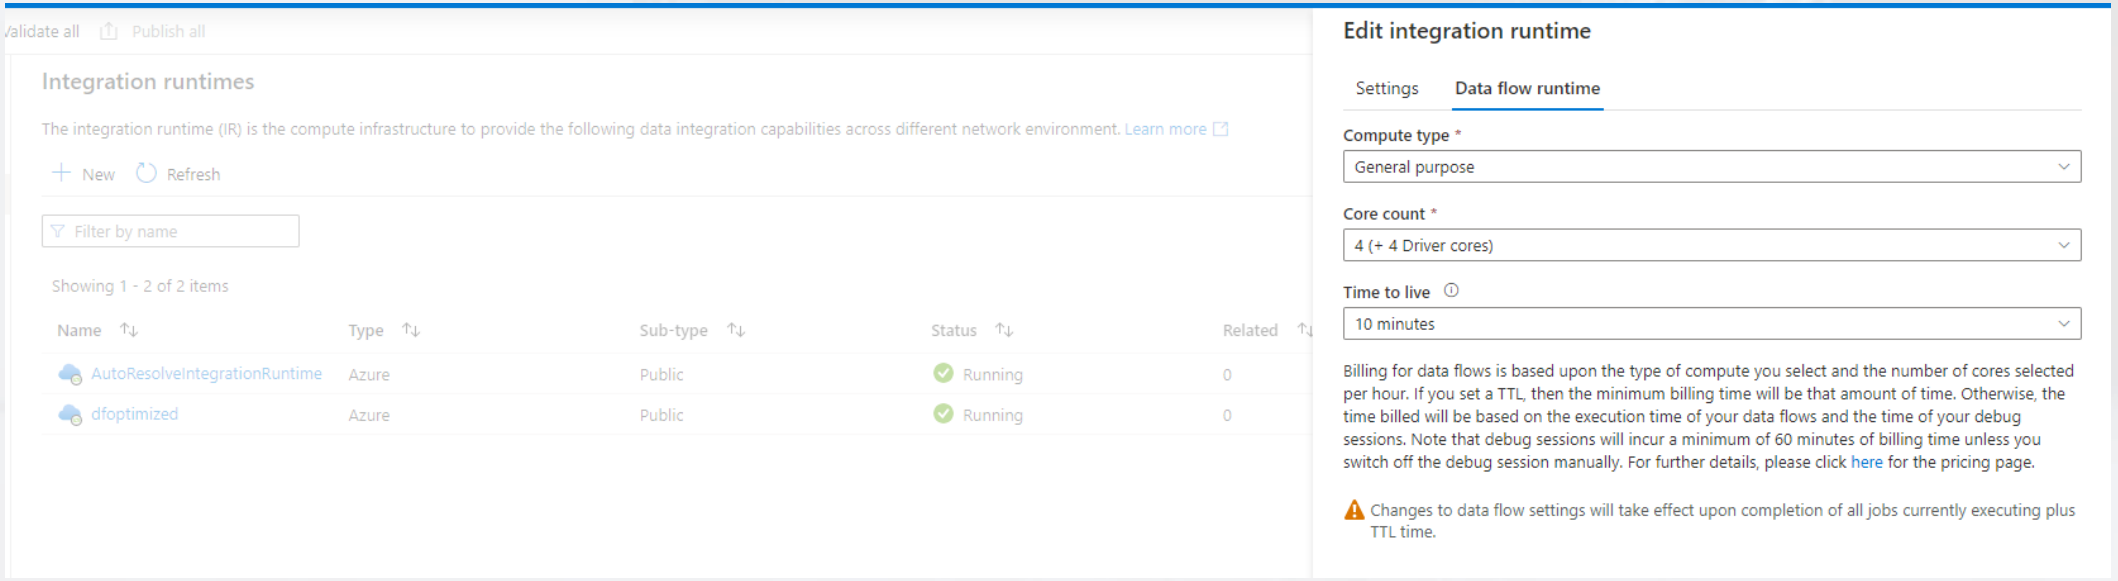

TTL integration runtime

Demo

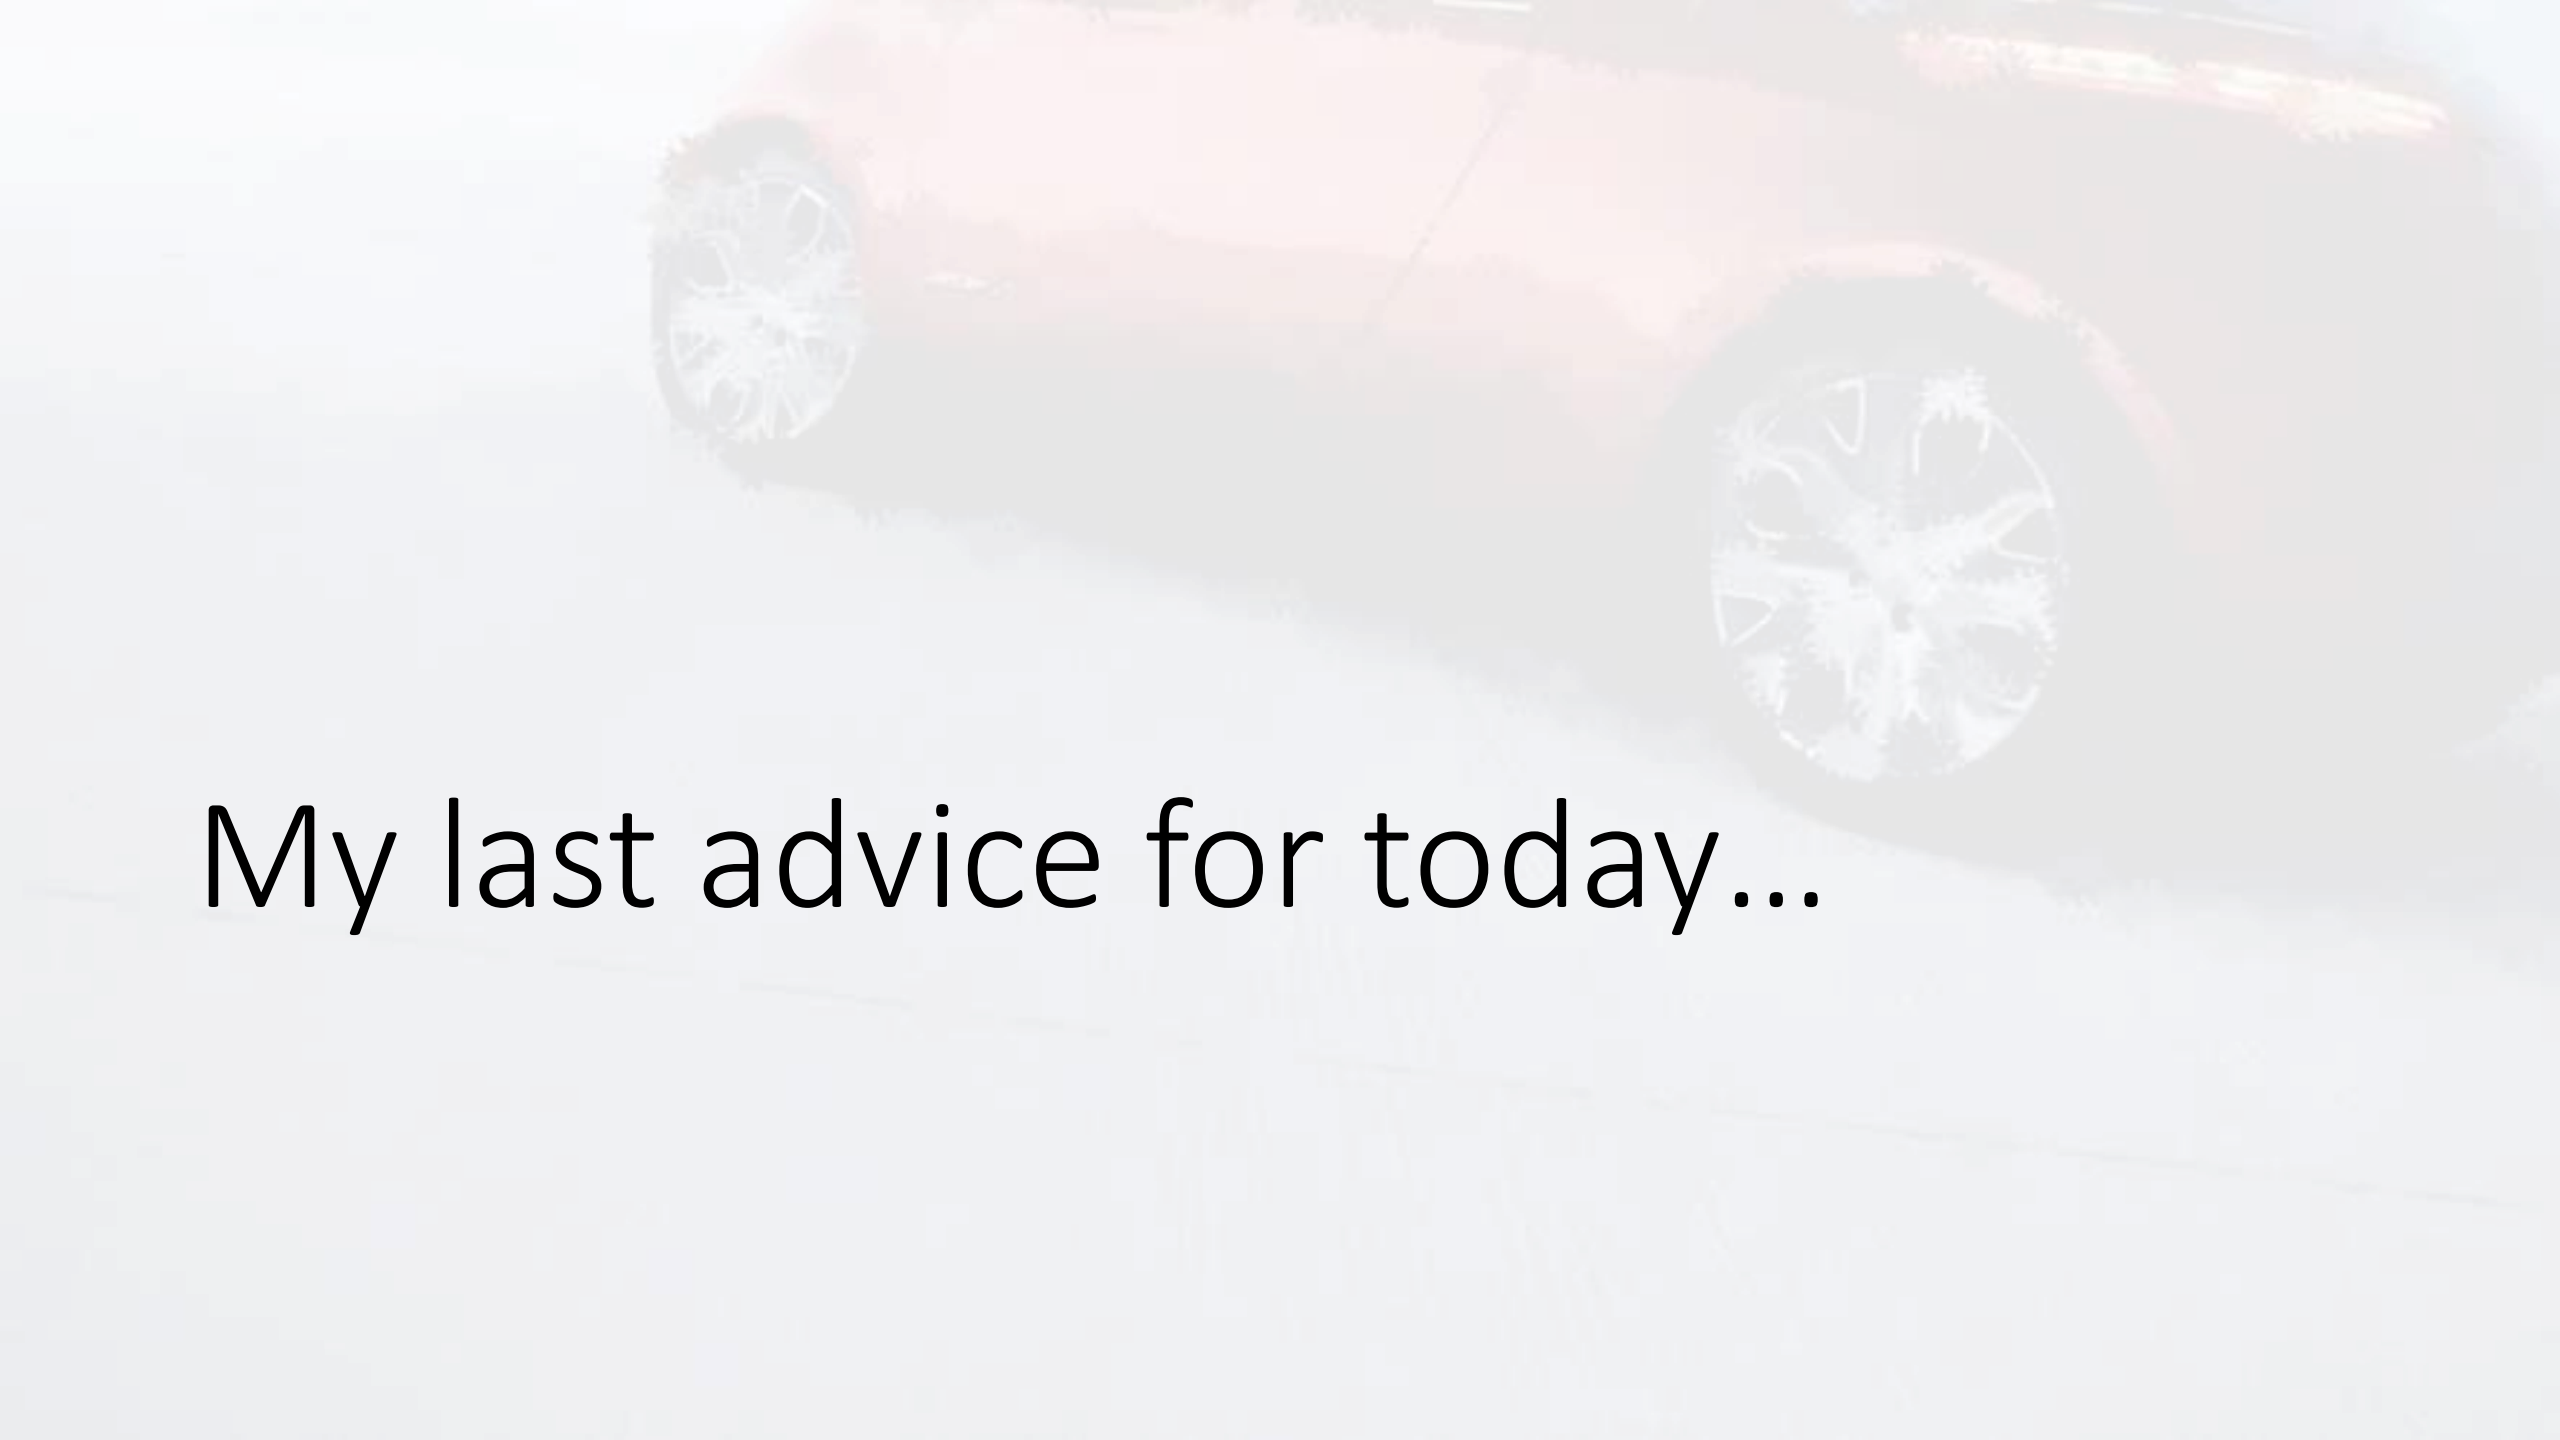

My last advice for today…

No demo sorry…

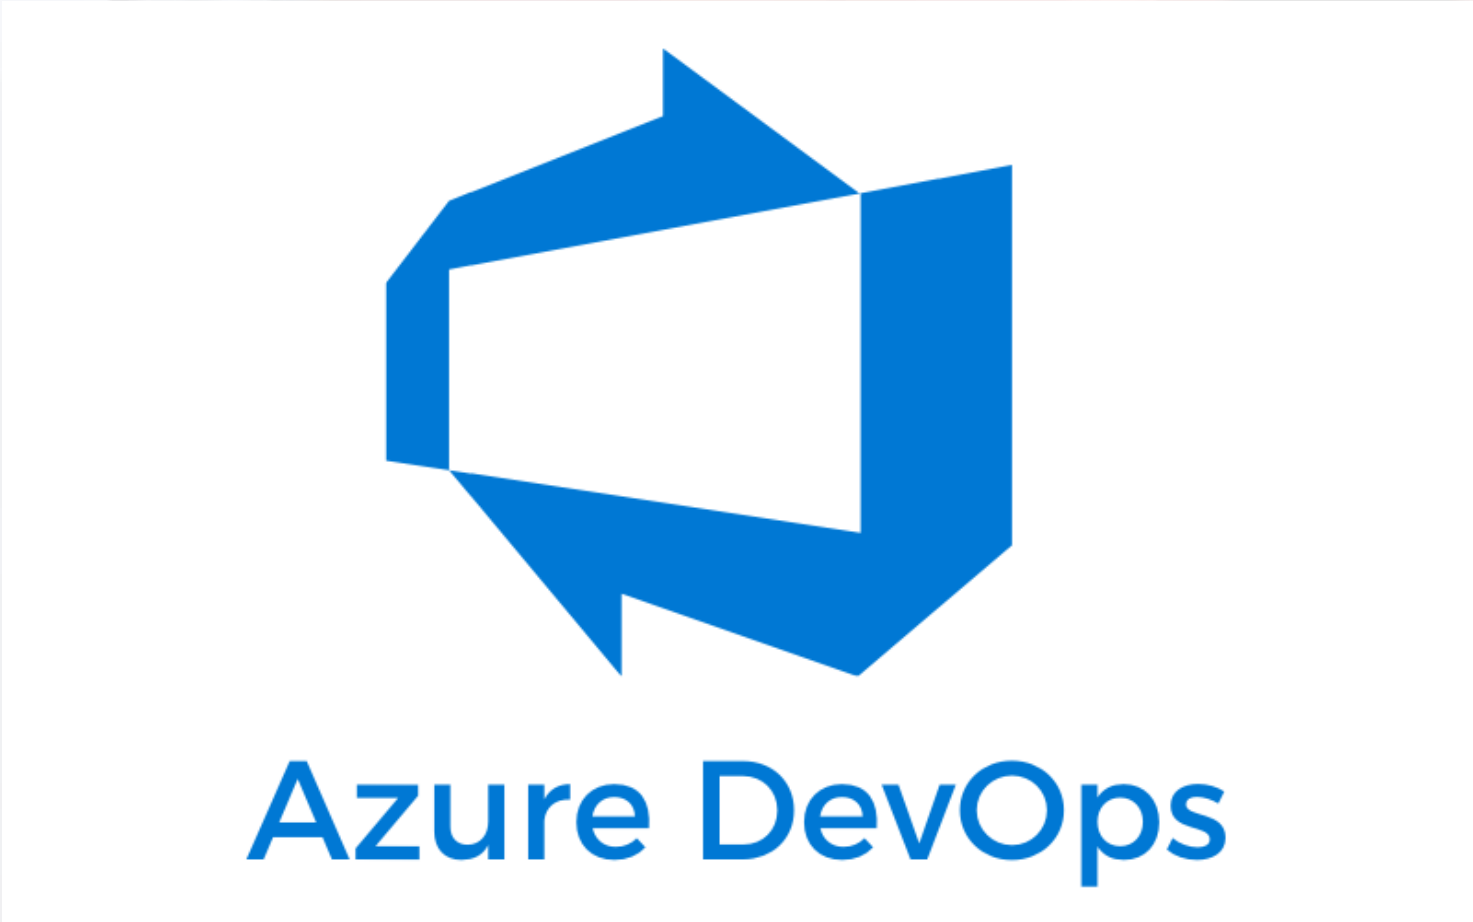

Useful Links

- ADF docs: http://tiny.cc/adfrw1
- Metadata driven out-of-the-box: http://tiny.cc/adfrw2
- Metadata driven by Paul Andrew: https://github.com/mrpaulandrew/procfwk
- Managed identities: http://tiny.cc/adfrw3
- Connect to ASQL using ADF's Managed Identity: http://tiny.cc/adfrw4


SQL START! 2022
Q&A

Thank you!