

Fabric I love you, I love you not

Riccardo Perico
BI & Power BI Engineer @Lucient Italia



Sponsor



Bio



Lucient[®]
ITALIA



Riccardo Perico
BI & Power BI Engineer



[riccardo-perico](#)



[R1k91](#)



[R1k91](#)



[riccardo-perico](#)



[riccardo-perico](#)

In this session I'll try not to “simply explain features”,
main goal is to **give a feedback** on them,
based on my personal experience with Fabric
doing tests and POCs.

DISCLAIMER



What to expect from this session?

- **Discover** things
- **Enforce** knowledge
- If you enter that door with doubts... will you get an **answer**?

Maybe

Maybe not...

- But I hope you'll get the **right doubts**

D E A D

DISCLAIMER

Maybe I'm wrong, or
you've a different point
of view...

Share your experience
while we go,
if you please



DISCLAIMER

I tried to keep my session as aligned as possible
with all the recent announcements.

Please forgive me if I lost something.

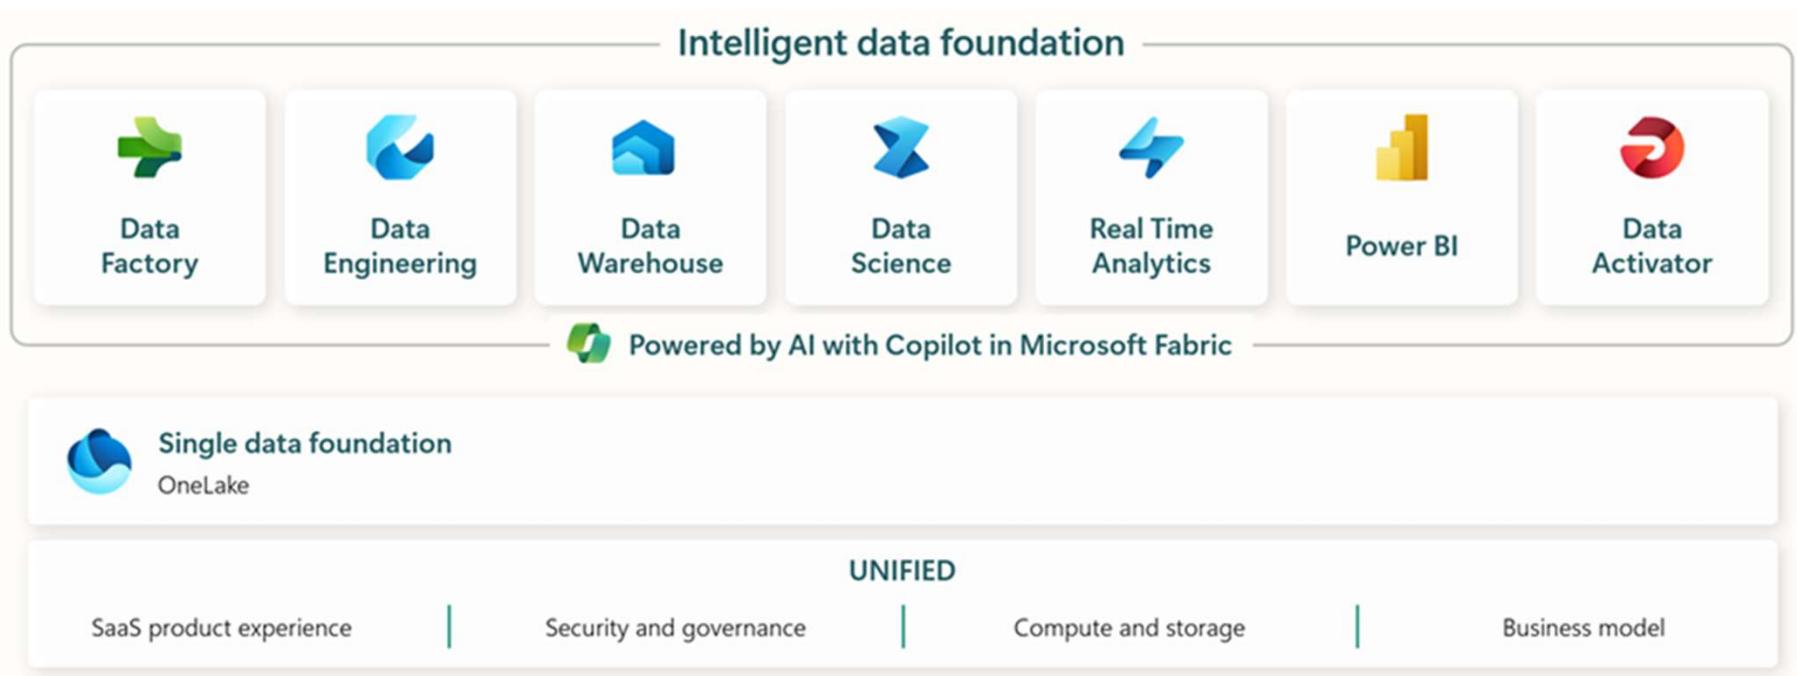
I feel this session very ambitious,
I hope to met your expectations.

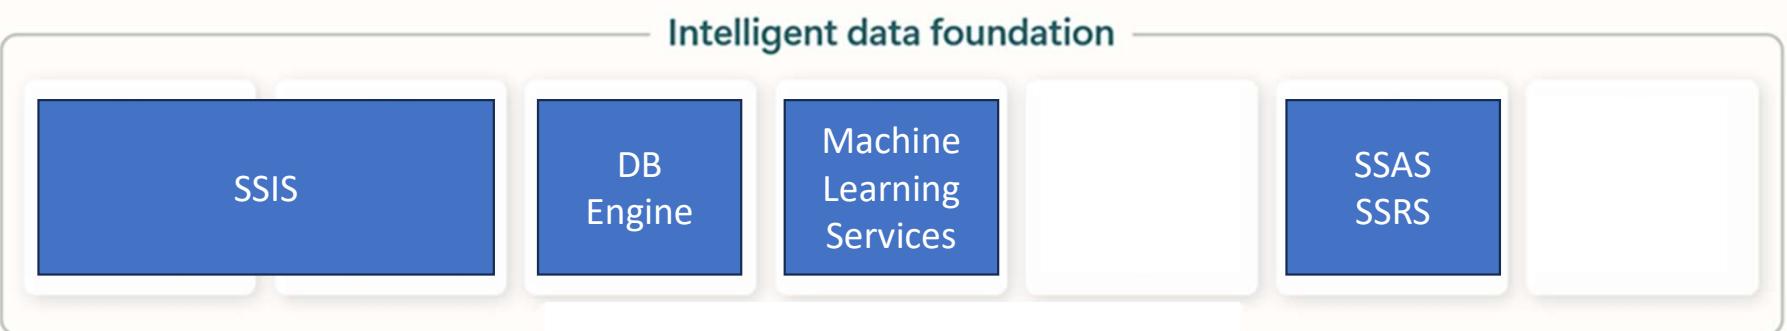
DISCLAIMER



What I like (maybe the most) is the idea

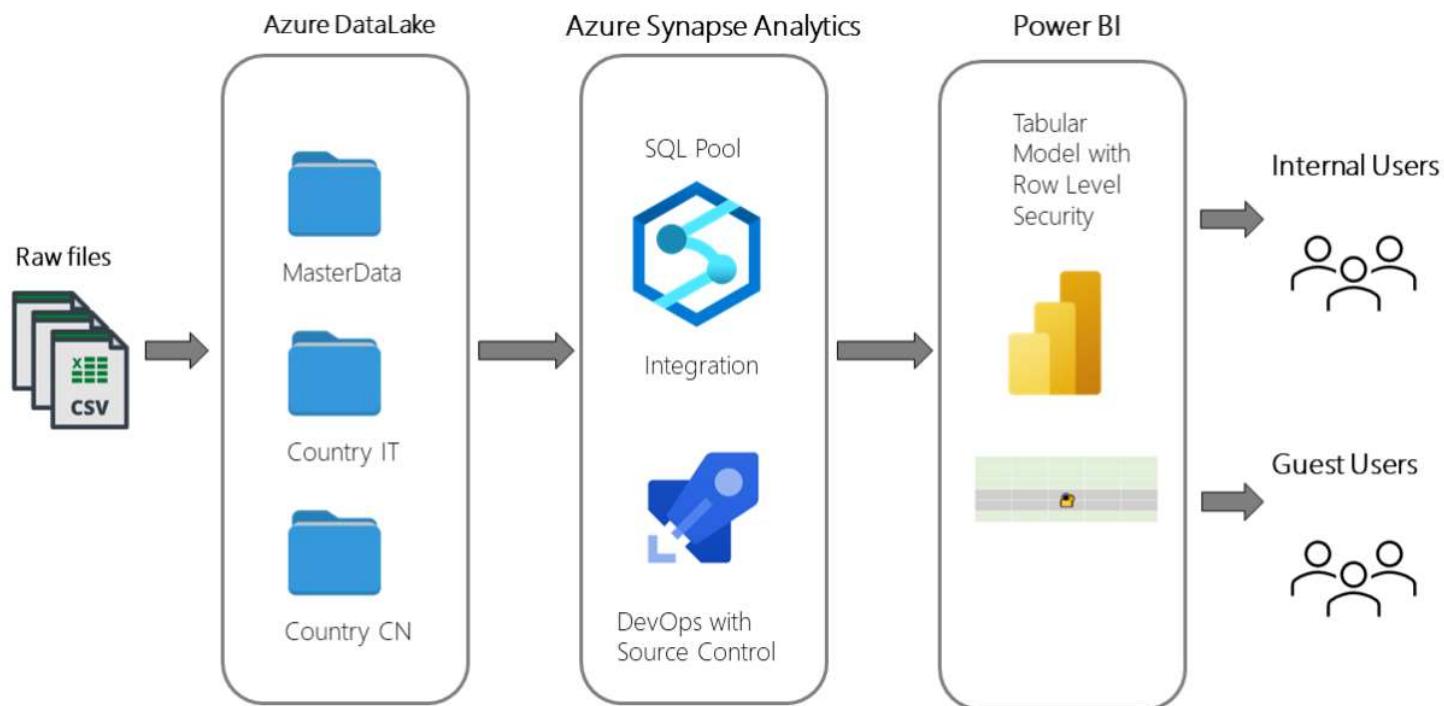
The great unifier





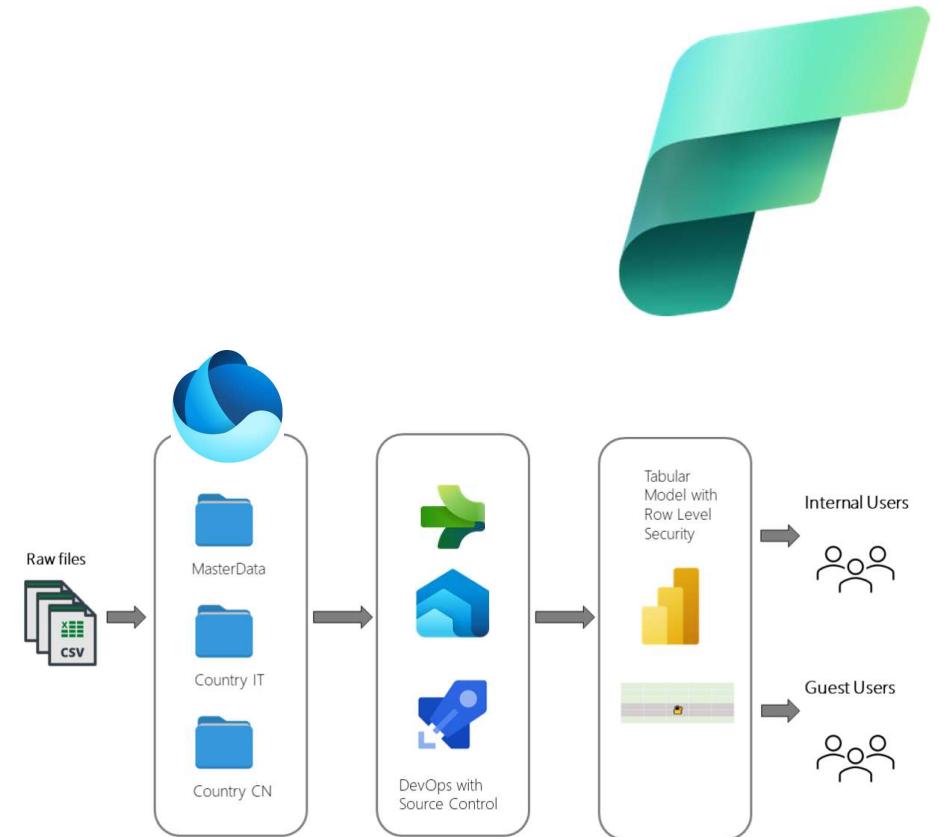
VM Storage

As it is



To be

- From **ADLS to OneLake folders**
- Users can use **Explorer Add-in**
- From **Synapse Pipelines to Data Pipeline**
- From **Dedicated SQL Pool to Warehouse**
- From **Power BI to Power BI**
- Therefore... all-in one... Workspace



OneLake idea is great

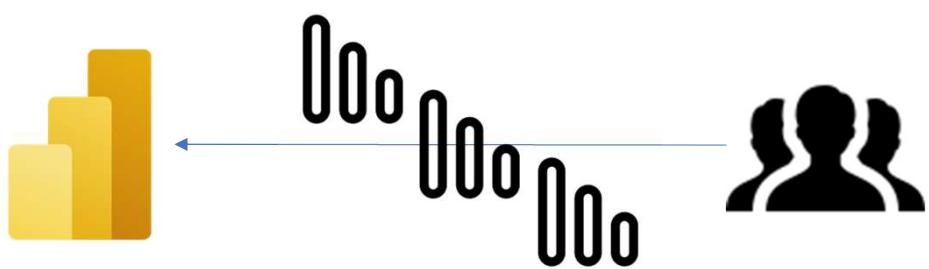
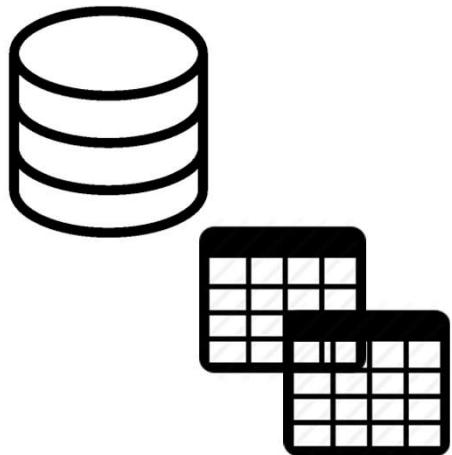
- If I'm the CDO, I can create a **single repository**
 - Structured
 - Unstructured
- Tell **all my platforms** to read and write there
 - Fabric workload
 - ADF pipelines
 - Databricks
 - Snowflake
- **Standard file format**



OneLake support for Import Mode Semantic Models



The usecase

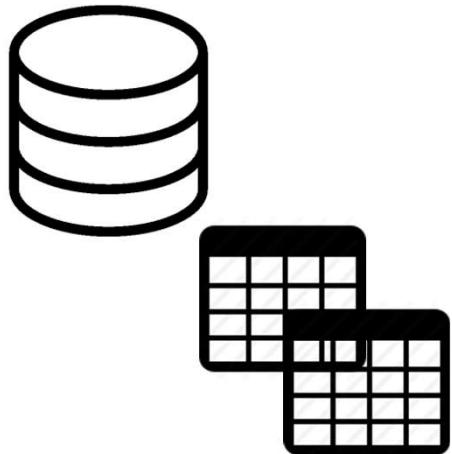


Business Layer / Semantic Layer

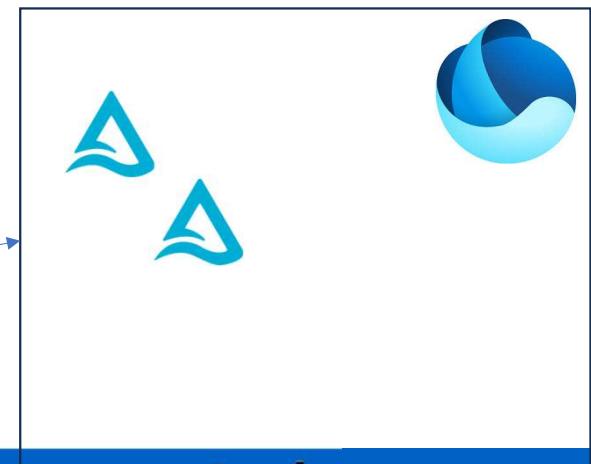
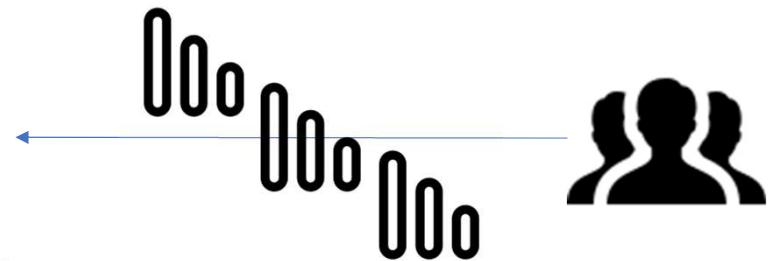
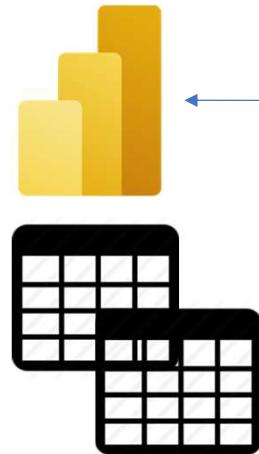
- + Cleansing
- + Naming Convention
- + Relationships
- + Calculated Columns
- + Calculated Tables

- Ask for CSV extract
- Use executeQueries API
- Go back to the source

A possible solution



- SQL
- R
- Scala
- Python



Sempy to leverage the “Golden Model”



Business Layer / Semantic Layer

- + Cleansing
- + Naming Convention
- + Relationships
- + Calculated Columns
- + Calculated Tables
- + Measures

```
%pip install semantic-link
```

```
import sempy.fabric as fabric
```

```
dataset = "Retail Analysis Sample PBIX"
```

```
fabric.evaluate_measure
```

```
(dataset, measure="Average Selling Area Size")
```

Use cases for Sempy

- Data Analysts / Data Scientists searching for **data already prepared**
- **Golden Model enrichment**
- Improved **data governance**
- **Strict data pipeline**
- Programmatic interface for **administrative tasks**
 - Models and partitions **refresh**
 - Dataset migration to **Direct Lake**
 - **Monitoring**
 - **Tenant management**
 - **API calls** in Python (alternative to PowerShell)
 - ...

OneLake loves other platforms



It's not a closed environment



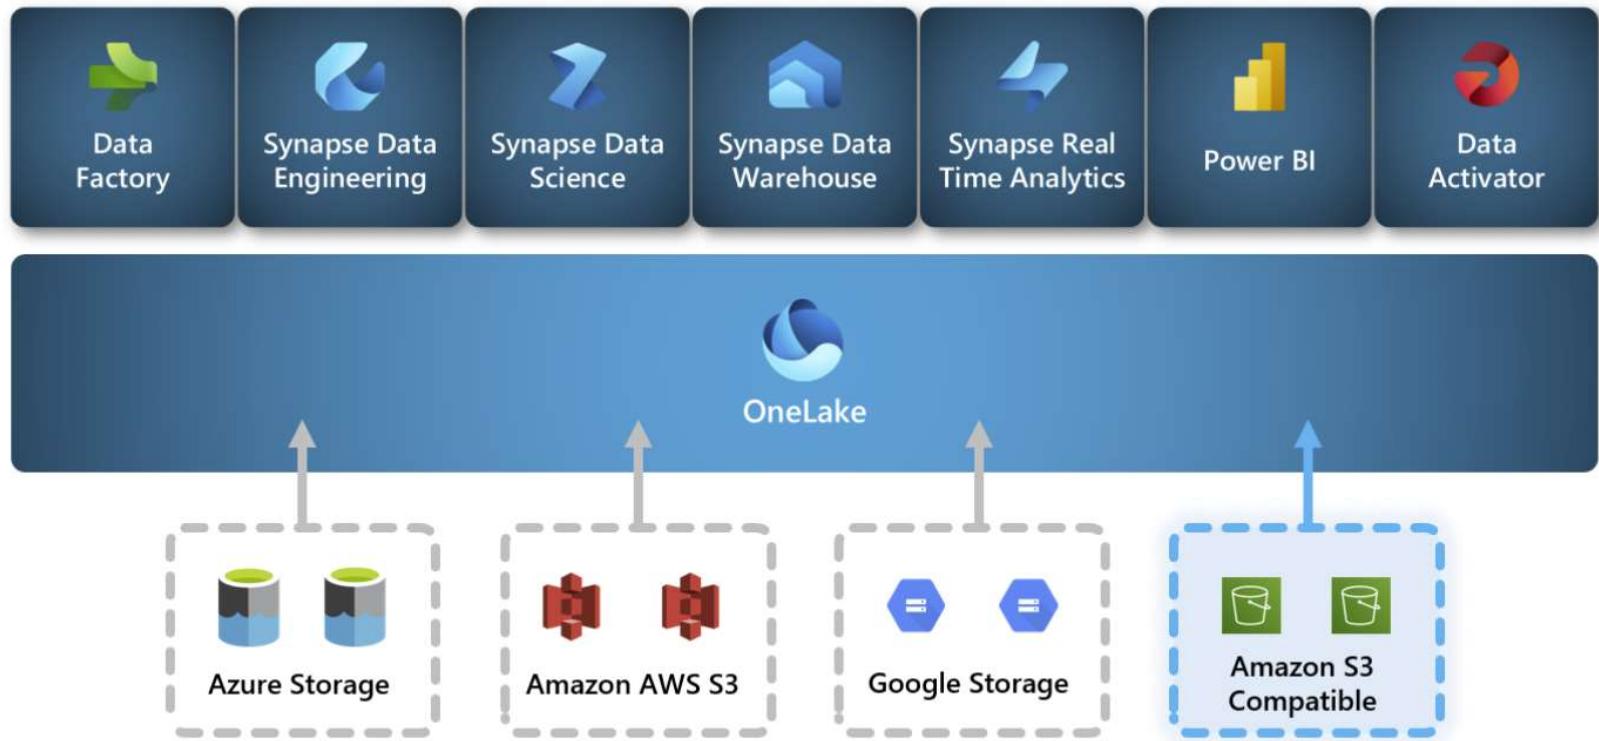
- OneLake support for **Iceberg**
- OneLake **Shortcuts over Snowflake**
- Snowflake can **store data in OneLake**
- Snowflake can **read data from OneLake**



- OneLake **shortcut for Databricks Unity Catalog**
- Federated Access **OneLake from Databricks**

One storage to rule them all





T-SQL is alive and well



Clone

- **Metadata only** copy (zero copy)
- **Time travel** back up to 30 days
- Good for **development and testing**
- Good for **backup and restore**
- **Archiving** system

```
CREATE TABLE  
    dbo.nyctaxi_clone  
AS CLONE OF  
    dbo.nyctaxi;  
OPTION  
(FOR TIMESTAMP AS OF  
'yyyy-mm-  
ddTHH:MM:SS.SSS');
```

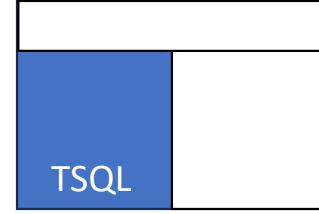
Time Travel (preview)

```
SELECT *
FROM dbo.Top10CustomersView
OPTION (
    FOR TIMESTAMP AS OF '2024-04-24T20:59:06.097'
);
```

Plan migration carefully



Migration is not straight forward



Not Supported

- SET XACT_ABORT ON
- MERGE
- PRIMARY KEY
- UNIQUE
- IDENTITY
- SET IDENTITY_INSERT
- Scalar Function
- ~~TRUNCATE TABLE~~

Latin1_General_100_BIN2_UTF8

```
SELECT *  
FROM dbo.MyTable;
```

<>

```
SELECT *  
FROM dbo.mytable;
```

```
SELECT COUNT(*)  
FROM dbo.MyTable  
WHERE Code = 'abc';
```

Id	Code
1	abc
2	abc
3	abc

Id	Code
1	Abc
2	abc
3	ABC

NVARCHAR

As of today, not supported.

Not a big deal since **Parquet manage automatically encoding** and collation is UTF-8.

View-table pipeline looses types

```
CREATE VIEW dbo.MyView AS
```

```
    SELECT
```

```
        AnIntField = CONVERT(int, t.F1)
```

```
        ,AString = CONVERT(varchar(20), t.F2)
```

```
        ,AndADateField = CONVERT(date, t.F3)
```

```
    FROM dbo.MyTable;
```

```
GO
```

```
CREATE TABLE dbo.MyNormalizedTable
```

```
(
```

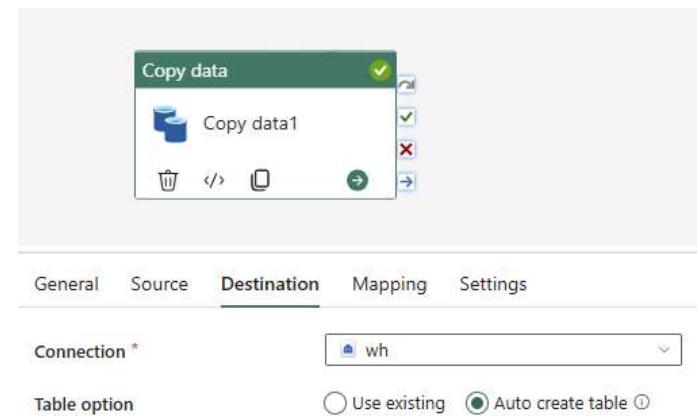
```
    AnIntField int NULL,
```

```
    Astring varchar (8000) NULL,
```

```
    AndADateField date NULL
```

```
)
```

```
GO
```



Workaround

```
SELECT *
INTO dbo.MyNormalizedTable
FROM dbo.MyView;
```

Or

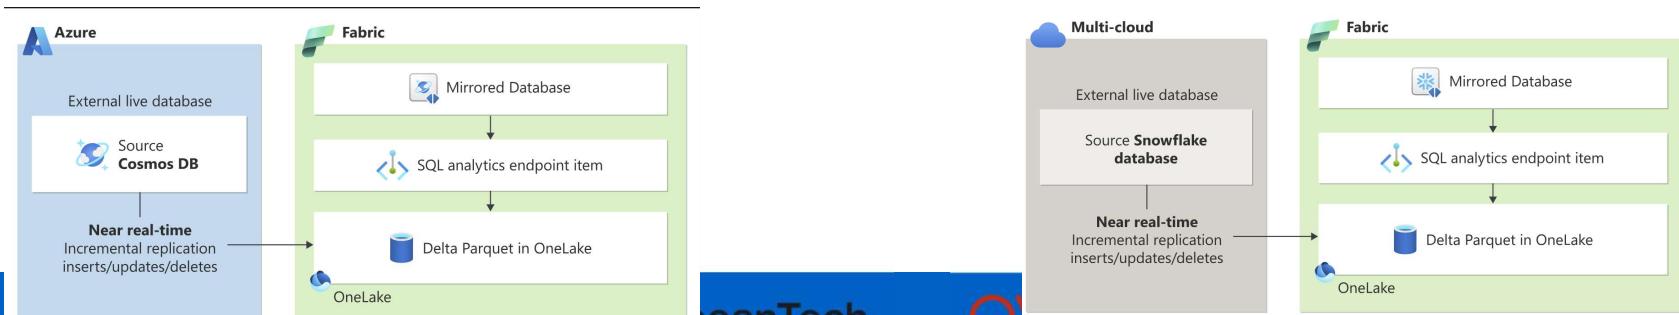
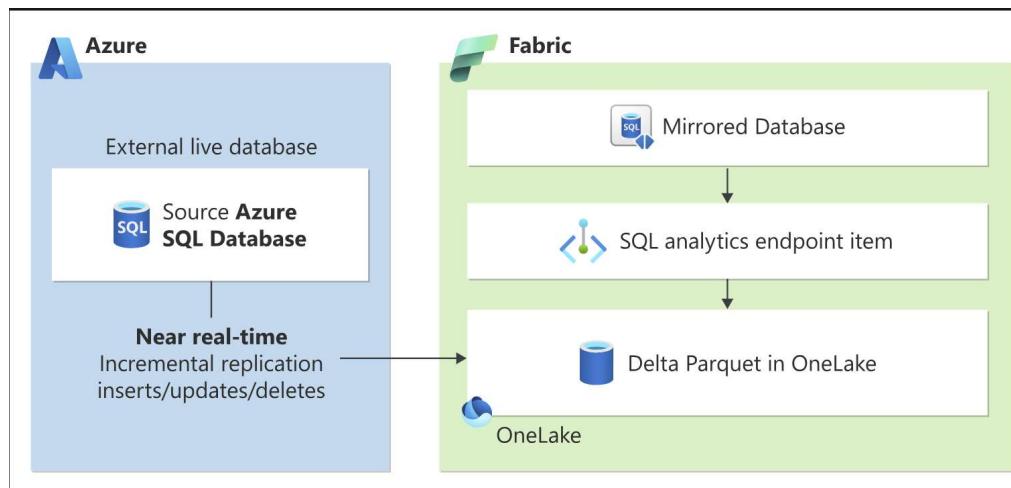
```
CREATE TABLE dbo.MyNormalizedTable AS  
SELECT *  
FROM dbo.MyView;
```



Pump data before playing



Mirroring to Fabric



Pros & Cons (On the SQL side)

- Managed ELT
- Near real-time replication
- Multi platform
- No extra cost 1TB per CU
- Preview feature
- No On-Prem SQL Server
- Public access only supported
- Cross tenant not supported
- ... it's not magic

Cons... continue

- A table cannot be mirrored if it does not have a **primary key rowstore clustered index**
- Following DDL operations aren't allowed on source tables
 - Switch/Split/Merge partition
 - Alter primary key
 - ~~Drop table~~
 - **Truncate table**
 - ~~Rename table~~
- When there is **DDL change**, a **complete data snapshot** is restarted for the changed table, and data is reseeded



Snowflake has less limitations

Another definitive Great Unifier...



Simply unmatched, truly limitless: Announcing Azure Synapse Analytics

By Rohan Kumar, Corporate Vice President, Azure Data

Posted on November 4, 2019
3 min read

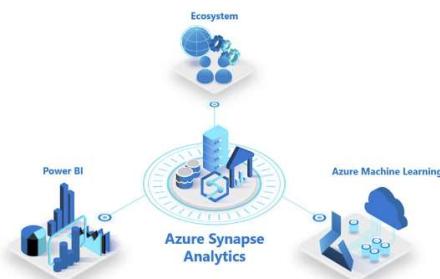


Big Data

Today, businesses are forced to maintain two types of analytical systems, data warehouses and data lakes. Data warehouses provide critical insights on business health. Data lakes can uncover important signals on customers, products, employees, and processes. Both are critical, yet operate independently of one another, which can lead to uninformed decisions. At the same time, businesses need to unlock insights from *all* their data to stay competitive and fuel innovation with purpose. Can a single cloud analytics service bridge this gap and enable the agility that businesses demand?

Azure Synapse Analytics

Today, we are announcing Azure Synapse Analytics, a limitless analytics service, that brings together enterprise data warehousing and Big Data analytics. It gives you the freedom to query data on your terms, using either serverless on-demand or provisioned resources, at scale. Azure Synapse brings these two worlds together with a unified experience to ingest, prepare, manage, and serve data for immediate business intelligence and machine learning needs.



Explore

Let us know what you think of Azure and what you would like to see in the future.

[Provide feedback](#)

Build your cloud computing and Azure skills with free courses by Microsoft Learn.

[Explore Azure learning](#)

What expert says:



[Microsoft Build 2023 Keynote Live Blog:
Introducing Fabric - Brent Ozar Unlimited®](#)

“...Microsoft’s data warehousing strategy that just can’t maintain focus for 3 years straight. From DATAAllegra to Parallel Data Warehouse to Hadoop to Analytics Platform System to Azure SQL Data Warehouse to Azure Synapse Analytics to Big Data Clusters, there’s something broken about the leadership vision here.”

“...I have this extra data that I need to join to my reports right now, and I don’t have the time to wait for the Microsoft Fabric admins to bring it in, so I’m just going to put it in this one place for now...”

Being a consultant could be frustrating

I sold “SQL Server-based” BIs for 10 years and I switched to Synapse and then to Fabric in less than 4 years.

“ONE PERSON CAN’T LEARN OR DO EVERYTHING”

Fabric is a team sport — DATA GOBLINS (data-goblins.com)



It seems a building site yet

- **Tons of features** every months
- Many features stay in **preview for a very long time** (even core)
 - GIT integration
 - Invoke Pipeline
 - Folders
 - Data Activator
 - Mirroring
 - PBIP and TMDL, New Card, Field Parameters...
- Flagship features only announced i.e. **OneSecurity**
- ~~Security features premium+ (+F64)~~



Database project support



DACPACs



Lakehouse or not Lakehouse



What experts say



Jovan Popovic

Principal Program Manager at Microsoft, working on Microsoft Fabric Warehouse. Worked on Azure Synapse, Azure SQL Azure SQL Managed Instance, and SQL Server.

Published Aug 10, 2023

[Choosing between Lakehouse and Warehouse in Microsoft Fabric \(linkedin.com\)](#)

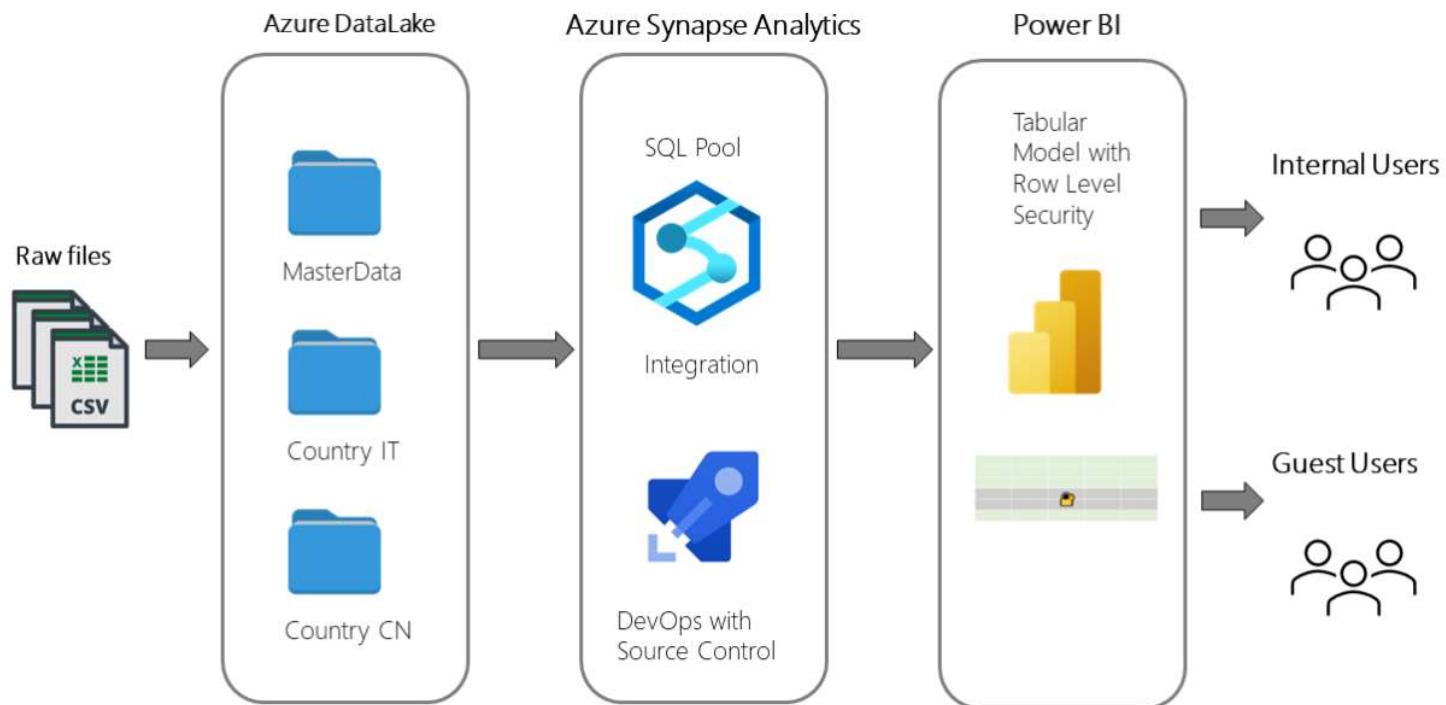
- **Language preference for data management** - The choice between Lakehouse and Warehouse depends on the preferred programming languages of your team. This includes both **syntax** and **functionalities** that are available in the language that you choose. If your team leans towards PySpark or Scala for data management, Lakehouse is the natural choice. On the other hand, Warehouse caters to those who favor T-SQL. This distinction may be less relevant for visual, low-code, or citizen developers. However, it's important to consider the features associated with different languages. For instance, if you need ML/regular expressions, Lakehouse is the choice for you because these features are available in Python/Scala. On the other hand, if you need multi-table transactions, some enterprise data management features, or fine-grained permissions on every object in the database with GRANT/DENY, RLS, data masking you should choose T-SQL language in the Warehouse.
- **Data format requirements** - If your data exists primarily in Delta format with relational structure, Warehouse seamlessly handles your needs. However, if you work with diverse formats like CSV, Parquet, or JSON, or you are using non-structured data, Lakehouse proves to be the more versatile solution.
- **Migration scenarios** - If your existing data solutions are implemented on SQL Server, Azure SQL, Synapse warehouse or other RDBMS systems, or involve a significant T-SQL code base that you wish to retain while transitioning to Fabric, the Warehouse is the preferable option. If you're migrating from a Spark/Databricks and have already implemented significant data processing logic in PySpark, Scala or SparkSQL notebooks, Lakehouse provides an easier path to migration within the Fabric ecosystem.

Even in official docs

Microsoft Fabric decision guide: choose a data store				
Article • 06/05/2024 • 10 contributors				
In this article				
Data store properties				
Scenarios				
Related content				
Use this reference guide and the example scenarios to help you choose a data store for your Microsoft Fabric workloads.				
Data store properties				
Expand table				
	Warehouse	Lakehouse	Power BI Data Mart	Eventhubs
Data volume	Unlimited	Unlimited	Up to 100 GB	Unlimited
Type of data	Structured	Unstructured, semi-structured, structured	Structured	Unstructured, semi-structured, structured
Primary developer persona	Data warehouse developer, SQL engineer	Data engineer, data scientist	Citizen developer	Citizen Data scientist, Data engineer, Data scientist, SQL engineer
Primary developer skill set	SQL	SparkSQL, PySpark, Spark SQL, R	No code, SQL	No code, KQL, SQL
Data organized by	Databases, schemas, and tables	Folders and files, databases, and tables	Database, tables, queries	Databases, schemas, and tables
Read operations	T-SQL, Spark (supports reading from tables using shortcuts, doesn't yet support accessing views, stored procedures, functions etc.)	Spark, T-SQL	Spark, T-SQL, Power BI	KQL, T-SQL, Spark, Power BI
Write operations	T-SQL	SparkSQL, PySpark, Spark SQL, R	Dataflows, T-SQL	KQL, Spark, connector ecosystem
Multi-table transactions	Yes	No	No	Yes, for multi-table ingestion. See update policy.
Primary development interface	SQL scripts	Spark notebooks, Spark job definitions	Power BI	KQL, Queryset, KQL Database
Security	Object level (table, view, function, stored procedure, etc.), column level, row level, DDL/DML, dynamic data masking	Row level, table level (when using T-SQL, none for Spark)	Built-in RLS editor	Row-level Security

[Fabric decision guide - choose a data store – Microsoft Fabric | Microsoft Learn](#)

Is it really a choice?



You'll likely need both...



- No Mirroring
- No cross-table transaction
- No T-SQL CRUD

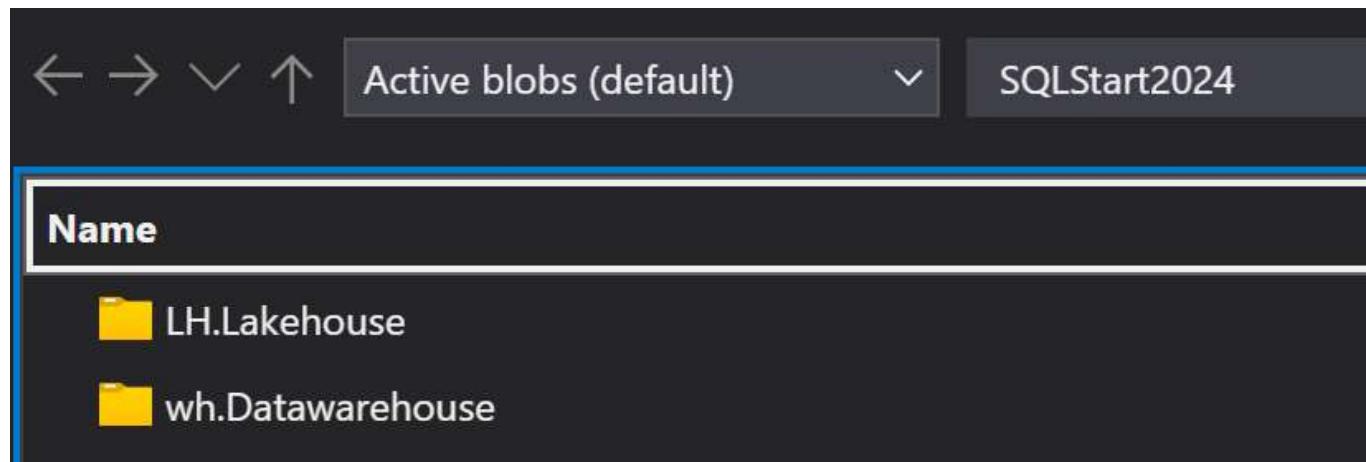


- No Shortcut
- No OneLake interoperability
- No security features
 - ADLS firewalled access

Dealing with Delta in a different way



Under the hood



Lakehouse style

```
df = spark.sql("SELECT 1 AS col")  
  
for i in range(1, 11):  
    df.write.format("delta").mode("append").  
    saveAsTable("TestVacumLH")
```

📁	_delta_log	6/9/2024 5:07:14 PM	Folder	11 item
📄	part-00000-29e9efd5-5c0f-4bff-884b-06ee4a4bb9ee-c000.snappy.parquet	6/9/2024 5:07:26 PM	parquet	781 B
📄	part-00000-3e33182e-74ad-4d19-91b1-8b2e6a5b74e2-c000.snappy.parquet	6/9/2024 5:07:28 PM	parquet	781 B
📄	part-00000-47b3d85f-2d10-409b-8594-0ae8fb4ebdef-c000.snappy.parquet	6/9/2024 5:07:30 PM	parquet	781 B
📄	part-00000-4d680093-de40-4e73-9218-bb7b55f690ad-c000.snappy.parquet	6/9/2024 5:07:15 PM	parquet	781 B
📄	part-00000-5ec6b95a-7c36-4a86-a13a-33780505117b-c000.snappy.parquet	6/9/2024 5:07:32 PM	parquet	781 B
📄	part-00000-5f6d19a6-4a95-4581-8184-b53d5e69558a-c000.snappy.parquet	6/9/2024 5:07:22 PM	parquet	781 B
📄	part-00000-7438f8bc-418e-436a-bdc2-b4fd22ccfffc3-c000.snappy.parquet	6/9/2024 5:07:18 PM	parquet	781 B
📄	part-00000-8a391aab-84cc-4516-9ad7-ebcfcd7ee16-c000.snappy.parquet	6/9/2024 5:07:24 PM	parquet	781 B
📄	part-00000-b0f803f5-7e1f-4302-94a3-271cbbc7b3b3-c000.snappy.parquet	6/9/2024 5:07:34 PM	parquet	781 B
📄	part-00000-e43b0ed5-9c09-4b99-b230-588fc0b30b68-c000.snappy.parquet	6/9/2024 5:07:20 PM	parquet	781 B

← testvacumlh (file view) > _delta_log

Name	Date modified	Type	Size
📄 000000000000000000000001.json	6/9/2024 5:07:15 PM	json	1 KB
📄 000000000000000000000002.json	6/9/2024 5:07:18 PM	json	713 B
📄 000000000000000000000003.json	6/9/2024 5:07:20 PM	json	713 B
📄 000000000000000000000004.json	6/9/2024 5:07:22 PM	json	713 B
📄 000000000000000000000005.json	6/9/2024 5:07:24 PM	json	713 B
📄 000000000000000000000006.json	6/9/2024 5:07:26 PM	json	713 B
📄 000000000000000000000007.json	6/9/2024 5:07:28 PM	json	713 B
📄 000000000000000000000008.json	6/9/2024 5:07:30 PM	json	713 B
📄 000000000000000000000009.json	6/9/2024 5:07:32 PM	json	713 B
📁 _temporary	6/9/2024 5:07:34 PM	Folder	0 item

Warehouse style

```
SELECT 1 AS col
INTO dbo.TestVacuumWH;
GO

DECLARE @i INT = 1;

WHILE @i < 10
BEGIN
    INSERT INTO
        dbo.TestVacuumWH
    VALUES(1);
    SET @i = @i + 1;
END;
```

The screenshot shows two tables in the Azure Storage Explorer:

Active blobs (default)

Name	Access Tier	Access Tier Last Modified	Last Modified
1340			6/9/2024 5:16 PM
1337			6/9/2024 5:16 PM
1334			6/9/2024 5:16 PM
1331			6/9/2024 5:16 PM
1328			6/9/2024 5:16 PM
1325			6/9/2024 5:16 PM
1322			6/9/2024 5:16 PM
1319			6/9/2024 5:16 PM
1317			6/9/2024 5:16 PM
1310			6/9/2024 5:16 PM

_delta_log

Name	Access Tier	Access Tier Last Modified	Last Modified
000000000000000009.checkpoint.parquet	Hot (inferred)		6/9/2024 5:16 PM
_last_checkpoint	Hot (inferred)		6/9/2024 5:16 PM
0000000000000000000009.json	Hot (inferred)		6/9/2024 5:16 PM
0000000000000000000004.json	Hot (inferred)		6/9/2024 5:16 PM
0000000000000000000005.json	Hot (inferred)		6/9/2024 5:16 PM
0000000000000000000006.json	Hot (inferred)		6/9/2024 5:16 PM
0000000000000000000007.json	Hot (inferred)		6/9/2024 5:16 PM
0000000000000000000008.json	Hot (inferred)		6/9/2024 5:16 PM
0000000000000000000000.json	Hot (inferred)		6/9/2024 5:16 PM
0000000000000000000001.json	Hot (inferred)		6/9/2024 5:16 PM
0000000000000000000002.json	Hot (inferred)		6/9/2024 5:16 PM
0000000000000000000003.json	Hot (inferred)		6/9/2024 5:16 PM

Last checkpoint

```
{  
  "version": 11,  
  "size": 15,  
  "sizeInBytes": 13909,  
  "numOfAddedFiles": 10  
}
```

Checkpoint.parquet

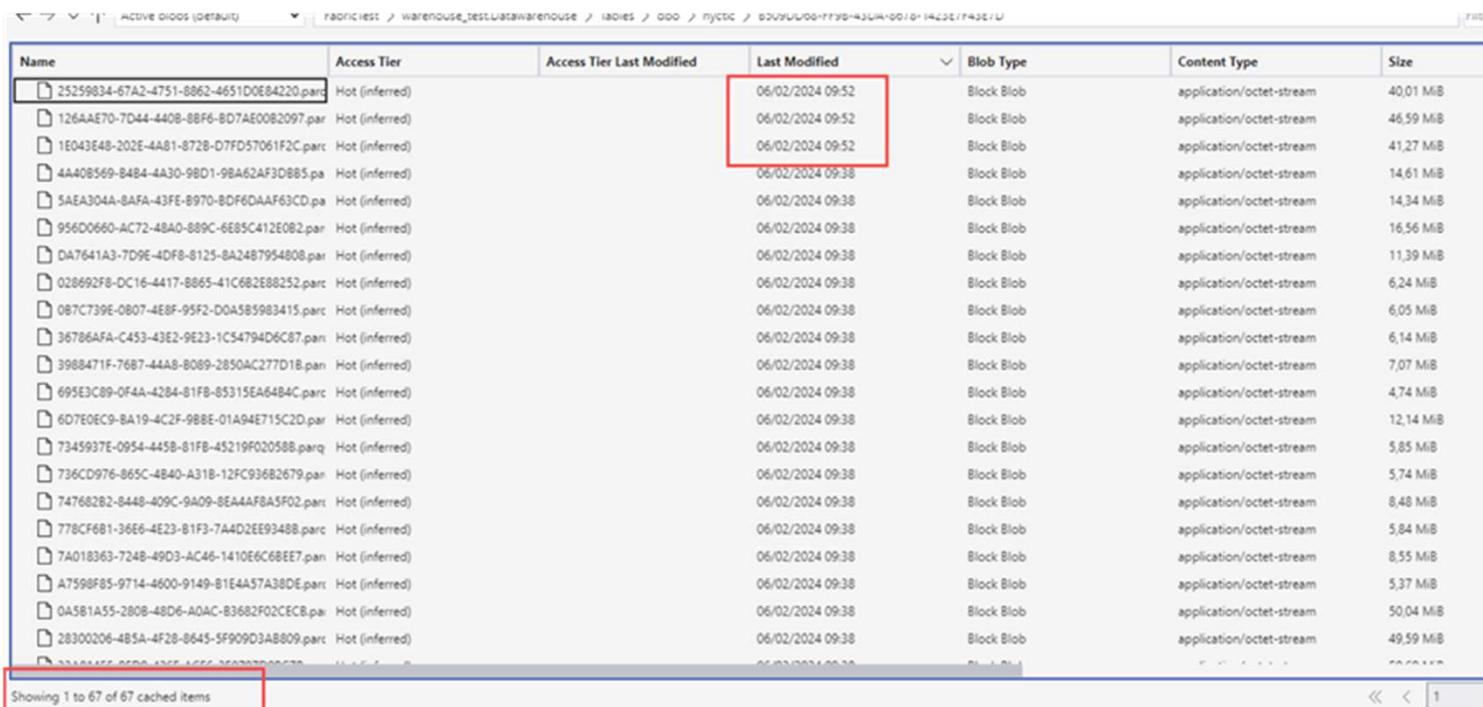
```
{"txn": {"appId": "424c3456-c3dc-e5f8-530f-135019cd5974", "version": 1530, "lastUpdated": null}, "protocol": null, "metaData": null, "add": null, "remove": null}  
{"txn": null, "protocol": null, "metaData": null, "add": {"path": "E7285DBA-A770-473F-8ED7-D7D849AF3BB3/0/1530/3893A81A-6437-43CB-B1FE-CB1D4BE98D84.parquet"},  
{"txn": null, "protocol": null, "metaData": null, "add": {"path": "E7285DBA-A770-473F-8ED7-D7D849AF3BB3/0/1524/60547DA2-DBFC-44CD-91E6-3E13517ACC05.parquet"},  
{"txn": null, "protocol": null, "metaData": null, "add": {"path": "E7285DBA-A770-473F-8ED7-D7D849AF3BB3/0/1521/9A4F0D97-5F62-469A-A454-6F8DD2571F4A.parquet"},  
{"txn": null, "protocol": null, "metaData": null, "add": {"path": "E7285DBA-A770-473F-8ED7-D7D849AF3BB3/0/1518/EA152942-9558-433D-BD7D-BE66C8907563.parquet"},  
{"txn": null, "protocol": null, "metaData": null, "add": {"path": "E7285DBA-A770-473F-8ED7-D7D849AF3BB3/0/1515/6A4A7BC2-CCA8-48A5-827E-256D8082190B.parquet"},  
{"txn": null, "protocol": null, "metaData": null, "add": {"path": "E7285DBA-A770-473F-8ED7-D7D849AF3BB3/0/1512/C34BE407-2192-4C72-920C-8FA21C85EC09.parquet"},  
{"txn": null, "protocol": null, "metaData": null, "add": {"path": "E7285DBA-A770-473F-8ED7-D7D849AF3BB3/0/1509/CF0614B4-DD85-43E8-B53D-AD95F991A481.parquet"},  
{"txn": null, "protocol": null, "metaData": null, "add": {"path": "E7285DBA-A770-473F-8ED7-D7D849AF3BB3/0/1506/7E96F95B-AF37-441E-92F7-DE95561EFAAC.parquet"},  
{"txn": null, "protocol": null, "metaData": null, "add": {"path": "E7285DBA-A770-473F-8ED7-D7D849AF3BB3/0/1504/32D7FE8B-3CDD-49FF-A0AE-FC2280D11C2E.parquet"},  
{"txn": null, "protocol": {"minReaderVersion": 1, "minWriterVersion": 1, "readerFeatures": null, "writerFeatures": null}, "metaData": null, "add": null, "remove": null}  
{"txn": null, "protocol": null, "metaData": {"id": "E7285DBA-A770-473F-8ED7-D7D849AF3BB3", "name": "TestVacumWH", "description": null, "format": {"provider": "parquet", "compression": "none", "blockSize": 1048576}}, "add": {"path": "E7285DBA-A770-473F-8ED7-D7D849AF3BB3/0/1497/421E29B3-8296-4E5D-91CF-9302B295CD71.parquet"},  
|
```

Later on...

The screenshot shows a file listing interface. At the top, there are navigation icons (left arrow, right arrow, down arrow, up arrow), a dropdown menu set to "Active blobs (default)", and a breadcrumb path "SQLStart2024 > wh.Datawarehouse". Below this is a table with two columns: "Name" and "Access Tier". The "Name" column lists various file names, and the "Access Tier" column shows "Hot (inferred)" for all entries. The file "00000000000000000011.json" is highlighted with a red border.

Name	Access Tier
00000000000000000010.json	Hot (inferred)
00000000000000000011.json	Hot (inferred)
00000000000000000009.checkpoint.parquet	Hot (inferred)
_last_checkpoint	Hot (inferred)
00000000000000000009.json	Hot (inferred)
00000000000000000005.json	Hot (inferred)
00000000000000000006.json	Hot (inferred)
00000000000000000007.json	Hot (inferred)
00000000000000000008.json	Hot (inferred)
00000000000000000000.json	Hot (inferred)
00000000000000000001.json	Hot (inferred)
00000000000000000002.json	Hot (inferred)
00000000000000000003.json	Hot (inferred)
00000000000000000004.json	Hot (inferred)

WH - Automatic Data Compaction



Name	Access Tier	Access Tier Last Modified	Last Modified	Block Type	Content Type	Size
25259834-67A2-4751-8862-4651D0E84220.parc	Hot (inferred)		06/02/2024 09:52	Block Blob	application/octet-stream	40,01 MiB
126AAE70-7D44-4408-8BF6-BD7AE00B2097.parc	Hot (inferred)		06/02/2024 09:52	Block Blob	application/octet-stream	46,59 MiB
1E043E48-202E-4A81-8728-D7FD57061F2C.parc	Hot (inferred)		06/02/2024 09:52	Block Blob	application/octet-stream	41,27 MiB
4A40B569-B4B4-4A30-9BD1-9BA62AF3DBB5.parc	Hot (inferred)		06/02/2024 09:38	Block Blob	application/octet-stream	14,61 MiB
5AEE304A-8AFA-43FE-B970-BDF6DAAF63CD.parc	Hot (inferred)		06/02/2024 09:38	Block Blob	application/octet-stream	14,34 MiB
956D0660-AC72-48A0-889C-6E85C412E082.parc	Hot (inferred)		06/02/2024 09:38	Block Blob	application/octet-stream	16,56 MiB
DA7641A3-7D9E-4DF8-8125-8A2487954808.parc	Hot (inferred)		06/02/2024 09:38	Block Blob	application/octet-stream	11,39 MiB
028692F8-DC16-4417-8865-41C6B2E88252.parc	Hot (inferred)		06/02/2024 09:38	Block Blob	application/octet-stream	6,24 MiB
0B7C739E-0807-4E8F-95F2-D0A5B5983415.parc	Hot (inferred)		06/02/2024 09:38	Block Blob	application/octet-stream	6,05 MiB
36786AFA-C453-43E2-9E23-1C54794D6C87.parc	Hot (inferred)		06/02/2024 09:38	Block Blob	application/octet-stream	6,14 MiB
3988471F-76B7-44AB-B089-2850AC277D1B.parc	Hot (inferred)		06/02/2024 09:38	Block Blob	application/octet-stream	7,07 MiB
695E3C89-0F4A-4284-81FB-85315EA64B4C.parc	Hot (inferred)		06/02/2024 09:38	Block Blob	application/octet-stream	4,74 MiB
6D7E0EC9-BA19-4C2F-9BBE-01A94E715C2D.parc	Hot (inferred)		06/02/2024 09:38	Block Blob	application/octet-stream	12,14 MiB
7345937E-0954-4458-81FB-45219F020588.parc	Hot (inferred)		06/02/2024 09:38	Block Blob	application/octet-stream	5,85 MiB
736CD976-865C-4B40-A31B-12FC93682679.parc	Hot (inferred)		06/02/2024 09:38	Block Blob	application/octet-stream	5,74 MiB
747682B2-8448-409C-9A09-8EA4AF8A5F02.parc	Hot (inferred)		06/02/2024 09:38	Block Blob	application/octet-stream	8,48 MiB
778CF6B1-3666-4E23-B1F3-7A4D2EE9348B.parc	Hot (inferred)		06/02/2024 09:38	Block Blob	application/octet-stream	5,84 MiB
7A018363-7248-49D3-AC46-1410E6C68EE7.parc	Hot (inferred)		06/02/2024 09:38	Block Blob	application/octet-stream	8,55 MiB
A7598F85-9714-4600-9149-B1E4A57A38DE.parc	Hot (inferred)		06/02/2024 09:38	Block Blob	application/octet-stream	5,37 MiB
0A581A55-2808-48D6-A0AC-B3682F02CECB.parc	Hot (inferred)		06/02/2024 09:38	Block Blob	application/octet-stream	50,04 MiB
28300206-4B5A-4F28-8645-5F909D3AB809.parc	Hot (inferred)		06/02/2024 09:38	Block Blob	application/octet-stream	49,59 MiB
32A8AEEC-0E9A-47E6-8E03-37037000C780.parc	Hot (inferred)		06/02/2024 09:38	Block Blob	application/octet-stream	49,59 MiB

%%scal

OPTIMIZE <table|fileOrFolderPath>;

```
{\"txn\":null,\"add\":null,\"remove\":[\"path\":\"part-00000-4d680093-de40-4e73-9218-bb7b55f690ad-c000.snappy.parquet\",\"d  
{\"txn\":null,\"add\":null,\"remove\":[\"path\":\"part-00000-e43b0ed5-c09-4b99-b230-588fc0b30b68-c000.snappy.parquet\",\"d  
{\"txn\":null,\"add\":null,\"remove\":[\"path\":\"part-00000-5fdff196-aa45-4951-8184-b5d35ef6958a-c000.snappy.parquet\",\"d  
{\"txn\":null,\"add\":null,\"remove\":[\"path\":\"part-00000-5ec6b95a-7c36-4a86-a13a-33780505117b-c000.snappy.parquet\",\"d  
{\"txn\":null,\"add\":[\"path\":\"part-00000-85d7f75a-3a61-4a87-b9e7-92e1dc2142e-c000.snappy.parquet\"],\"partitionValues  
{\"txn\":null,\"add\":[\"path\":\"part-00000-3e33182e-74a4-4d19-91b1-82b2ea65b74e2-c000.snappy.parquet\",\"d  
{\"txn\":null,\"add\":null,\"remove\":[\"path\":\"part-00000-299ef5d-5c0f-4bf7-884b-06ee4a4bb9ee-c000.snappy.parquet\",\"d  
{\"txn\":null,\"add\":null,\"remove\":[\"path\":\"part-00000-47bd38fs-2df-409b-8594-08fb4ebdec-c000.snappy.parquet\",\"d  
{\"txn\":null,\"add\":null,\"remove\":[\"path\":\"part-00000-7438fb8c-418e-436a-bdc2-b4fd22cfdfc3-c000.snappy.parquet\",\"d  
{\"txn\":null,\"add\":null,\"remove\":[\"null\", \"metaData\"],\"protocol\":{\"minReaderVersion\":1, \"minWriterVersion\":2, \"read  
{\"txn\":null,\"add\":null,\"remove\":[\"null\", \"metaData\"],\"id\":10453dd2-8228-4af5-bca0-80af45377cf\", \"name\":null, \"descri  
{\"txn\":null,\"add\":null,\"remove\":[\"path\":\"part-00000-b0f803f5-7e1f-4302-94a3-271ccb7b3b3-c000.snappy.parquet\",\"d  
{\"txn\":null,\"add\":null,\"remove\":[\"path\":\"part-00000-8a391aab-84cc-4516-9ad7-ebcfcdf7ee16-c000.snappy.parquet\",\"d
```

```
≡ 0000000000000000000010.checkpoint  
1  [{"col":1}]  
2  {"col":1}  
3  {"col":1}  
4  {"col":1}  
5  {"col":1}  
6  {"col":1}  
7  {"col":1}  
8  {"col":1}  
9  {"col":1}  
10 {"col":1}
```

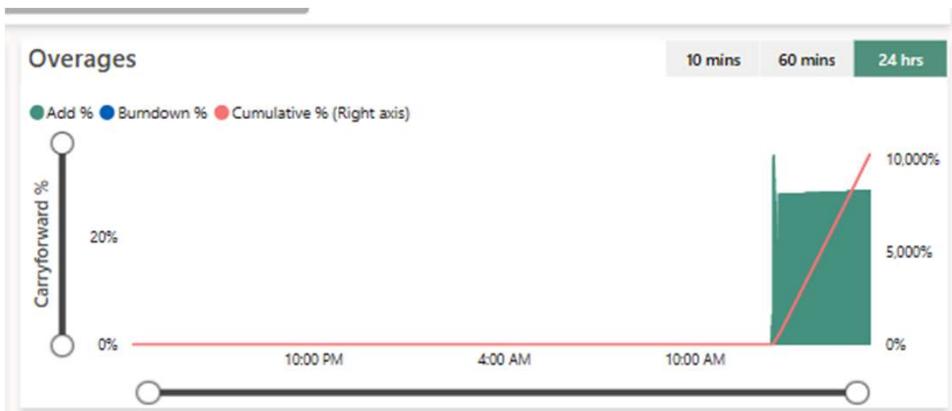
A close-up, low-angle shot of Spider-Man's upper body and arms. He is wearing his iconic red and blue suit with the spider emblem on the chest. His right hand is pointing towards the viewer, while his left hand is clenched in a fist. The background is a plain, light color.

With great power comes
great responsibility

Burns as much as it can

Fabric SKU	Equivalent Premium SKU	Baseline Capacity Units (CU)	Burstable Scale Factor
F2		2	1x - 32x
F4		4	1x - 16x
F8		8	1x - 12x
F16		16	1x - 12x
F32		32	1x - 12x
F64	P1	64	1x - 12x
F128	P2	128	1x - 12x
F256	P3	256	1x - 12x
F512	P4	512	1x - 12x
F1024	P5	1024	1x - 12x
F2048		2048	1x - 12x

And then...

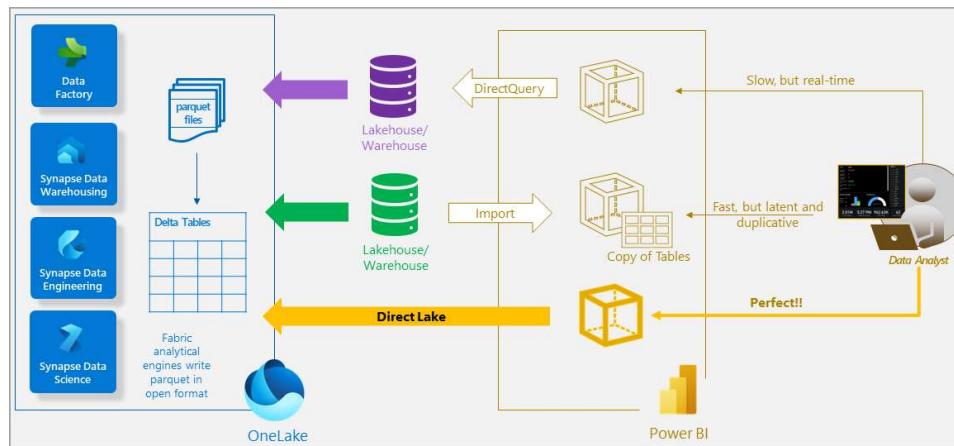


Other things I like

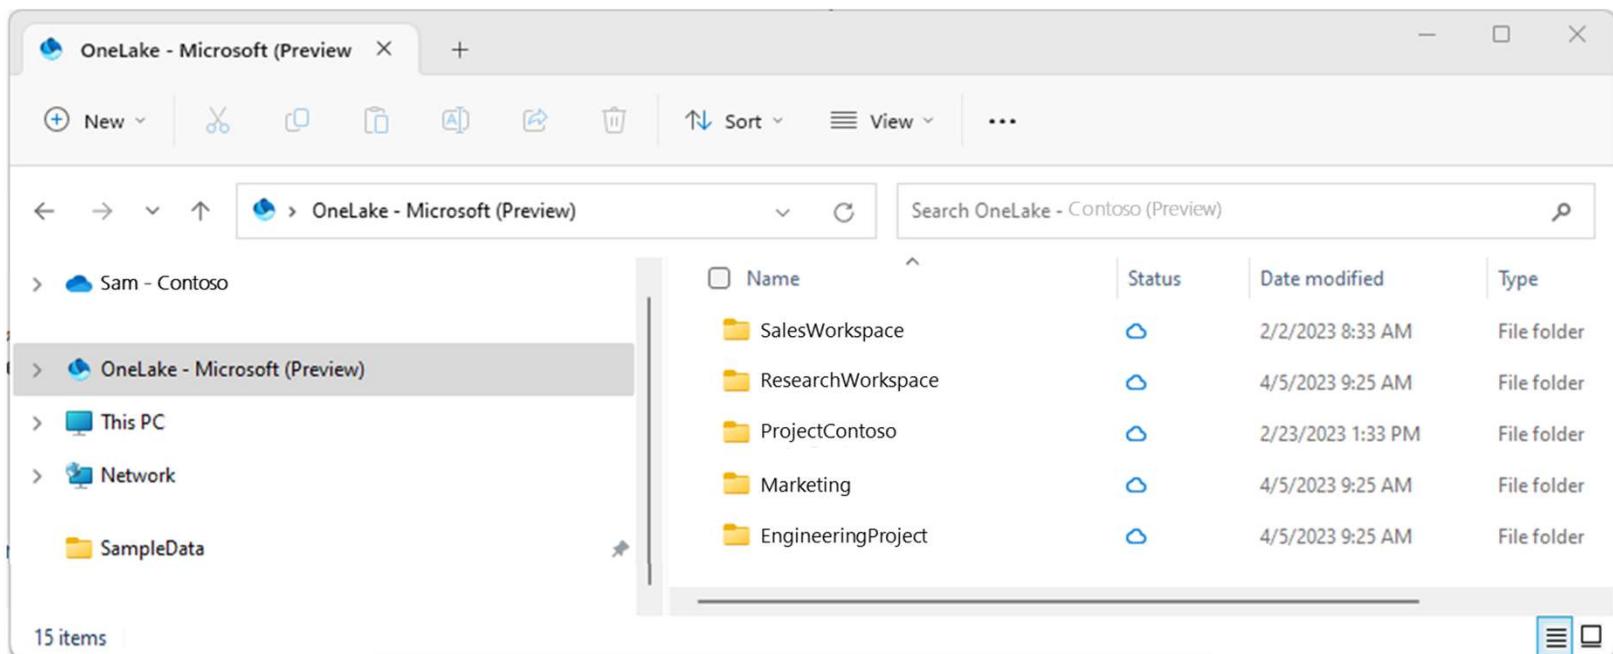


Direct Lake

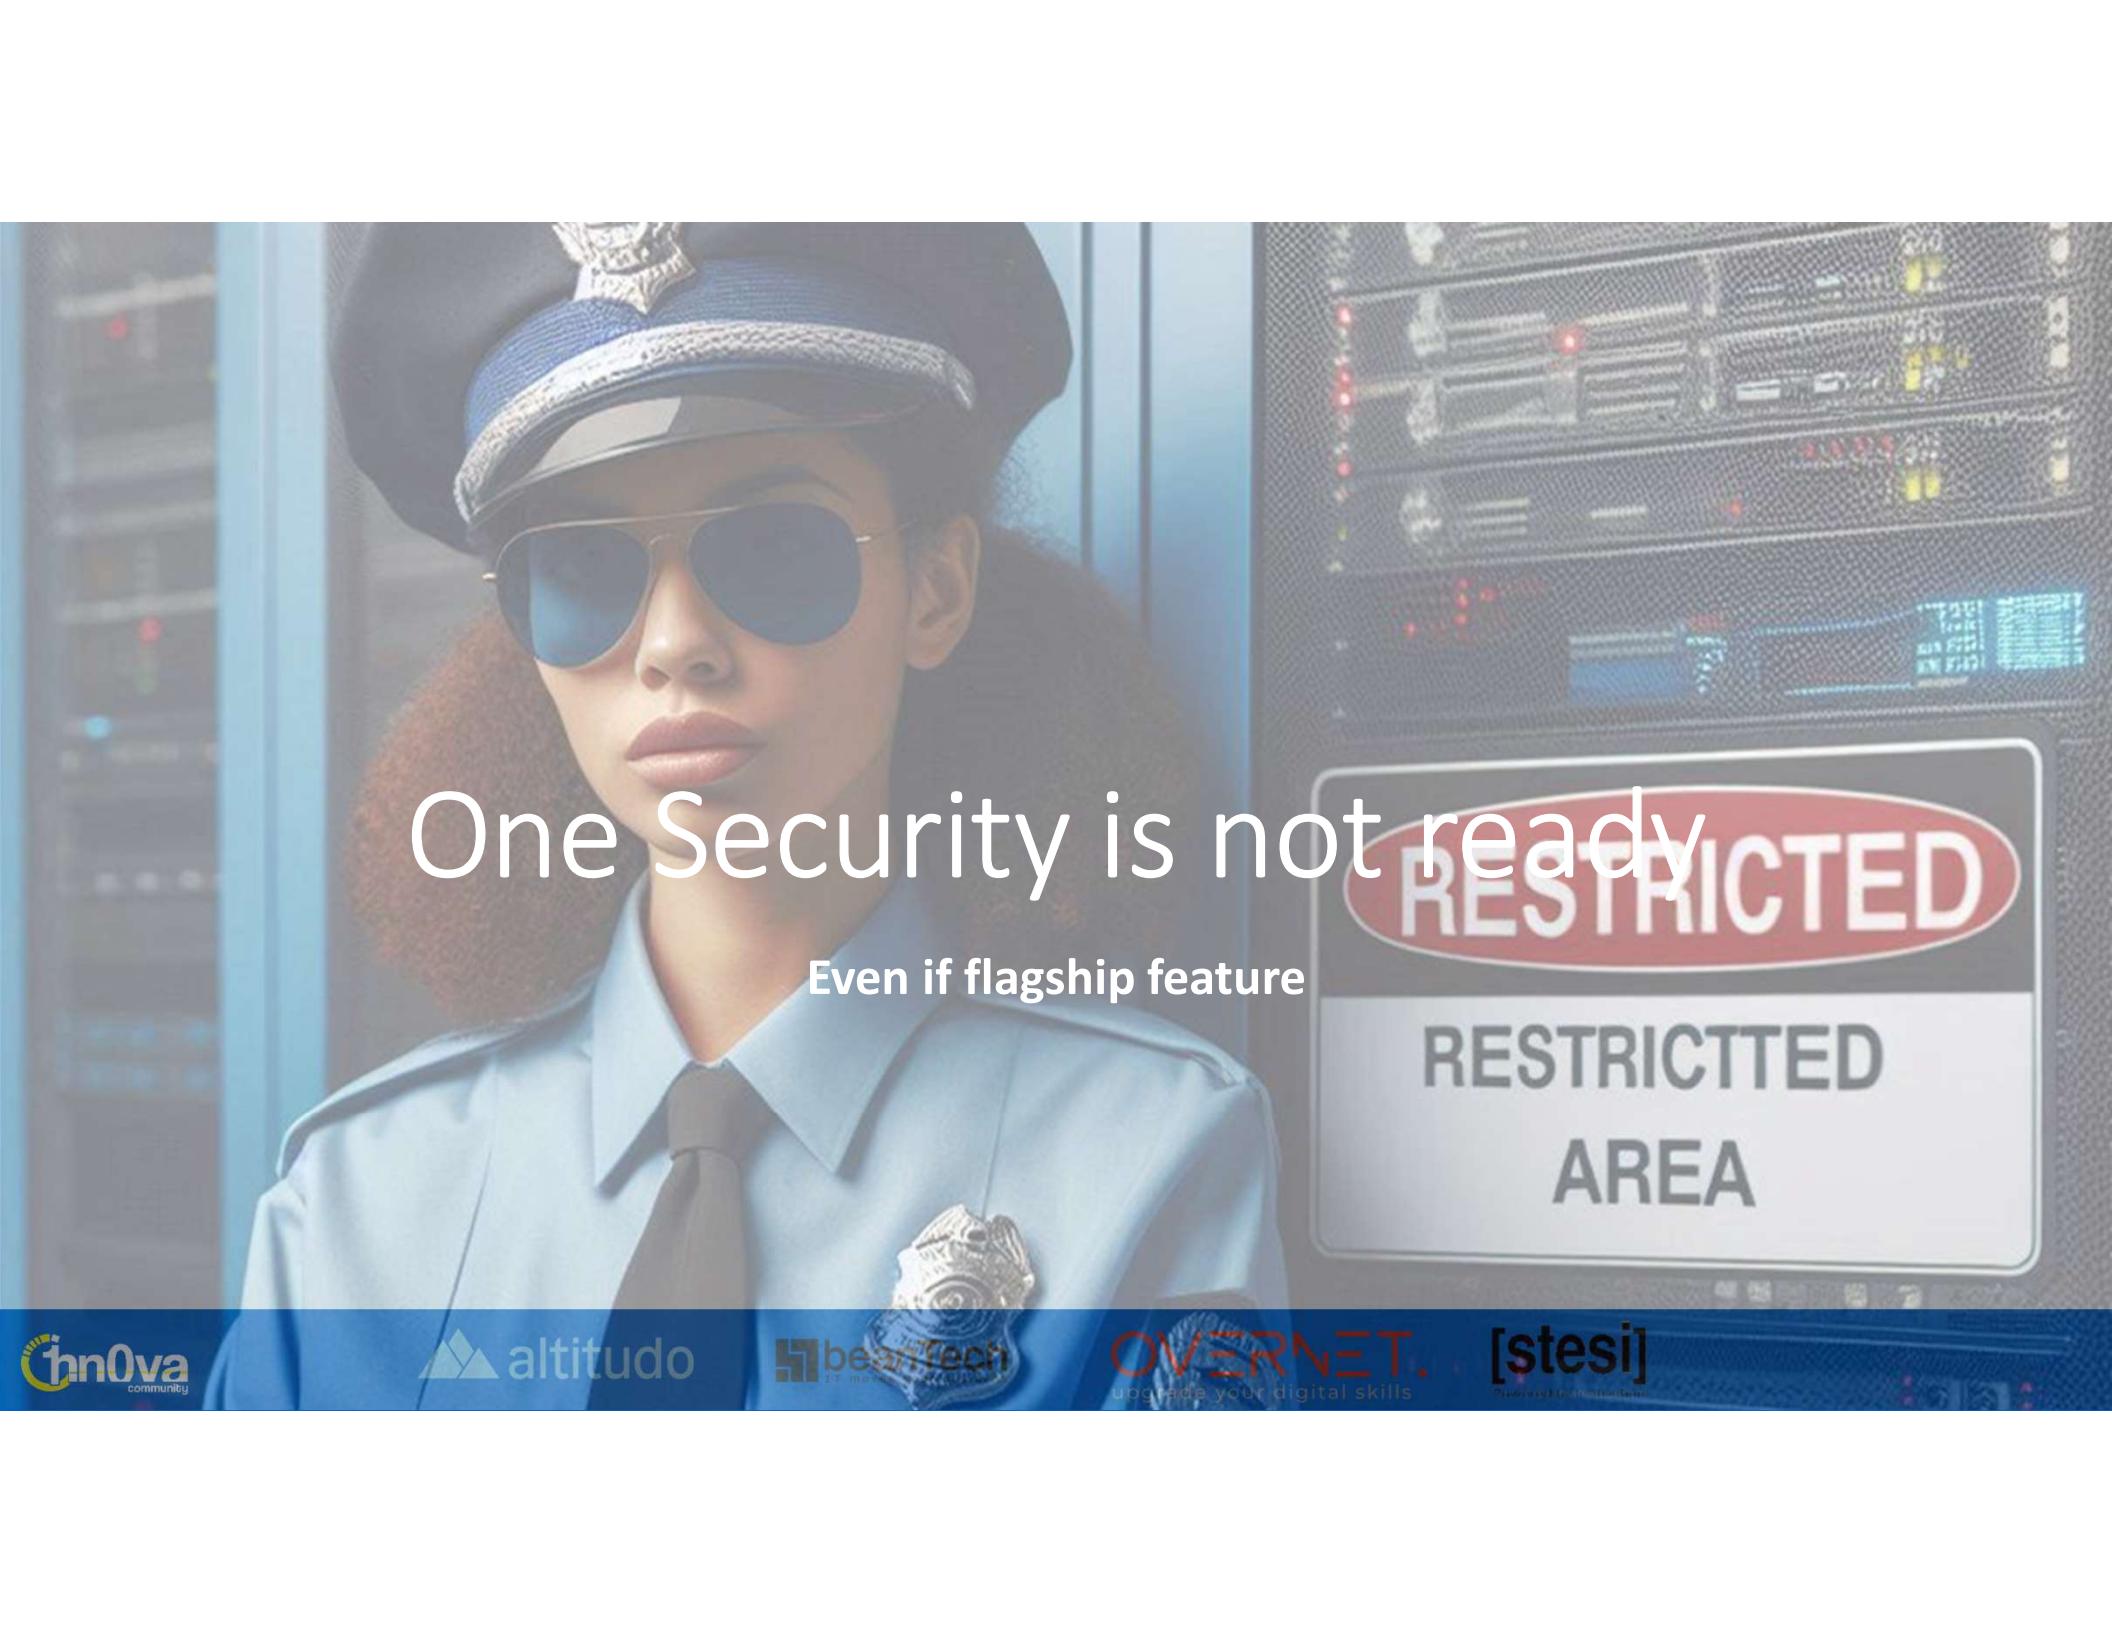
- Really promising
- No need to refresh when star schema is ready in store
- Could fallback to DQ
- Import could generate better plans



OneLake Windows Explorer Add-in







One Security is not ready
Even if flagship feature

WORK



IN PROGRESS...

CICD

Useful Links

- [Access Fabric data locally with OneLake file explorer](#)
- [Transitioning from ADLS to OneLake](#)
- [Datasets - Execute Queries - REST API \(Power BI Power BI REST APIs\)](#)
- [Exporting Power BI Reports And Sharing With Granular Access Control In Fabric](#)
- [Limitations for Fabric mirrored databases](#)
- [Create shortcuts to on-premises data](#)
- [Microsoft Idea Limit Bursting](#)
- [Apache XTable™ \(Incubating\)](#)

Questions?



Thank You!



Like

- Clone table wh
- Mirroring even if limited
- Concept
 - Great unifier
 - Simplify deployment and connection
 - A new sql server monolite
 - Onelake idea
- Shortcuts
- Start stop capacity
- Direct lake
- Import mode support for onelake
- Onelake and azure storage explorer
- Onelake addin win explorer

Don't like

- Another umbrella definitive data product
- Cantiere in movimento
- Visual studio integration with WH
 - Schema compare see everything views without access db without access
 - VS SSDT and ADS SSDT
- Private endpoint on ADLS only for spark with shortcut
- Hidden engine behaviour (bursting smoothing)
<https://ideas.fabric.microsoft.com/ideas/idea/?ideaid=22575b35-86eb-ee11-a73e-00224855381e>
- Need to do stuff differently in lh and wh (auto compact/optimize vacuum, zordering vordering managed/unmanaged)
- Direct lake fallback to dr
- Create Shortcuts from wh not supported – not clear the role of the lh
- Different possibilities between workloads
- Cicd
- One security not already there

Demo OneLake

- See warehouse and lakehouse behind the scenes
- See what happens if you save a power bi dataset in onelake
- Shortcut slow to a far distance
- Shortcut on-prem

Demo: OneLake Import Mode and Sempy

- Show a semantic model import mode extracting and the result
- Show sempy doing a query
- Show sempy