

Torino Data Meeetup

IT

In collaborazione con TOOLBOX

Via Agostino da Montefeltro, 2, 10134 TO



PBIUG ITALY

18:30

Novità su Power BI e FABCON

con RICCARDO PERICO & VICTOR RIVAS

19:00

Power Query: The Dark Knight

con LORENZO VERCELLATI

20:00

Chiacchiere

momento di condivisione



Follow us on:



Completamente **GRATIS**

Riccardo Perico



riccardo.perico@lucient.it

 github.com/R1k91

 [@r1k91.bsky.social](https://bsky.app/profile/@r1k91.bsky.social)

 [@R1k91](https://twitter.com/R1k91)

 [linkedin.com/in/riccardo-perico](https://www.linkedin.com/in/riccardo-perico)


 medium.com/riccardo-perico




BI Architect @ **Lucient**

+15 years in Microsoft “Data Realm”

Giving my little contribution to the community



Victor Rivas



victormrivasc@gmail.com

[in linkedin.com/in/victormrivasc](https://www.linkedin.com/in/victormrivasc)

[Web https://okviz.com/synoptic-panel/#in-action](https://okviz.com/synoptic-panel/#in-action)




Content Creator @




Teacher @




Politecnico
di Torino




0001




0008




0014




0023




0028




0033




0041




0048




0052




0056




0062




0068




0072




0076




0079




0082




0086




0096




0099




0102




0111




0116




0123



0129




0132



132 announcements!!!



Fabric March 2026 Feature Summary




March 18, 2026 by [Katie Murray](#) 103,681 Views




[Fabric March 2026 Feature Summary | Microsoft Fabric Blog | Microsoft Fabric](#)

Welcome to the **Fabric March 2026 Feature Summary**—and welcome to **FabCon!**

As we kick off FabCon, this update captures the momentum we're seeing across the Fabric platform and the conversations happening with customers and partners right now. March brings a wide range of enhancements across governance, data engineering, real-time



Fabric Platform



CICD

Branched workspace with Git integration (Preview)

Branched workspace is a new developer experience designed to help you manage workspaces during a branch-out flow. With clearer visual cues, you can better understand workspace relationships and work more confidently on new features. This feature will be released by the end of March 2024.

- Foxtrot Unit
- > Beta Company
- Yankee Team
- Feature-alpha
- Alpha-advanced
- Alpha-optimized

Selective branching with Git integration (Preview)

Fabric Git Integration Branch-out with selective branching introduces a more focused branch-out experience in Fabric. Developers can select only the items they need for a feature, reducing clutter in the target workspace, improving reliability, and accelerating time-to-code. By working with a smaller, purpose-built workspace, developers can iterate faster and with greater confidence.

Select items for your workspace

Latest items from Git branch: contoso-main. You can always add more items after creating the new branch. [Learn more](#)

- Search
- Item ↑
- ☐ Annual Summary Report
 - ☐ Data Lake V2 for Finance
 - ☐ Draft Financial Notebook
 - ☒ ETL Pipeline for Financial Data
 - ☐ Financial Analysis Notebook
 - ☐ Financial Data Model
 - ☒ Investment Experiment Notebook

Connection reference item type in Variable Library (Preview)

The new connection reference item type in Variable Library introduces a new way to manage external data connections in Microsoft Fabric. This new variable type lets you reference existing connections—such as Azure SQL or Snowflake—by storing a connection ID in the Variable Library, instead of embedding static connection strings in code.

Variables_1

Variables

☐ Name *

Note

Type

☐ demoCloudConnection

Note

Type

Connection reference

Default value set *

Active

Select connection

Basic

Boolean

DateTime

Guid

Integer

Number


String

Other

✓ Connection reference (preview)

CI/CD deployments from the CLI—deploy workspaces in One Command

A new **deploy** command integrates the Fabric CI/CD Python library directly into the Fabric CLI, enabling full workspace deployments—including item rebinding and configuration—from a single command. Teams can run deployments from their terminal, GitHub Actions, or Azure DevOps pipelines. Combined with Service Principal authentication and federated credentials for GitHub OIDC, this enables zero-touch, Git-based promotion workflows that fit modern DevOps practices—no custom scripts or additional tools required.



Agents



Fabric MCP AI code assistants (Generally Available)

The Fabric Local MCP is an open-source MCP server that runs on your machine. This solution integrates AI coding assistants with the comprehensive Fabric API, offering OpenAPI specifications, best-practice guidelines, item definition schemas, and example payloads to enable agents to produce precise, production-ready code while minimizing errors.

OneLake tools enable live file operations including upload, download, table inspection, and item creation. This update introduces integrated authentication, automatic retry, production SLAs, and telemetry. Install via `npx @microsoft/fabric-mcp` in any MCP-compatible client—it works with VS Code, Claude Desktop, Cursor, and more. Fabric Local MCP on GitHub.Fabric

Fabric Remote MCP Server: AI agents operate directly in your Fabric environment

Fabric Remote MCP is a cloud-hosted MCP server that allows AI agents to perform real operations in your Fabric environment—create workspaces, manage permissions, work with item definitions, and more. No local installation is required. Agents authenticate via Entra ID and operate within your existing RBAC boundaries, with every tool invocation recorded in audit logs.


The preview launches with capabilities spanning workspace management, item CRUD and definitions, and permission management. It works with any MCP-compatible client, including GitHub Copilot, Cursor, and Claude Desktop.

Learn more in this blog post: [Introducing Fabric MCP \(Preview\)](#).


Fabric CLI as an Execution Layer for AI Agents

Fabric CLI is designed to work well with AI agents. A structured agent instructions file and a dedicated Fabric CLI Skill provide AI assistants like GitHub Copilot and Claude with the context they need to generate correct CLI commands from natural language. Improved error messages with actionable guidance help agents self-correct, and the interactive REPL mode enables persistent terminal sessions for multi-step agent workflows.



Using a CLI as the execution layer for AI agents is an emerging industry pattern—instead of agents calling raw REST APIs (which require extensive token-heavy context about endpoints, auth, and payloads), agents issue concise CLI commands that encapsulate that complexity, making AI-driven Fabric automation more practical and reliable.




OneLake



OneLake file explorer (Generally Available)




Data Engineering



With Custom Live Pools, Fabric Data Engineering now enables you to:

- Create dedicated Spark pools and schedule them tailored to your workload needs.
- Reduce session startup overhead by keeping pools warm with libraries preinstalled.
- Run reliably within Managed VNets and Private Link-enabled environments.



Z-order and liquid clustering support in the Native Execution Engine


With the Native Execution Engine, Fabric Data Engineering continues to raise the bar on price-performance leadership for large-scale analytics. Beyond execution-time optimizations, the engine now includes native support for Z-Order and Liquid Clustering, allowing advanced data layout techniques to fully benefit from vectorized, C++-based execution paths.

Dynamic session sharing limit up to 50 for high concurrency

Fabric High Concurrency Spark sessions enable both interactive exploration and large-scale, pipeline-driven notebook execution, supporting parallel, scheduled, and event-driven workloads at enterprise scale.


Customers often achieve higher density by packing notebooks into a shared High Concurrency (HC) session using session tags, effectively fitting up to five notebooks per session to control startup overhead and cost. While effective, this approach relies on static limits and manual tuning.

With this update, Fabric Data Engineering allows the maximum number of notebooks attached to a High Concurrency session to be increased up to 50, enabling dynamic session sharing at much higher scale.



Support for multiple schedules in Fabric materialized lake views

MLVs now support multiple named schedules per lakehouse. Previously, all MLVs shared a single schedule, and teams needing different refresh timings resorted to notebook-triggered refreshes. This workaround bypasses dependency management, centralized error reporting, and retry logic; failures can persist for weeks undetected. Each named schedule now targets a specific subset of views. A finance pipeline can refresh hourly while an analytics pipeline refreshes daily. When a schedule fires, Fabric refreshes the views in parallel, surfaces errors centrally, and sends notifications.



PySpark support for Fabric materialized lake views (Preview)

MLVs now support PySpark authoring (Preview), letting data engineers create, refresh, and replace MLVs from Fabric notebooks using the DataFrameWriter API. Previously, teams wrote definitions in Spark SQL, which made custom cleansing logic, UDFs for business rules, and procedural transformations harder to express. With PySpark authoring, MLVs gain access to the entire Python ecosystem.


A gold-layer MLV can score transactions against a fraud detection model, standardize addresses using a geocoding library, or validate records against external regulatory rules. All existing MLV capabilities, including data quality constraints, table properties, and scheduled refreshes, work identically with PySpark-authored definitions. Full refresh only today; optimal refresh is coming soon.

Python


Copy

```
from pyspark.sql import functions as F


# Simple helper function: concatenate columns as string
def concat_name_age(df):
    return df.withColumn(
        "name_age",
        F.concat(F.coalesce(F.col("name"), F.lit("")),
                  F.lit("-"),
                  F.coalesce(F.col("age").cast("string"), F.lit("")))
    )
```



Data Science and AI



Data Agent in GA



Advanced security and governance in data agents (Preview)


Data agents in Microsoft Fabric (Preview) include capabilities that elevate security and governance standards. Through integration with Microsoft Purview, organizations gain access to comprehensive auditing, eDiscovery, data lifecycle management, communications compliance, and classification. These tools capture prompt and response telemetry along with user context, supporting enterprise protection and regulatory compliance.

Additionally, we are introducing outbound access protection support for the Data Agent artifact to help mitigate sensitive data exfiltration risks and adhere to strict security policies at the individual workspace level.


With these updates, organizations can monitor, control, and safeguard all data agent interactions more effectively.

Data source enhancements for data agents (Preview)

The latest preview brings significant enhancements to data agent source capabilities in Microsoft Fabric. Users can now connect **Graph as a data source**, enabling them to model and analyze complex relationships within their data for richer, AI-driven insights. Additionally, support for KQL user-defined functions (UDFs) and SQL functions is available, allowing for more sophisticated and efficient querying in KQL- and SQL-enabled sources. These enhancements make data agents more flexible and powerful, supporting faster analytics and expanded scenario coverage.



Warehouse



COPY INTO and OPENROWSET support for OneLake sources (Generally Available)

Previously, this capability supported Lakehouse sources only. With this release, we are expanding support to all OneLake items (except Warehouses).

This enables much more flexible ingestion scenarios, including:

- Using partner workloads such as COPY Jobs.
- Using staging areas across different Fabric items.
- Loading data stored anywhere in OneLake-backed it

Outbound Access Protection (OAP) support for Warehouses

Outbound Access Protection for Fabric Data Warehouse provides protection for enterprise environments.

Warehouse now supports connector rules that allow only approved warehouse can access.

Customers can define rules to allow access to:


- Specific Azure Data Lake Storage Gen2 accounts
- Other Fabric workspaces
- Approved external connectors

Fabric Data Warehouse recovery (Preview)

In fast moving production environments, an item gets dropped accidentally due to an incorrect script. Suddenly, critical reports are broken, and teams are asking the same question: *"How fast can we recover?"*


With dropped warehouse recovery in Microsoft Fabric, a deleted warehouse no longer means starting over.

You can now restore a dropped warehouse together with everything that goes with it—data, schemas, snapshots, permissions, and saved queries—in minutes, without rebuilds, re-ingestion, or complex restore workflows.




SQL Audit Logs (Generally Available)

SQL Audit Logs for Fabric Data Warehouse enable organizations to capture and analyze database activity for security monitoring, compliance, and forensic analysis.



Analyze unstructured text using T-SQL AI functions (Preview)



SQL pools (preview)

SQL pools divide your compute for query processing. Query runs are assigned to pools using a classifier. Choose the setup you want: built-in (automatic) or custom (you control the compute divisions and allocations) for all warehouses and SQL analytics endpoints in your workspace.

[Learn more about SQL pools](#)

SQL pool setup **Using custom setup**

[Revert to built-in setup](#)

The built-in setup divides your compute evenly into two pools. Query runs are classified (assigned) to pools based on the query statement type (SELECT vs non-SELECT). To change the classifier (from "query statement type" to "query submitting application") and divide your compute differently, set up custom pools.

[Save](#) [+ New SQL pool](#) [Delete](#) [Classifier type](#)

Workspace SKU is FT1. 100% of compute resources are allocated.



■ Compute for **Adhoc**: 10% ■ Compute for **Pipelines**: 40%
■ Compute for **Reports**: 50% ■ Unused: 0%


☐	Pool name	App name (exact name) ⓘ	Compute %	Optimize for reads	Default ▲
☐	Adhoc		10%	☐	☒
☐	Pipelines	ETL ×	40%	☐	☐
☐	Reports	Daily Reports × Exec Dashboard ×	50%	☒	☐



PBIUG ITALY



Databases




HUB and SQL DB

Introducing the Database Hub in Fabric: An agentic control plane for your database estate

As data estates grow, so does operational complexity. Most organizations manage a mix of relational and NoSQL databases across edge, PaaS, and SaaS environments, often through fragmented tools, portals, and management experiences. As AI places even greater demands on data estates, unifying databases under a single access point and control plane has become essential.

Microsoft Fabric is expanding its role as the central access point for enterprise data to help address this challenge. We are introducing the Database Hub in Fabric, now available in [early access](#), as a unified database management experience that brings together databases across edge, PaaS, and Fabric into a single, coherent view. Database teams now have one place to explore, observe, govern, and optimize their entire database estate, including Azure SQL, Azure Cosmos DB, Azure Database for PostgreSQL, SQL Server (enabled by Azure Arc), Azure Database for MySQL, and Fabric databases, without changing how each service is deployed.





SQL database in Fabric


Since reaching general availability in November 2025, SQL database in Fabric has seen rapid customer adoption as organizations modernize SQL workloads with less operational overhead and tighter integration with analytics and AI. Guided by customer feedback, the platform emphasizes simplicity, autonomy, security, and AI optimization. We are introducing a set of improvements and new features that make it easier to migrate, manage, and optimize SQL workloads in Fabric:

- **Simplified migration with new assistant:** The Migration Assistant in public preview helps SQL developers move SQL Server and Azure SQL workloads into Fabric by importing schemas, assessing compatibility, and guiding migration with minimal manual effort.
- **Configurable autonomous management:** While maintaining a SaaS-first approach, new options allow database-level control over vCore scaling, expanded compatibility levels, enhanced T-SQL features, and settings that ease application transitions without code changes.
- **Support for all collations:** All Azure SQL database collations are supported when creating a new SQL database in Fabric for enhanced global data compatibility and app development flexibility. Collations control text sorting and comparison in SQL databases, impacting filtering, searches, and multilingual content management. Users can specify collations seamlessly during database creation via the REST API across deployment methods.

Check out the [How to set a different collation for SQL database in Fabric](#) demo and explore the [sample code in Git repo](#).



Data Factory



Copy Job





Richer Change Data Capture (CDC) with Oracle, Fabric DW, and SCD Type 2

Copy job continues to improve the no-code CDC experience with richer, enterprise-ready replication patterns. This release introduces a set of enhancements in Copy job in Microsoft Fabric Data Factory that make CDC replication more powerful and easier to use without writing code:

- **Oracle CDC source**—Capture changes directly from Oracle databases.
- **Fabric Data Warehouse sink**—Replicate CDC data into Fabric Data Warehouse.
- **SCD Type 2**—Preserve full history with valid dating, and handle deletes as soft deletes.

With built-in SCD Type 2 and soft delete handling, Copy job automatically preserves every version of a record as it changes over time, instead of overwriting history. This makes it easy to answer point-in-time questions, support regulatory audits, and run accurate historical analytics—capabilities that traditionally require complex MERGE logic or custom code.



Dataflow Gen2



Fabric Variable Library integration (Generally Available)

Throughout the last couple of months, we've removed some of the limitations such as:


- **Variable limit:** the previous limit was 50 variables. You can now reference as many variables as you need in your Dataflow.
- **Power Query editor support and using a default value:** you can now see how the variable gets evaluated within the actual Power Query editor.

Not only that but within the Dataflow Gen2 experience, after enabling the input widgets through the Options menu, you will also have a new way that simplifies how you can reference variables with a complete no code experience fully integrated into the experiences and dialogs that you love from Dataflow Gen2:

New data destinations

Dataflow Gen2 continues to expand where curated data can land, supporting both lake-first architectures and hybrid data estates. With new and updated destinations, teams can publish transformed outputs in the formats and platforms that best fit their downstream consumers—whether that's open lake storage, lakehouse files, enterprise warehouses, or business-friendly file formats.

The screenshot shows the 'Filter rows' dialog box in Power Query. It has a title bar 'Filter rows' with a help icon. Below the title, it says 'Apply one or more filter conditions to the rows in this table.' There are two radio buttons: 'Basic' (selected) and 'Advanced'. Under 'Basic', there are two filter conditions. The first condition is 'Keep rows where "Column1" is greater than or equal to' followed by a dropdown menu showing 'Enter a value'. The second condition is 'is less than or equal to' followed by a dropdown menu showing 'Enter a value'. There are also radio buttons for 'and' (selected) and 'or'. At the bottom, there are 'OK' and 'Cancel' buttons.



Data Pipeline

Virtual network data gateway supports up to nine instances

This update enables greater scalability and higher throughput for enterprise workloads. With expanded instance capacity, organizations can handle increased data movement and processing, improve parallel job performance, and enhance reliability for mission-critical data pipelines. It also provides more flexibility to scale gateway infrastructure in line with growing business requirements.

Advanced options

High availability and load balancing for virtual network data gateways. [Learn more.](#)

Time interval of inactivity before auto-pause



30 minutes

Number of member gateways

9

SSIS Pipeline Activity (Preview)

SQL Server Integration Services (SSIS) has been a cornerstone of enterprise data integration for decades, powering mission-critical ETL workloads across thousands of organizations worldwide. Invoke SSIS Package activity in Data Factory in Microsoft Fabric (Preview), provides the power of your existing SSIS investments directly into Fabric's unified SaaS analytics platform.



Mirroring



Mirroring for SharePoint List (Preview)

Mirroring for SharePoint Lists enables continuous replication of SharePoint Lists and Document Libraries into OneLake without building custom ETL pipelines. This capability keeps SharePoint data automatically synchronized in near real time, ensuring analytics in Fabric stay aligned while SharePoint remains the system of record. When mirrored, both list tables and document library metadata land in OneLake in an analytics-ready format, with document libraries replicated via shortcuts and converted into Delta Lake tables.

New mirrored SharePoint Online List

New source



Connection settings

Site URL * ⓘ

Implementation

> Advanced options


Connection credentials

Connection
 ⓘ


Connection name

Authentication kind

You are currently signed in as:



Power BI



Dataflow Gen1



ANNOUNCEMENTS

Dataflows: Thank you for eight years of Gen1—and why Gen2 is the future



Miguel Llopis

Group Program Manager - Data Integration

📅 April 2, 2026

[in](#) Share


[🐦](#) Tweet

[f](#) Like


For more than eight years, customers have relied on **Power BI Dataflows (Gen1)** as a core part of their analytics solutions. We're grateful for the trust you placed in the Power Query experience to build reusable, low-code data preparation pipelines that power reports, semantic models, and downstream analytics.

Now, we're sharing an update on the future of Dataflows.

Dataflows Gen2 builds on everything you know from Gen1—preserving the familiar Power Query authoring experience—while delivering major improvements in scale, flexibility, cost efficiency, and manageability. Going forward, **all new Dataflow innovation will land only in Dataflows Gen2.**



Q&A



Thanks

